

# Information Retrieval

Deep Learning with Pytorch

# PyTorch

---

- ❖ PyTorch is a powerful, yet easy-to-use mathematical library with automatic differentiation for Python, mainly used in Deep Learning.
  - PyTorch has a Python front-end component that offers a high-level programmable interface to users, while most of the code runs on a backend in C/C++ (CUDA/MPS/ROCm) for maximum efficiency.

# PyTorch - Tensors

---

- ❖ The library itself contains multiple abstraction layers, where the most low-level data representation is called a **tensor**, on which multiple mathematical operations can be performed.
  - Tensors are a specialized data structure that are very similar to arrays and matrices.

```
1 x = [0,1,2,3,4]
```

```
2 x_tensor = torch.tensor(x)
```

```
3 x_tensor, x_tensor.shape, x_tensor.dtype, x_tensor.device
```

```
(tensor([0, 1, 2, 3, 4]), torch.Size([5]), torch.int64, device(type='cpu'))
```

# PyTorch - Tensors

---

- ❖ Tensors can be allocated or sent to any compute device available in your system.

```
x_cuda_tensor = x_tensor.to("cuda")  
x_cuda_tensor, x_cuda_tensor.shape, x_cuda_tensor.dtype, x_cuda_tensor.device  
  
(tensor([0, 1, 2, 3, 4], device='cuda:0'),  
 torch.Size([5]),  
 torch.int64,  
 device(type='cuda', index=0))
```

# PyTorch - Tensors

- ❖ Tensors can be allocated or sent to any compute device available in your system.

```
big_matrix = torch.rand((10000,1000))  
big_matrix_gpu = big_matrix.to("cuda")
```

On CPU

```
%time big_matrix @ big_matrix.T
```

```
CPU times: user 2.44 s, sys: 178 ms, total: 2.62 s  
Wall time: 2.72 s
```

On GPU

```
%time big_matrix_gpu @ big_matrix_gpu.T
```

```
CPU times: user 975 µs, sys: 993 µs, total: 1.97 ms  
Wall time: 2.06 ms
```

# PyTorch – Neural Networks

- ❖ Neural networks are composed of a set of layers/modules that perform operations on data.
  - PyTorch `torch.nn` namespace provides all the building blocks you need to build your own neural network.
- ❖ Every layer/module in PyTorch subclasses the `nn.Module`. A neural network is a module itself that consists of other modules (layers).

## torch.nn

These are the basic building blocks for graphs:

### torch.nn

- Containers
- Convolution Layers
- Pooling layers
- Padding Layers
- Non-linear Activations (weighted sum, nonlinearity)
- Non-linear Activations (other)
- Normalization Layers
- Recurrent Layers
- Transformer Layers
- Linear Layers
- Dropout Layers
- Sparse Layers
- Distance Functions
- Loss Functions
- Vision Layers
- Shuffle Layers
- DataParallel Layers (multi-GPU, distributed)
- Utilities
- Quantized Functions
- Lazy Modules Initialization
  - Aliases

# PyTorch – Neural Networks

## ❖ With special interest for IR:

- Embedding Layer (e.g. `torch.nn.Embedding`)
- Convolution Layers (e.g. `torch.nn.Conv2d`)
- Recurrent Layers (e.g. `torch.nn.LSTM`)
- Linear Layers (e.g. `torch.nn.Linear`)
- Non-linear Activations (e.g. `torch.nn.ReLU`)

## torch.nn

These are the basic building blocks for graphs:

### torch.nn

- Containers
- Convolution Layers
- Pooling layers
- Padding Layers
- Non-linear Activations (weighted sum, nonlinearity)
- Non-linear Activations (other)
- Normalization Layers
- Recurrent Layers
- Transformer Layers
- Linear Layers
- Dropout Layers
- Sparse Layers
- Distance Functions
- Loss Functions
- Vision Layers
- Shuffle Layers
- DataParallel Layers (multi-GPU, distributed)
- Utilities
- Quantized Functions
- Lazy Modules Initialization
  - Aliases

# PyTorch – Neural Networks

---

```
class AbstractModel(torch.nn.Module):  
  
    def __init__(self):  
        super().__init__()  
        # Instantiation of layers and parameters  
  
    def forward(self, x):  
        # operations that use the previous created layers over the inputs  
        logits = x  
        return logits
```



# PyTorch – Neural Networks

---

```
class AbstractModel(torch.nn.Module):
```

```
def __init__(self):  
    super().__init__()  
    # weights  
    self.weights = torch.nn.Parameter(torch.tensor([2, 4, 6]))
```

```
def forward(self, x):  
    # move the model to the GPU  
    self.model_on_gpu = self.to("cuda")  
    # move the inputs to the GPU  
    self.model_on_gpu(torch.tensor([1, 2, 3]).to("cuda"))  
    return self.model_on_gpu
```

the inputs

# PyTorch – Dataset and DataLoaders

- ❖ PyTorch also offers a high-level API for handling data processing and for feeding data to your model.
  - `torch.utils.data.Dataset`
  - `torch.utils.data.DataLoader`
- ❖ Dataset class acts as an abstraction that holds your data and knows how to fetch it.

```
class SimpleDataset(torch.utils.data.Dataset):
    def __init__(self, num_samples=1000):
        super().__init__()
        # load your data here

        # dummy data
        self.features = [torch.rand(1, random.randint(3, 5)) for _ in range(num_samples)]
        self.labels = [float(torch.mean(f) > 0.5) for f in self.features]

    def __len__(self):
        # return the number of samples in the dataset
        return len(self.features)

    def __getitem__(self, idx):
        # get the sample corresponding to index "idx"
        return self.features[idx], self.labels[idx]

ds = SimpleDataset()
print(ds[10])
print(ds[31])

(tensor([[0.5195, 0.7827, 0.2147, 0.2094, 0.7416]]), 0.0)
(tensor([[0.6587, 0.7761, 0.9588, 0.7888]]), 1.0)
```

# PyTorch – Dataset and DataLoaders

❖ The DataLoader offers an abstraction that effectively selects which sample will be fed to the model.

- Selects a batch of data
- Converts that data to tensors
- Repeat

```
def collate_fn(batch):
    features, labels = zip(*batch)
    max_len = max(f.size(1) for f in features)
    padded = [torch.cat([f, torch.zeros(1, max_len - f.size(1))], dim=1) for f in features]
    return torch.cat(padded), torch.tensor(labels)

# Create dataloader
dataloader = torch.utils.data.DataLoader(
    ds,
    batch_size=4,
    shuffle=True,
    collate_fn=collate_fn
)

next(iter(dataloader))

(tensor([[0.7736, 0.3263, 0.6309, 0.0000],
        [0.0070, 0.0373, 0.2937, 0.0000],
        [0.1397, 0.7423, 0.2542, 0.0000],
        [0.4568, 0.6127, 0.1965, 0.0904]]),
 tensor([1., 0., 0., 0.]))
```

# Goal for Today!!!

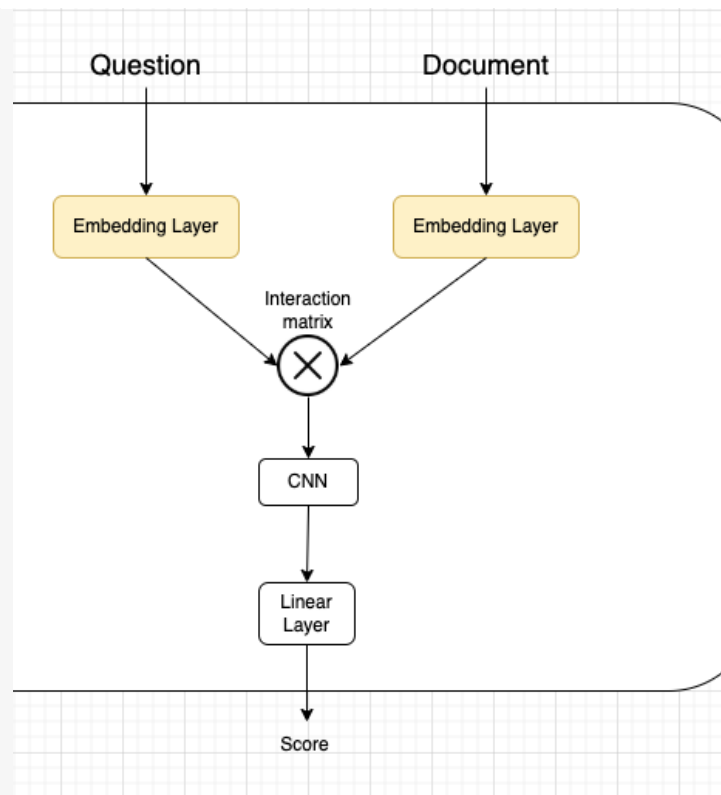
❖ Implement the following interaction-based Information Retrieval Model

```
class CNNInteractionBasedModel(torch.nn.Module):
```

```
def __init__(self, vocab_size):
    super().__init__()
    self.vocab_size = vocab_size
    # Instantiation of layers and parameters
    # - embedding layer
    # - Convolution and pooling layers
    # - Activation functions
    # - Linear layer for scoring

def forward(self, query, document):
    print(query, document)
    # flow of the computation
    # - convert the query and document ids to document
    #   vectors with the embedding layer
    # - create an interaction matrix
    # - apply convolution and max pooling over the matrix
    # - apply linear layer to the resulting feature maps
    # - return the final logits
    # - optionally convert the logits to probabilities with sigmoid function

    return logits
```



---

❖ Start code:

<https://colab.research.google.com/drive/1HpQbXQ5be3MQCOtYQjntCjDeioGXwply?usp=sharing>