

Evolving Piranha CMS for ”Contents’R’Us”:

Event-Driven Extensions & Secure Integrations

Daniel Madureira - 107603

João Andrade - 107969

José Gameiro - 108840

Tomás Victal - 109018

Professor Cláudio Teixeira

Software Architectures

DETI

Universidade de Aveiro

May 2025

Contents

1	Introduction	3
1.1	System Analysis	3
1.2	Project vision	3
2	Architectural Design Methodology	5
2.1	Advantages	5
2.2	Alternatives	5
2.3	Application of ADD in Our Context	6
2.3.1	Step 1: Identify Architectural Drivers	6
2.3.2	Step 2: Establish Iteration Goal	6
2.3.3	Step 3: Choose Architectural Elements to Decompose	6
2.3.4	Step 4: Select Design Concepts	7
2.3.5	Step 5: Instantiate Architectural Elements	7
2.3.6	Step 6: Verify and Refine Requirements	7
2.3.7	Step 7: Repeat Steps 2-6 for Next Element	7
3	Design Iterations	8
3.1	Iteration 0	8
3.1.1	Functional Requirements	8
3.1.2	Quality Attributes	8
3.1.3	System Constraints	9
3.2	Iteration 1	9
3.3	Iteration 2	11
3.4	Iteration 3	12
3.5	Iteration 4	13
4	Core Domains	16
4.1	Core Domains	16
4.1.1	Event Publishing Domain	16
4.1.2	Subscription Management Domain	17
4.1.3	External Integration and Inbound Event Handling Domain	18
4.1.4	Security and Access Control Domain	18
4.2	Supporting Subdomains	19
4.2.1	Monitoring & Observability Subdomain	19
4.3	Generic Subdomains	20
4.3.1	User Management Subdomain	20
4.3.2	CMS Platform	21
5	Cross-Cutting Concerns	22
5.1	Security	22
5.2	Logging and Auditing	22
5.3	Error Handling and Resilience	23
5.4	Observability and Monitoring	23
5.5	Performance and Scalability	24
5.6	Compliance and Data Protection	24

6	Proposed Architecture and Roadmap	25
6.1	The Components and their responsibilities	26
6.1.1	Controllers	27
6.1.2	Webhooks Dispatcher	27
6.1.3	Subscriptions Manager	27
6.1.4	Publish Receiver	27
6.1.5	Message Broker	27
6.1.6	Event Manager Service	28
6.1.7	Piranha CMS Core	28
6.1.8	Key Manager	28
6.2	Communication/Interfaces Between Components	29
6.3	Road Map	29
7	Implementation and Demonstration	31
7.1	The Road-map	31
7.2	Implementation	31
7.2.1	Tags	31
7.2.2	Subscription Management Module	32
7.2.3	Event Publishing Infrastructure	37
7.2.4	Key management module	39
7.3	External Publishing	42
7.3.1	Webhook Delivery System	44
7.3.2	Demo Web Application	46
7.4	Architecture	49

1 Introduction

1.1 System Analysis

Piranha CMS is a lightweight, cross-platform content management system built on **.NET Core**. It provides essential CMS features such as modular content types, intuitive page editing, media management, and extensibility through plugins and APIs. Its architecture is clean, and the core design principles favor modularity and developer-friendliness, making it a solid foundation for modern applications.

However, as we evaluate the CMS in the context of Scenario 3: Event-Driven Extensions and Secure Integrations, several critical gaps emerge:

- **Lack of Event Infrastructure:** The CMS does not natively support domain event emission (e.g., on content creation or deletion). There is no built-in event bus or observable mechanism for internal services or third parties to subscribe to.
- **No Inbound Event Handling:** Piranha CMS is designed as a pull-based system. It cannot receive or react to externally pushed events.
- **Security Limitations for Integrations:** Although Piranha provides some API-level extensibility, it lacks support for secure, configurable integration endpoints. There are no native mechanisms for key management, event signing, or authentication of external publishers.
- **Synchronous Design Bias:** The current design presumes user-initiated operations and lacks asynchronous processing capabilities (e.g., message queues, background event handling).

These limitations are in conflict with the requirements of a modern CMS that must operate in complex ecosystems involving CRMs, marketing automation tools, analytics platforms, and third-party content pipelines.

1.2 Project vision

The vision of this initiative is to transform Piranha CMS into a modern, asynchronous content platform that is integration-ready and resilient. The evolved system will act not just as a content manager, but also provide a real-time event-driven ecosystem.

This transformation will be driven by the principles of loose coupling, secure communication, and configurability, allowing external systems to push and pull data dynamically while maintaining trust, integrity, and performance.

With this in mind, three main strategic goals have been defined:

1. **Custom Publish/Subscribe Mechanism** - Introduce a flexible domain eventing system that allows administrators to configure which models (e.g., pages, posts, media) emit events and under what circumstances. The system must also support outbound event formatting, topic-based filtering, and dynamic subscription registration, signing, and validation.
2. **Inbound Integration via External Publishers** - Enable Piranha CMS to receive and process events from authorized third-party systems. This requires a configuration layer to register publishers, define accepted message types, and map incoming events to internal workflows or commands (e.g., content updates, status changes).

3. **Secure, Configurable Messaging Infrastructure** - Establish robust security practices across all event flows. This includes:

- Key and token management for authentication;
- Message signing and verification to ensure integrity;
- Endpoint configuration interfaces for defining allowed targets and actions;
- Logging and audit trails for traceability.

2 Architectural Design Methodology

In order to evolve the Piranha CMS into an event-driven platform with secure integration, we choose the Attribute-Driven Design (ADD) as our primary architectural design methodology.

2.1 Advantages

The Attribute-Driven Design provides several advantages that aligns specifically well with our chosen scenario.

ADD focuses on **quality attributes** in architectural decision-making, which corresponds directly with our emphasis on performance, scalability, security, and configurability. By placing these architectural drivers at the center of the design process, ADD enables the construction of a system that is not only functional but also robust, scalable, and future-proof.

Furthermore, ADD's **top-down** and **iterative** approach encourages gradual decomposition of functionality based on concrete scenarios, starting with high-level structures and progressively refining them through multiple iterations. This aligns well with evolving an existing system like Piranha CMS, where we must carefully integrate new capabilities while preserving existing functionality.

ADD provides a catalog of **architectural tactics** that directly address specific quality attributes. These proven design approaches offer valuable guidance for implementing event-driven patterns, security mechanisms, and scalability strategies.

Moreover, ADD complements Domain-Driven Design principles, allowing us to incorporate domain modeling while maintaining a focus on architectural quality attributes. This combination will help ensure that our event system aligns with meaningful business domains.

On top of all these benefits, this methodology also emphasizes documenting both the decisions made and the rationale behind them, creating a valuable reference for future development and maintenance.

2.2 Alternatives

When choosing the right architecture design methodology we considered some other alternatives that we later discarded. One of the discarded options is TOGAF (The Open Group Architecture Framework).

TOGAF is a well-established framework, especially in enterprise environments, where long-term planning, organizational modeling, and governance are priorities. However, for a focused, technical implementation like ours, centered on integrating specific components into an existing CMS, it turned out to be too broad and high-level. Using TOGAF would've added a lot of overhead in terms of documentation and formal process, without really helping us solve the technical challenges at hand.

This framework is more useful when working on large-scale systems across multiple departments or organizations, and since we were working on a tightly scoped extension to an open-source CMS, TOGAF felt like overkill. It simply didn't offer the hands-on tools we needed to handle things like validating events, managing retries, or isolating faults in a microservice-style design.

We also looked at ACDM (Architecture-Centric Design Method). It has a more structured, step-by-step approach that focuses on formal processes and repeated cycles of design and validation. ACDM is great for projects in highly regulated industries or systems that need rigorous documentation and oversight. But our project needed to move fast and stay flexible. We had to be able to experiment—try out different messaging flows, plug in new modules, and quickly adapt to feedback. ACDM, with its strict phases and heavier process, would have slowed us down. It also doesn't focus as much on designing around key quality attributes like resilience or extensibility, which were top priorities for us.

In the end, we chose ADD because it gave us the best of both worlds: enough structure to keep our architecture sound, but enough flexibility to evolve quickly. It let us design around real scenarios and quality goals without getting bogged down in bureaucracy, and that made it the right fit for building an event-driven extension that works seamlessly with the existing CMS.

2.3 Application of ADD in Our Context

We will apply the Attribute-Driven Design methodology through the following process:

2.3.1 Step 1: Identify Architectural Drivers

The first step involves identifying the key architectural drivers that will shape our design decisions. These include:

- Primary functionality requirements derived from the strategic goals
- Critical quality attributes identified in our system analysis
- Project constraints, including compatibility with the existing Piranha CMS architecture
- Technical environment considerations, such as deployment platforms and integration technologies

2.3.2 Step 2: Establish Iteration Goal

For each iteration of the design process, we will establish specific decomposition goals. Our initial iterations will focus on the high-level system structure, particularly the event infrastructure and integration gateways. Subsequent iterations will refine component designs and interface specifications.

2.3.3 Step 3: Choose Architectural Elements to Decompose

Based on the iteration goal, we will select specific elements for further decomposition. This process will begin with the core system modules and progressively refine their internal structures and relationships.

2.3.4 Step 4: Select Design Concepts

For each element being decomposed, we will select appropriate design concepts based on our architectural drivers. These concepts will include patterns such as:

- Event-driven architecture patterns for the publish/subscribe system
- Separation of concerns using a layered approach

2.3.5 Step 5: Instantiate Architectural Elements

We will instantiate specific architectural elements by applying the selected design concepts. This includes defining components, connectors, interfaces, and their properties. For each element, we will specify responsibilities, relationships, and quality attribute characteristics.

2.3.6 Step 6: Verify and Refine Requirements

After each iteration, we will verify that the decomposed elements satisfy the architectural drivers and refine our understanding of requirements as needed. This feedback loop ensures that our architecture remains aligned with project goals throughout the design process.

2.3.7 Step 7: Repeat Steps 2-6 for Next Element

We will repeat this process for each architectural element, progressively refining the overall system design through multiple iterations.

3 Design Iterations

Following the ADD methodology, our architecture was developed through five main iterations, which are described below.

3.1 Iteration 0

To kickstart the ADD process, we first needed to figure what our primary functionalities were as well as our constraints, quality attributes and architectural concerns. On the following tables we'll expose our main functional requirements, quality attributes and constraints, and explore how they evolved alongside the iterative process.

3.1.1 Functional Requirements

To better organize the functional requirements, we identified two primary system actors: the **Administrator** and the **External Publisher**, each associated with specific responsibilities and interactions within the system.

ID	Requirement	Actor(s)
FR-1	Define possible actions for external publishers	Admin
FR-2	Publish new content	External Publisher
FR-3	See all the existing events	Admin
FR-4	Decide if a given content will be published	Admin
FR-5	Subscribe and unsubscribe to existing events	External Publisher
FR-6	Define authorized external publishers	Admin
FR-7	See new published events	Admin & External Publisher
FR-8	Needs to integrate easily with external services	Admin & External Publisher

3.1.2 Quality Attributes

Our quality attributes focus on 9 main categories that are presented below:

ID	Attribute	Category
QA-1	Event processing (publishing and inbound) must occur with a maximum latency of 1 second under normal load	Performance
QA-2	The system must support up to 10,000 concurrent event subscriptions and 1,000 inbound events per minute without service degradation	Scalability
QA-3	All event communications (inbound and outbound) must be authenticated and signed to ensure integrity and authenticity	Security
QA-4	The system should achieve 99% up time for the Event Manager and Inbound Handler modules	Availability
QA-5	Code for new modules should follow clean architecture principles and allow easy extension	Maintainability

ID	Attribute	Category
QA-6	If an event fails to deliver, the system should try again up to three times with increasing delays before sending it to a dead letter queue.	Reliability
QA-7	The admin configuration interface must allow an event to be fully configured (published or subscribed) within 5 clicks	Usability
QA-8	All event transactions (successful, failed, retried) must be logged with timestamp, event type, and endpoint information	Logging & Monitoring
QA-9	If sensitive data is part of the payload, messages must comply with GDPR (i.e., encryption at rest and in transit)	Compliance

3.1.3 System Constraints

The table below summarizes the key technical and organizational constraints that guided our system design.

ID	Constraint
CON-1	Our work must be based on the existing Piranha CMS (.NET Core, C#) framework
CON-2	Must use open-source or .NET-compatible libraries (no paid 3rd-party tools without prior approval)

With these artifacts in mind, we were finally able to fully start the iterative process within the ADD methodology.

3.2 Iteration 1

The initial design focuses on introducing a foundational event-driven architecture layered on top of the existing Piranha CMS core. This version introduces several new components to support publish/subscribe functionality.

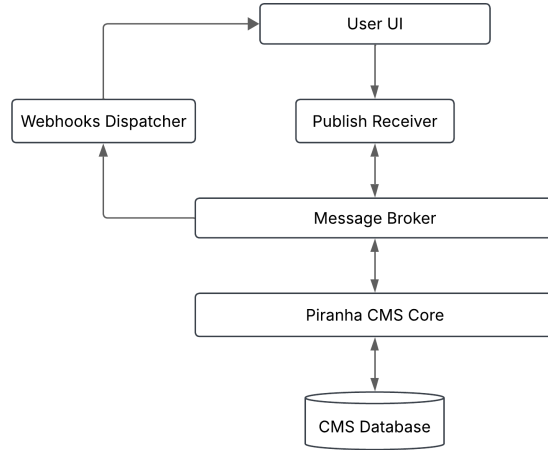


Figure 1: Iteration 1 Architecture

The core components in this iteration include the following:

- **User UI:** Acts as the primary interface through which users trigger content actions (like publishing, updating or deleting content);
- **Publisher Receiver:** Captures publishing events from the UI and passes them to the message processing system;
- **Message Broker:** A decoupled messaging layer that asynchronously routes events between components, ensuring scalable and reliable delivery;
- **Piranha CMS Core:** The existing CMS logic which processes the event data and updates the database;
- **CMS Database:** The persistent storage used by Piranha CMS;
- **WebHooks Dispatcher:** Handles outbound events by forwarding messages to external subscribers (third-party systems).

This iteration focuses on basic decoupling and allows the system to be observed externally, referencing WebHooks as described by Microsoft[1].

Limitations and Observations

With this initial architecture in mind, our team found some aspects that were missing:

- **Security:** While the architecture introduces integration points for external services, it does not yet address message integrity, endpoint authentication, or permission controls;
- **Asynchronous layer:** After taking another look at our quality attributes, we noticed that this architecture can't handle asynchronous event delivery, as it needs that the event consumers to be online to receive the said event;
- **Subscription Management:** A component is missing to allow for management subscriptions. Without it, system administrators or external clients have no way to configure which events they want to subscribe to;

- **No User-Defined Configurations:** There isn't a way for administrators to configure endpoints, keys, or choose which models should publish or subscribe to events.

3.3 Iteration 2

In our next iteration, we focused on the **asynchronism** of the system and on **subscription management**. This step introduced the logic to handle event routing when subscribers are offline, enabling the system to operate asynchronously. After that, we also reviewed the event flow from the previous iteration and improved it by applying the publish/subscribe (pub/sub) pattern.

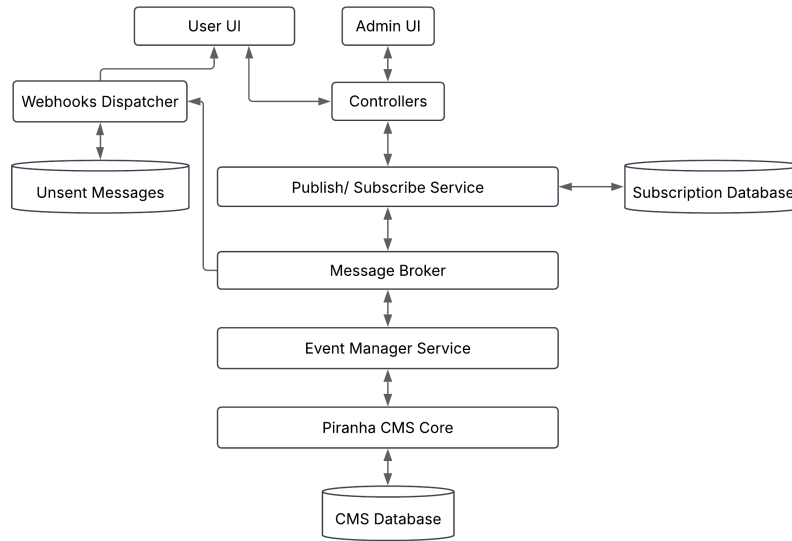


Figure 2: Iteration 2 Architecture

With these goals in mind, this iteration introduced some core components to our architecture:

- **Message Broker:** This component acts as a central channel for communication between producers (like the CMS core or external event generators) and consumers (such as webhook dispatchers or subscription managers). It abstracts the complexities of direct communication and enables asynchronous messaging.
- **Publish/Subscribe Service:** This service registers consumers interested in specific event types and routes messages accordingly. This introduces a scalable pattern where multiple services can react to CMS events without tight coupling to the CMS logic.
- **Subscription Database:** A new persistent storage layer is introduced to keep track of subscriber preferences. This enables dynamic registration and ensures resilience during service restarts or failures.
- **Controllers and UI Refinement:** The controller layer is now formally introduced to interface between the Admin/User UIs and the publish/subscribe layer, ensuring a clear separation between presentation and logic layers.

- **Webhooks Dispatcher Integration:** The webhooks dispatcher now consumes messages from the broker. If a webhook cannot be sent, it is stored in the Unsent Messages database for retries, introducing reliability and fault-tolerance.

With these new components we were able to introduce some architectural patterns and reinforce some others. On one hand, the broker pattern ensures the loose coupling and the asynchronous processing of the events while the publish/subscribe pattern allows for extensibility, enabling easy on-boarding of new consumer services such as analytics, monitoring, or real-time notifications.

Limitations and Observations

Now with these changes applied to the architecture our group still noticed that some components were missing:

- **Security:** A security module was still not addressed to evaluate the integrity of all the messages exchanged, endpoint authentication or permission controls;
- **User-Defined Configurations:** Another limitation that was also not addressed in this iteration was the lack of allowing administrators to configure endpoints, keys, or choose which models should be published or subscribe to events.

3.4 Iteration 3

For the third iteration, we focused on improving the publish/subscribe system by decoupling responsibilities into distinct services and establishing internal communication protocols. The goal was to enhance maintainability, enable more granular control over event flow, and begin addressing foundational security concerns.

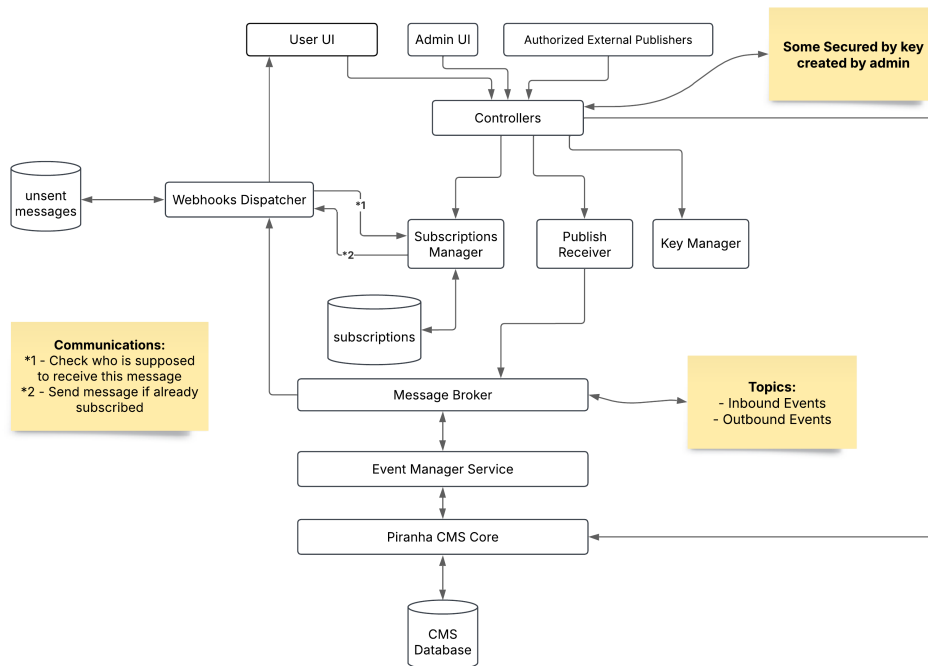


Figure 3: Iteration 3 Architecture

As we refined the architecture, we added some more core components:

- **Service Separation:** The publish and subscribe functionalities are now handled by separate services. This separation enables independent scaling, clearer responsibilities, and easier evolution of each component.
- **Subscription Manager & Webhooks Dispatcher Integration:** A communication channel was established between the Subscription Manager and the Webhooks Dispatcher. This allows the dispatcher to:
 - Know which subscribers should receive messages for a given event.
 - Receive notifications when a previously subscribed user reconnects, triggering the dispatch of any unsent messages.
- **Unsent Message Handling:** With the new coordination between services, unsent messages can now be reliably delivered when conditions allow (e.g., subscriber becomes available again), improving delivery guarantees.
- **Authentication Service:** A new service was introduced to manage security credentials. It supports:
 - Generating authentication keys for publishers.
 - Validating these keys before allowing message publication.

This iteration improved system modularity through clearer service boundaries and enhanced reliability by integrating subscription management with message delivery. Early authentication mechanisms were introduced to support secure event publishing, and the system became more resilient to temporary delivery failures, ensuring better service for subscribers.

Limitations

While this iteration introduced key architectural improvements, the authentication service is still in an early stage and lacks fine-grained access control. Additionally, message de-duplication and rate-limiting are not yet implemented, which may pose issues in high-traffic environments.

Another limitation we identified was the lack of detail regarding the Piranha CMS Core. In this iteration, the component is represented in an abstract and high-level manner, which made it difficult to determine which specific internal modules or services our implementation would need to interact with.

3.5 Iteration 4

In our final iteration, we took the feedback provided from our presentation and focused on the most referred topic: the unclear integration with the Piranha CMS core and how external systems can see content from our CMS.

The goal in this iteration was to address how our system would interact with the existing Piranha CMS Core given its poorly described architecture. In the end, we achieved this architecture:

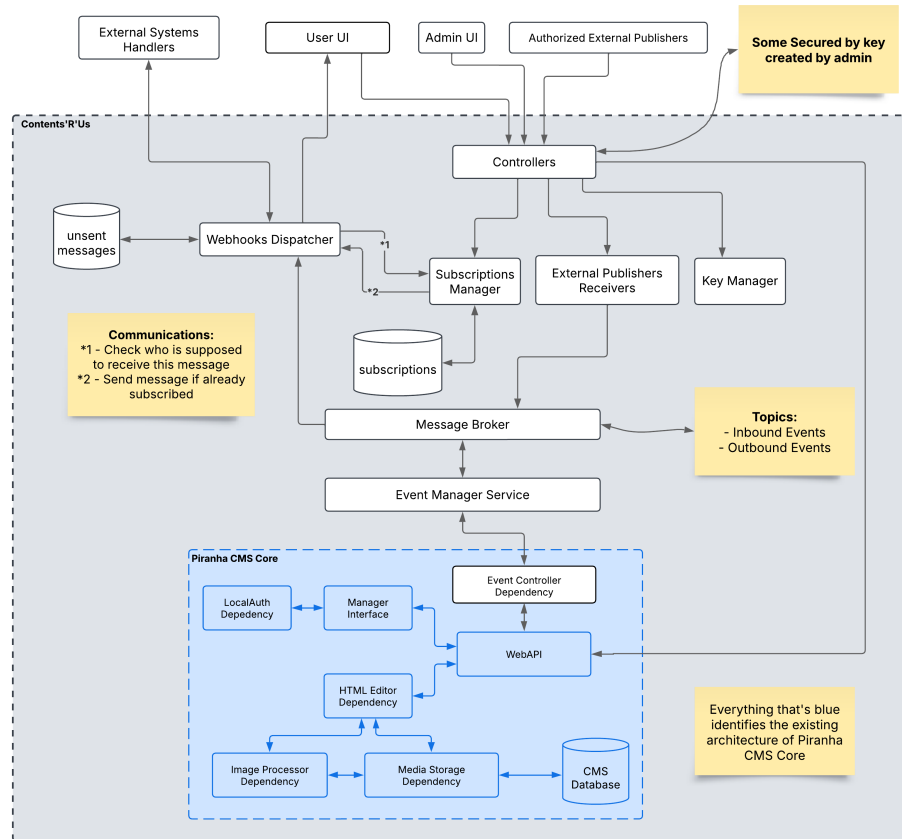


Figure 4: Iteration 4 Architecture

This last iteration added more components to our architecture:

- **Piranha CMS Core components:**

- WebApi;
- HTML Editor Dependency;
- Image Processor Dependency;
- Media Storage Dependency;
- Manager Interface;
- LocalAuth Dependency;
- CMS Database.

- **Event Controller Dependency:** Captures and forwards internal CMS events to the Event Manager Service, acting as the bridge between traditional CMS operations and the event-driven infrastructure.

This iteration allowed us to rethink how we integrate the newer architectural components into the existing core application. With this new architecture we explicitly mapped the dependencies to existing Piranha CMS Core (blue components), (LocalAuth, WebAPI, HTML Editor, Media Storage, etc). We also added a new dependency to extend the core CMS and effectively enabling allows the communication between our new components and the ones present the Piranha CMS core.

This final iteration significantly enhances the architecture by explicitly defining the boundaries and dependencies of the Piranha CMS Core, resolving previous ambiguities. By introducing the Event Controller Dependency, the system now supports seamless event propagation from internal CMS actions to external services, enabling robust publish/subscribe capabilities.

4 Core Domains

To implement a modern, event-driven architecture in Piranha CMS, we applied **Domain-Driven Design (DDD)** to identify, model, and modularize the key areas of the system. Our goal is to isolate core domains that encapsulate distinct responsibilities, maintain high cohesion, and remain loosely coupled to facilitate extensibility, integration, and testability.

4.1 Core Domains

We identified the following **four core domains**, each reflecting a critical capability required to achieve the project vision:

- **Event Publishing Domain**
- **Subscription Management Domain**
- **External Integration and Inbound Event Handling Domain**
- **Security and Access Control Domain**

These domains represent distinct bounded contexts. They communicate via well-defined interfaces and event-driven mechanisms, ensuring that changes in one domain have minimal impact on others.

4.1.1 Event Publishing Domain

This domain handles the production, structuring, and broadcasting of outbound domain events related to CMS actions. It ensures decoupled, asynchronous communication with external systems.

Key Responsibilities

- Define **event contracts** for core CMS actions (e.g., `ContentCreated`, `ContentUpdated`, `ContentDeleted`);
- Manage a central **Message Broker** to decouple producers (CMS Core/Event Manager) from consumers (external systems);
- Coordinate with the **Webhook Dispatcher** to deliver events only to valid subscribers;
- Tag each event with metadata (e.g., content type, timestamp, origin) for filtering/routing;
- Support topic-based filtering and load balancing for scalability.

Key Entities/Services

- **Event**: Domain object representing a business-relevant action;
- **EventTopic**: Logical grouping of similar events;
- **MessageBrokerService**: Handles asynchronous event queuing and delivery;
- **WebHookDispatcher**: Retrieves pending messages and invokes external webhooks;

Technical Considerations:

- Use a reliable queueing mechanism like **Kafka**.
- Guarantee **at-least-once** or **exactly-once** delivery semantics

4.1.2 Subscription Management Domain

This domain is responsible for handling how external consumers subscribe to CMS events. It encapsulates the logic and policies for managing subscriptions and their lifecycle.

Key Responsibilities

- Manage and store subscription data, topics, endpoints, filters, and permissions;
- Provide an admin interface to register/update/remove subscriptions.
- Validate new subscriptions (e.g., endpoint health check, authentication);
- Coordinate with the WebHook Dispatcher to ensure correct message routing.

Key Entities/Services

- **Subscription**: Domain entity representing a single subscription instace (e.g., topic, endpoint, secret);
- **SubscriptionManagerService**: Core service that verifies and maintains the subscription catalog;
- **API Gateway**: Exposes public/admin-facing endpoints for subscription management;
- **subscriptions data store**: Persistent storage for registered subscription.

Technical Considerations

- Rate-limit and throttle endpoints to prevent overload;
- Support retry policies and failure handling for undelivered messages;
- Encrypt sensitive data like webhook secrets or callback URLs.

4.1.3 External Integration and Inbound Event Handling Domain

Handles secure authenticated intake of domain events initiated by external systems (like blogs, social media, etc). It translates these events into internal CMS actions.

Key Responsibilities

- Accept inbound event through a secure API interface;
- Validate incoming messages using API keys and digital signatures;
- Forward validated events to the Message Broker for processing by the **Event Manager Service**
- Support a wide range of actions triggered externally, like, content publication, status updates, etc.

Key Entities/Services

- **InboundEvent**: Represents externally sourced events;
- **Publisher**: Represents an authorized external entity allowed to send incoming events;
- **PuublishReveicer**: Component that receives, authenticates, and validates incoming events;
- **EventProcessor**: Applies CMS-side logic based on event type (example: update content in DB).

Technical Considerations

- Enforce validation, signature verification, and IP whitelisting;
- Log and audit all incoming events for traceability and debugging;
- Allow external publishers to manage configurations through the admin interface.

4.1.4 Security and Access Control Domain

A cross-domain, foundational layer that ensures only authenticated and authorized actors interact with the system, protecting both event publishing and subscription mechanisms.

Key Responsibilities

- Generate and manage API keys for all publishers and subscribers;
- Authenticate incoming and outgoing requests via keys or tokens;
- Define role-based access control (RBAC) for internal users managing subscriptions;
- Encrypt/decrypt secure fields and enforce HTTPS-only communication.

Key Entities/Services

- **Key**: Represents credentials issued to a publisher or subscriber;
- **KeyManager**: Service responsible for creating, validating, rotating, and revoking keys;
- **AuthService**: Provides token validation, signature checks, and scopes;
- **Permission**: Defines what each key can do, like, publish only, subscribe only.

Technical Considerations

- Implement JWT-base verification;
- Provide an admin dashboard for managing secrets;
- Include audit trails and anomaly detection for suspicious usage.

4.2 Supporting Subdomains

For **Supporting Subdomains** our group only found and defined one, **Monitoring & Observability**

4.2.1 Monitoring & Observability Subdomain

This domain is responsible for ensuring operational transparency across the system. It provides real-time and historical insights into the delivery pipeline of events and helps diagnose system issues, ensures Service Level Agreements (SLAs), and support auditing for compliance.

Key Responsibilities

- Log **delivery attempts**, including timestamps, endpoint targets, and payload metadata;
- Track **delivery status**: successful deliveries, retries, and failures;
- Store and expose **metrics** like latency, throughput, error rates, retry counts, and subscriber-specific health;
- Maintain **audit trails** of event flows for compliance and debugging (for example, GDPR, or HIPAA);
- Alert anomalies like repeated failures or unexpected latency spikes.

Key Entities/Services

- **DeliveryLog**: Records each outbound delivery attempt and its outcome;
- **RetryMonitor**: Tracks and schedules retries for failed deliveries;
- **AuditTrail**: Captures immutable records of key actions (e.g., key creation, publisher registration);

- **MonitoringService**: Collects and aggregates system metrics;
- **AlertingSystem**: Generates alerts based on metric thresholds or error patterns.

Technical Considerations

- Integrate with tools like Prometheus, Grafana, Jaeger, or ELK stack;
- Ensure logs are structured and filterable by event type, endpoint, and status;
- Implement retention policies and anonymization if sensitive data is logged;

4.3 Generic Subdomains

In terms of **Generic Subdomains** only identified 2 main ones:

- **User Management**
- **CMS Platform**

4.3.1 User Management Subdomain

Leverages the existing identity and access control infrastructure provided by Piranha CMS. It manages users who interact with admin-facing features like key management, subscriptions, and integration configuration.

Key Responsibilities

- Authenticate and authorize admin users;
- Manage user roles and permissions, like, system admin, integration manager, etc;
- Enable secure login and session management;
- Support multi-user operations, audit user actions, and enforce RBAC policies.

Key Entities/Services

- **User**: Represents a system user with login credentials and metadata;
- **Role**: Encapsulates a permission group, that can manage keys, or view logs, etc;
- **AuthenticationService**: Handles sign-in and session validation;
- **AuthorizationService**: Enforces what actions a user is permitted to take.

Technical Considerations

- Extend Piranha's built-in identity provider where necessary;
- Enforce password and session security best practices;
- Log user actions related to integration and security settings.

4.3.2 CMS Platform

This subdomain consists of the foundational Piranha CMS capabilities. It manages content models, templates, media, and editorial workflows, acting as the *content engine* that event-driven extensions are layered on top of.

Key Responsibilities

- Handle core content creation, editing, versioning, and publishing workflows;
- Serve API responses and web views using Piranha’s headless and hybrid CMS features;
- Act as data source for outbound events, like when a new page is published;
- Enable extension through modules, middleware, and API endpoints.

Key Entities/Services

- **Page, Post, Block, Region:** Piranha’s core content entities;
- **ContentService:** Handles business logic around content lifecycle;
- **CMS API:** Serves data to front-end clients or headless consumers;
- **Piranha Middleware:** Allows hooks and pipelines for integrating event triggers.

Technical Considerations

- Ensure decoupled triggers so event publishing doesn’t block content updates;
- Use versioned APIs for external systems consuming content;
- Apply performance tuning and caching for scalable API response times.

5 Cross-Cutting Concerns

In the context of evolving Piranha CMS to support event-driven communication and secure external integrations, managing cross-cutting concerns is critical to ensure the system is not only functional but also secure, maintainable, and observable. These concerns span multiple architectural layers and influence both runtime behavior and system resilience.

Below are the primary cross-cutting concerns relevant to Scenario 3, along with our strategies for addressing them:

5.1 Security

Security is the most critical concern in this scenario. The system will expose endpoints for receiving and sending sensitive content-related events, potentially across organizational boundaries.

Strategies

- **Authentication and Authorization:** Implement token-based (e.g., JWT) authentication for external publishers and internal consumers. Role-based access control (RBAC) will restrict access to critical operations like subscription management or content updates;
- **Message Integrity and signing:** All messages exchanged between systems will be signed using HMAC or asymmetric encryption (e.g., RSA) to ensure tamper resistance. Verification is enforced at both ends before processing;
- **Key Management:** Admin users can configure and rotate keys through a secure UI. Key storage will be encrypted and auditable;
- **HTTPS and TLS Enforcement:** All communication channels must be encrypted via HTTPS. Self-signed certificates will be disallowed for production environments;
- **Replay Attack Protection:** Timestamps will be included in message payloads to avoid replay attacks. Messages with expired timestamps will be rejected.

In the previous section, we laid out three key features we aimed to build over the course of four weeks: the Event Publishing Infrastructure, the Webhook Delivery System, and the Subscription Management Module. By the end of the development period, we not only completed all of these as planned but also went a step further by adding a new module for key generation and management. This extra feature strengthens how clients are identified and how access is controlled across the system.

Looking back at the roadmap, we're proud to say that we hit all our goals on time—and even managed to deliver a bit more than we originally set out to do.

5.2 Logging and Auditing

Observability is key in an event-driven system, particularly when dealing with asynchronous workflows and third-party integrations. Logs are essential for debugging, auditing, and security analysis.

Strategies

- **Structured Logging:** Use structured logging frameworks (e.g. Serilog) to emit logs in a consistent JSON format for easy parsing by log aggregators (e.g., ELK stack, Seq);
- **Correlation IDs:** Every event or API call will carry a correlation ID to trace its life-cycle across services, improving diagnostics and incident response;
- **Audit Trails:** Sensitive actions (e.g., publishing events, accepting inbound messages, key rotation) will be logged in an immutable audit log;
- **Log Levels and filtering:** Logging granularity will be configurable. Sensitive data will be masked from logs, and error logs will include contextual information.

5.3 Error Handling and Resilience

Event-driven systems must handle failures gracefully to avoid cascading issues and data loss, especially with asynchronous message delivery.

Strategies

- **Retry Policies:** Failed outbound or inbound event processing will trigger exponential back-off retries, with configurable limits.
- **Unsent Messages DB:** After retry limits are reached, failed messages will be moved to a database.
- **Fallback Handlers:** Define fallback logic for common failure scenarios (e.g., timeouts from external services).

5.4 Observability and Monitoring

The system must offer real-time insights into event flow health, performance bottlenecks, and integration statuses.

Strategies

- **Health Checks and Dashboards:** Expose health endpoints for integration with monitoring tools (e.g., Prometheus, Grafana). These include event queue sizes, and key rotation status.
- **Real-Time Metrics Collection:** Track metrics such as event throughput, processing latency, and success/failure rates.
- **Alerts and Notifications:** Threshold-based alerts (e.g., high growth in the unsent messages DB, repeated signature verification failures) will notify administrators for quick remediation.

5.5 Performance and Scalability

Asynchronous systems that rely on messaging can experience bursts of traffic or load spikes. Poorly designed systems may suffer from message bottlenecks, thread exhaustion, or delayed content updates.

Strategies

- **Asynchronous Message Queues:** Implement queue-based architecture (e.g., RabbitMQ, Kafka) to decouple producers and consumers, ensuring smooth load balancing.
- **Back-pressure Mechanisms:** Introduce rate limiting and flow control to prevent overload on internal or external systems.

5.6 Compliance and Data Protection

When transmitting or processing content-related events, regulatory compliance with data protection laws (e.g., GDPR, HIPAA) becomes crucial.

Strategies

- **Data Minimization:** Ensure that only necessary data is included in messages. Personally Identifiable Information (PII) should be redacted unless explicitly required.
- **Consent and Access Logging:** Record user consent and track how and when data is shared with third parties.
- **Retention Policies:** Define clear retention rules for logs, messages, and audit data to comply with regulations and reduce storage costs.

6 Proposed Architecture and Roadmap

The system architecture is based on Piranha CMS, which serves as the foundation for managing core content and operations. The architecture then incorporates multiple components to implement an asynchronous publish-subscribe (pub-sub) system. To make this feasible, a webhook pattern is used by the clients, along with a message broker that stores and forwards events between Piranha and the system responsible for managing subscriptions.

To arrive at the final architecture, we followed an iterative process based on the Attribute-Driven Design (ADD) method. We began by designing the components necessary for sending events to external clients, focusing on the webhook delivery mechanism. In the next iteration, we extended the architecture to support the internal publication of events from the CMS, ensuring reliable and scalable event generation and distribution. Finally, we addressed the security requirements, particularly the management of API keys and access control, with emphasis on the role of the administrator who oversees key generation and subscription permissions.

Ultimately, we have developed an architecture consisting of eight components that work together to enable the event-driven CMS. In this section, we will explain each component in detail, along with their respective interfaces.

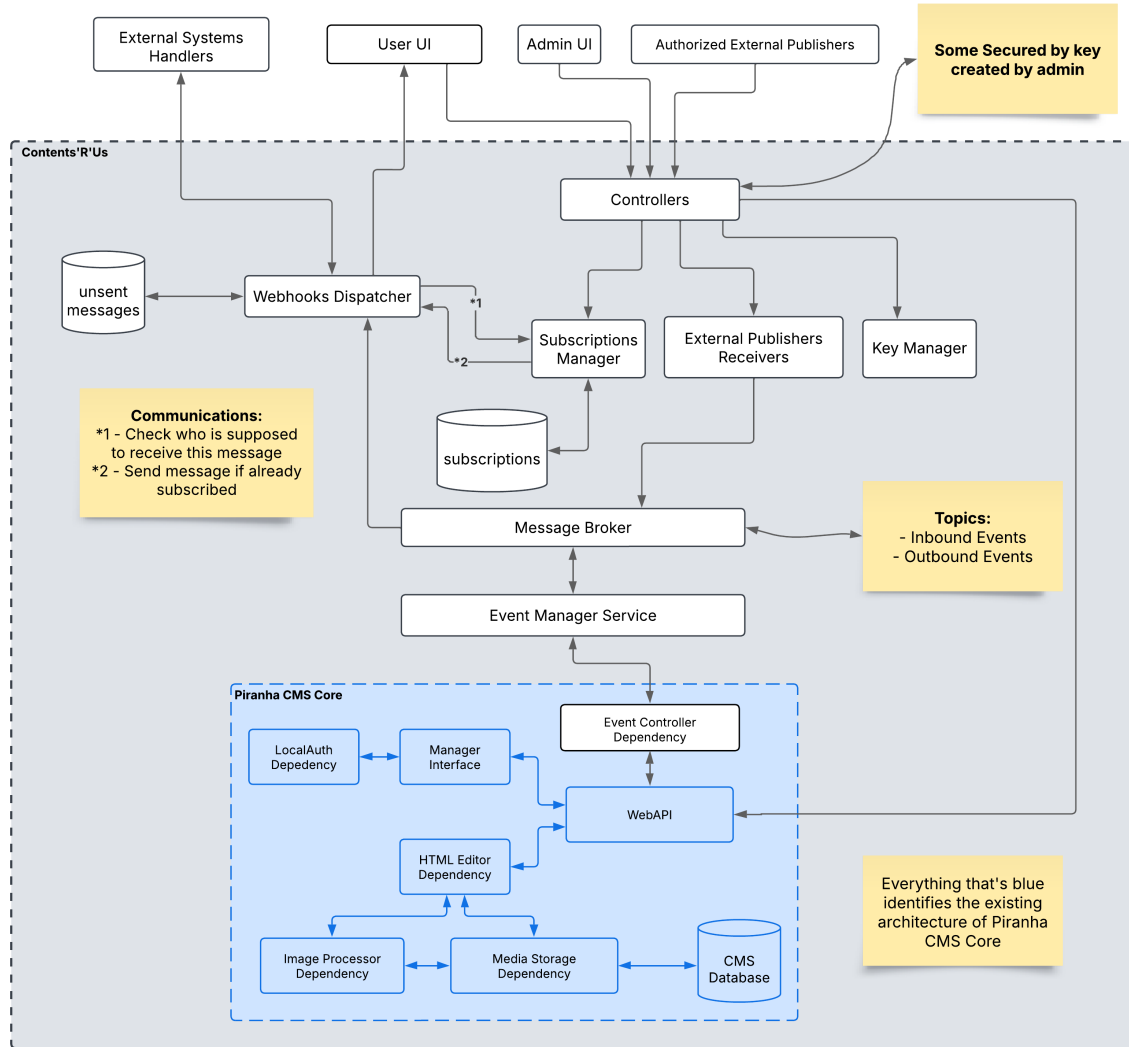


Figure 5: Proposed architecture

6.1 The Components and their responsibilities

The architecture comprises several essential components necessary for the system to operate effectively, as illustrated in the image above (Figure 16). These components are outlined as follows:

- Controllers
- Webhooks Dispatcher
- Subscriptions Manager
- Publish Receiver
- Message Broker
- Event Manager Service
- Piranha CMS Core

- Key Manager

6.1.1 Controllers

The controllers are the entry points where requests arrive in the system. They perform security verification based on the type of request and forward the information to the appropriate service for processing.

6.1.2 Webhooks Dispatcher

The Webhooks Dispatcher is the component that ensures the communication of events to the clients subscribed or that need to receive information from the CMS. It ensures that events generated within Piranha CMS are communicated to external systems or subscribers. It listens for events created in the Event Manager Service that are stored in the message broker. The dispatcher works by receiving event notifications, processing them, and sending them to the clients based on the subscriptions stored.

6.1.3 Subscriptions Manager

This component is responsible for managing relations between events and clients. It saves and queries over a list of active subscriptions, mapping events to multiple clients that have the interest of receiving notifications for the events they subscribed to. It ensures that only relevant information is given to the clients. It communicates closely with the Webhook Dispatcher and also handles the logic of creating new subscriptions and sending events that trigger the delivery of unsent events when an already registered client registers again. It also sends events to the Webhook Dispatcher to delete unread messages from a client that does not register.

6.1.4 Publish Receiver

The Publish Receiver is the component that receives event messages from external sources or other components in the system. It is ready to accept incoming events that need to be processed and published by the system. The Publish Receiver publishes events to the message broker, and its primary responsibility is to ensure that events are delivered to the appropriate component for further processing. It plays a key role in integrating external event sources such as updates and posts into the Piranha CMS ecosystem and ensures that the system remains reactive to external triggers.

6.1.5 Message Broker

The Message Broker is the central component that makes the infrastructure more scalable by allowing multiple instances of a single component. It facilitates communication between various services in an event-driven manner. It serves as a queue of messages and events for multiple services to consume. The Message Broker decouples components, allowing them to communicate asynchronously. This supports message persistence and enables scalability by distributing messages across multiple consumers.

6.1.6 Event Manager Service

The Event Manager Service is the core component responsible for managing the flow of events throughout the system. It monitors changes within the Piranha CMS Core (such as content creation, updates, or deletions) and triggers corresponding events based on these changes. Once an event is generated, it is placed in the message broker, ensuring that all other services are notified of the event. The Event Manager Service acts as the controller of event generation, ensuring that events are triggered at the right time and in the correct format, depending on the information received by the CMS.

6.1.7 Piranha CMS Core

At the heart of the event-driven architecture lies the Piranha CMS Core. This is the central system that powers the content management functionalities, such as creating, updating, and deleting content. When a change is made in the CMS, a trigger can be initiated through the Event Manager Service. The events are then propagated throughout the system, reaching components like the Webhook Dispatcher. The Piranha CMS Core is the origin of the events and is central to the operation of the event-driven system.

The CMS Core is composed of several key modules:

- **Manager Interface:** The administrative interface used for managing content and configuration. It serves as the main interaction point for users and connects to subsystems like the WebAPI, HTML Editor and manage keys.
- **WebAPI:** Provides programmatic access to CMS functionalities via HTTP. It connects with the Event Controller to notify the event system of content changes and receive content from Event Controller.
- **Event Controller Dependency:** Captures and forwards internal CMS events to the Event Manager Service, acting as the bridge between traditional CMS operations and the event-driven infrastructure.
- **LocalAuth Dependency:** Manages authentication and authorization within the CMS, ensuring secure access to admin users.
- **HTML Editor Dependency:** Enables rich text content editing within the Manager Interface and supports integration with image and media resources.
- **Image Processor Dependency:** Responsible for processing images and integrates with the media storage system.
- **Media Storage Dependency:** Handles the storage and retrieval of uploaded media files and interacts directly with the CMS database.
- **CMS Database:** Stores all persistent data including content, media assets, and user credentials. It serves as the backbone of the CMS's data layer.

6.1.8 Key Manager

The Key Manager enables administrators to securely manage API keys used to authenticate requests to the Publish Receiver. It ensures that only authorized systems or users can publish events into the CMS infrastructure, enhancing security and access control across the event-driven ecosystem.

6.2 Communication/Interfaces Between Components

The communication between these components is organized as a flow that starts with content changes within the Piranha CMS Core and ends with the dissemination of events to interested parties. Here's how the components interact:

1. The **Piranha CMS Core** detects a change in the system (e.g., new content created) and triggers an event that is sent to the Event Manager Service.
2. The **Event Manager Service** listens for these changes and publishes them to the **Message Broker** where they will be consumed.
3. The **Message Broker** distributes the events to the entire system. It is part of two main flow one where the information goes to the PiranhaCMS and other where the information goes to the Webhooks Dispatcher.
4. The **Subscriptions Manager** filters the events, ensuring that only the subscribers receive the event this information is consumed by the Webhook Dispatcher. The Controllers allows clients and admins to interact with this module changing subscriptions.
5. The **Webhooks Dispatcher** consumes events from the message broker and sends the event to external systems based on the subscriptions present in the Subscription Manager.
6. The **Publish Receiver** can listen to incoming request to publish information into the system, this component sends events to the message broker.
7. The **Key Manager** can manage the keys generated by the admin and secure login with key verification in requests.

This architecture ensures that Piranha CMS is responsive and scalable, publish and subscribe content manager system. Each component plays a vital role in maintaining the overall structure and functionality of the system.

6.3 Road Map

The development of the event-driven extension for Piranha CMS will follow a phased approach to ensure proper integration, testing, and functionality. The following are the three principal features to be built:

1. **Event Publishing Infrastructure**
Development of the internal event flow, including the Event Manager Service and integration with a Message Broker. This step ensures that events from the CMS Core are properly captured, and that events originating from authorized publishing clients can reach the CMS.
2. **Webhook Delivery System**
Construction of the Webhooks Dispatcher, responsible for delivering event payloads to external subscribers. It will check for subscriptions created in the Subscription Management Module.

3. **Subscription Management Module**

Implementation of a robust Subscriptions Manager capable of registering, updating, and deleting client subscriptions. This module will also handle logic related to event filtering and mapping subscribers to specific events. This is foundational to enabling selective notification in the publish-subscribe model.

Each feature will be developed iteratively and validated against real CMS interactions to ensure full compatibility and performance under expected system loads.

7 Implementation and Demonstration

In this section, we will present our implementation following the **three principal features** that we defined in section 6.3.

All the code developed and described in this section of this report is available in the following link: https://github.com/zegameiro/AS_Final_Assignment

7.1 The Road-map

In the previous section, we laid out three key features we aimed to build over the course of four weeks: the *Event Publishing Infrastructure*, the *Webhook Delivery System*, and the *Subscription Management Module*. By the end of the development period, we not only completed all of these as planned but also went a step further by adding a new module for *key generation and management*.

This additional feature emerged during a later iteration of the ADD cycle. As a team, we collectively decided that, in order to properly support the other features—while also addressing our cross-cutting concerns and meeting the quality attributes outlined in Section 5.1—the inclusion of a key management mechanism was necessary. It became a natural extension of the work, seamlessly fitting into our four-week development plan.

Looking back at the road-map, we're proud to say that we achieved all of our goals within the planned timeline — and even delivered more than we originally set out to do.

7.2 Implementation

7.2.1 Tags

The first change we made to the system was adding a way to tag content—specifically Pages and Files—within Piranha CMS. To implement this, we added a Tags field to the models for both Page and Media. We then created the necessary migrations so the database could support this new field. Finally, we updated the services and UI to ensure the tags were properly handled—passing the values from the UI through the request, into the model, and ultimately saving them in the database.

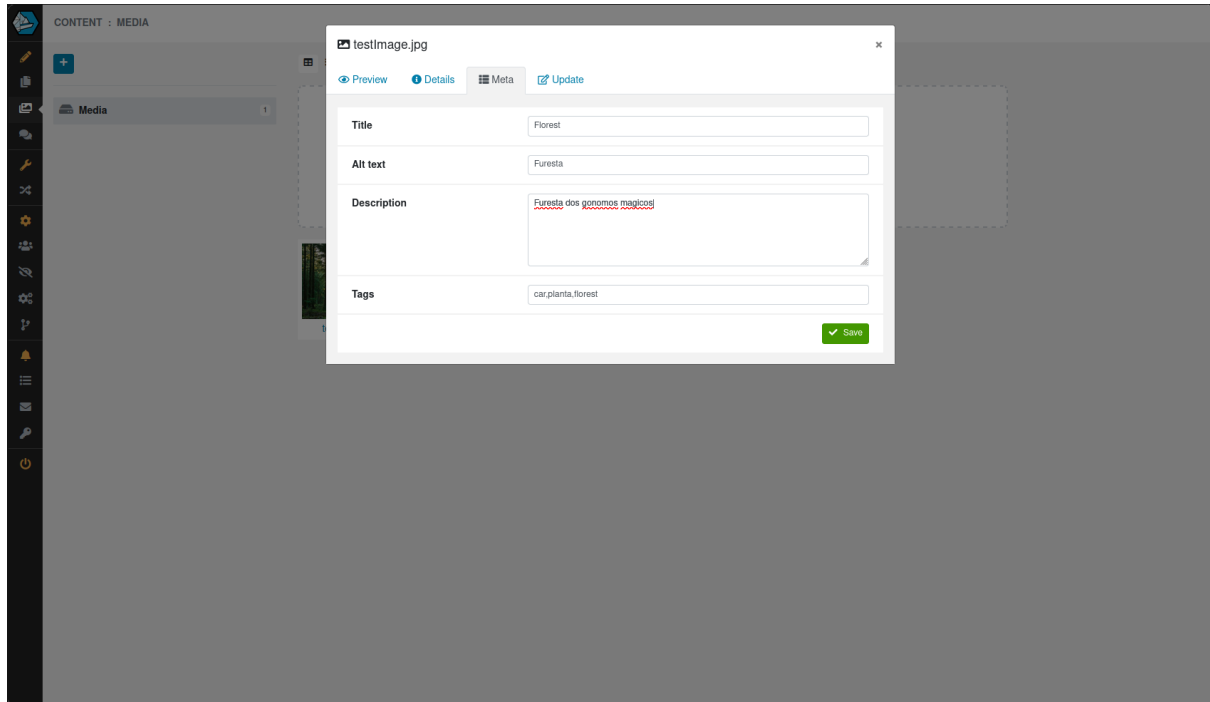


Figure 6: Tags added in the UI of metadata

7.2.2 Subscription Management Module

The main idea of this module is to allow the system to manage subscriptions by creating, updating or deleting them. To achieve this, we first needed to add a new table to the Piranha's database, called Subscriptions. We defined the table with the following properties

- **Id**: unique identifier for the subscription.
- **EventStatus**: Describes the current status of the event. For this an enumerate was created with the following items:
 - **Create**, identifies an event where a new resource is created;
 - **Update**, identifies an event where an existing resource is updated;
 - **Delete**, identifies an event where an existing resource is deleted;
 - **UpdateMetaData**, identifies an event where the metadata of an existing resource is updated.
- **EventType**: Specifies the type of the event the subscription is associated to. We also created an enumerate with the following items:
 - **Page**, represents events related to pages;
 - **Media**, represents events related to media items (such as images, pdf files, etc).
- **Tags**: A comma-separated list of tags or keywords related to the subscription.

- **CallbackUrl**: The URL that should be called when an event with the same type, status and tag has the subscription occurs.
- **Created**: The timestamp of when the subscription was created, automatically set to the current UTC time.

With this entity defined, the following step was to apply this change to the Piranha's database. For this it was required to create Migrations and apply them to the database. Our group had some problems with this step because to create new migrations it was required to have **DOTNET 8.0**, and to run the application it was **DOTNET 9.0**.

We add the to add the following lines of code to the file `Db.cs` to configure the `Subscription` entity using the Entity Framework Core Fluent API. This setup ensures that the entity is properly mapped to the database with the correct constraints and indexing.

```
mb.Entity<Models.Subscription>().ToTable("Piranha_Subscriptions");

mb.Entity<Models.Subscription>().Property(s => s.Id).IsRequired();

mb.Entity<Models.Subscription>().Property(s =>
    ↪ s.EventStatus).IsRequired();

mb.Entity<Models.Subscription>().Property(s =>
    ↪ s.EventType).IsRequired();

mb.Entity<Models.Subscription>().Property(s =>
    ↪ s.Tags).IsRequired().HasMaxLength(256);

mb.Entity<Models.Subscription>().Property(s =>
    ↪ s.CallbackUrl).IsRequired().HasMaxLength(512);

mb.Entity<Models.Subscription>().Property(s =>
    ↪ s.Created).IsRequired();

mb.Entity<Models.Subscription>().HasIndex(s => new
    ↪ {s.EventType,s.EventStatus,s.Tags}).IsUnique();
```

Some properties have a limit of characters specified, and all of them have the `.IsRequired()` to ensure that they must always be provided. An unique composite index across the `EventType`, `EventStatus` and `Tags` fields, this means that no two records can share the same combination of these three values.

Besides this, the following line was added to the `IDb.cs`:

```
DbSet<Models.Subscription> Subscriptions { get; set; }
```

In the Entity Framework Core, the `DbSet<T>` properties represent collections of entities that can be queried from or written into the database, meaning that this allows the application to query, create, update or delete the **Subscription Table**.

We decided to only create migrations in the `Piranha.Data.EF.SQLite` because a `.db` file is used in this project, and in the others existing (`Piranha.Data.EF.MySql`, `Piranha.Data.EF.PostGreSql`, `Piranha.Data.EF.SqlServer`) it required a more complex setup.

With this done, the next step was to create a **repository** to communicate with the database. For this repository two files were created `ISubscriptionRepository.cs` and `SubscriptionRepository.cs`, where the first file defines an interface and the second one the implementation of the interface. The methods created in these files were the following:

- `public SubscriptionRepository(IDb db)`: Initializes the repository with a database context. The constructor takes an `IDb` instance (which represents the EF database context) and assigns it to a private field `_db` for use in other methods;
- `public Task DeleteAllAsync()`: Deletes all subscription records from the database. It uses `ExecuteDeleteAsync()` for efficient bulk deletion directly in the database without loading entities into memory;
- `public Task DeleteAsync(Guid id)`: Deletes a specific subscription by its `Id`;
- `public async Task<IEnumerable<Subscription>> GetAllAsync()`: Retrieves all subscriptions, using `AsNoTracking()` for better performance when no entity tracking is required;
- `public async Task<IEnumerable<Subscription>> GetByEventTypeAsync(string eventType)`: Retrieves subscriptions filtered by `EventType`;
- `public async Task<IEnumerable<Subscription>> GetByEventStatusAsync(string eventStatus)`: Retrieves all subscriptions filtered by `EventStatus`;
- `public async Task<IEnumerable<Subscription>> GetByTagsAsync(string filter)`: Retrieves subscriptions filtered by tags;
- `public async Task<Subscription> GetByIdAsync(Guid id)`: Retrieves a single subscription by its unique `Id`;
- `public async Task<Subscription> SaveAsync(Subscription subscription)`: Creates a new subscription or updates an existing one. If the `Id` is empty, it treats the record as new, assigns a new ID and current UTC timestamp, and adds it to the context. If the `Id` exists, it finds the record, updates the fields, and throws an error if the record is not found.

With the repository, created the next step was to implement a service to communicate with the created repository. For this, two other files were created in the Piranha project, called `ISubscriptionService` and `SubscriptionService`. These also include functions similar to those inside the repository:

- `public SubscriptionService(ISubscriptionRepository repo)`: Initializes the service with an instance of `ISubscriptionRepository`. Follows **dependency injection principles** to keep the service loosely coupled and testable.

- **Delegated Methods** `GetAllAsync()`, `GetByIdAsync()`, `GetByEventTypeAsync()`, `GetByEventStatusAsync()`, `GetByTagsAsync()` and `DeleteAllAsync()`: These methods are thin wrappers that **delegate directly to the repository**. No additional logic is applied, maintains separation of concerns.
- `public async Task SaveAsync(Subscription subscription)`: Contains business logic to prevent duplicate subscriptions. Checks if a subscription with the same `EventType`, `EventStatus`, `Tags`, and a partially matching `CallbackUrl` already exists. If such subscription is found, it throws an `InvalidOperationException` to enforce uniqueness beyond what the database constraints allows.
- `public async Task DeleteAsync(Guid id)`: Performs a **safety check** to ensure the subscription exists before attempting to delete. Prevents silent failures by throwing an exception if the `id` does not match any existing subscription.

Now only two steps are missing to complete the Subscription manager implementation, which are the implementation of a controller that allows external users to create, retrieve and delete a subscription and a web page in the admin interface to allow an administrator to also manage subscriptions.

In terms of a Controller, we created a simple file called `SubscriptionController.cs` with the idea of having an API controller available to external systems so that they can create new valid subscriptions, delete and see them. For this three main endpoints were created:

- `GET /manager/api/subscription/{id}`: Retrieves a subscription with the specified ID. If none is found a code 404 is sent, else the subscription is sent with a 200 code;
- `POST /manager/api/subscription/`: Saves a new subscription that's present in the body of the request. If the subscription that's present in the body has missing or wrong information a **Bad Request** is sent, else it saves the new subscription with success;
- `DELETE /manager/api/subscription/{id}`: Deletes a subscription with the specified ID. If the given ID isn't connected to any subscription then a 404 error is sent, else the subscription with the ID is deleted with success.

In the manager application, a new tab was added to the side menu, that has 3 options **Events List**, **Subscriptions** and **Keys**:

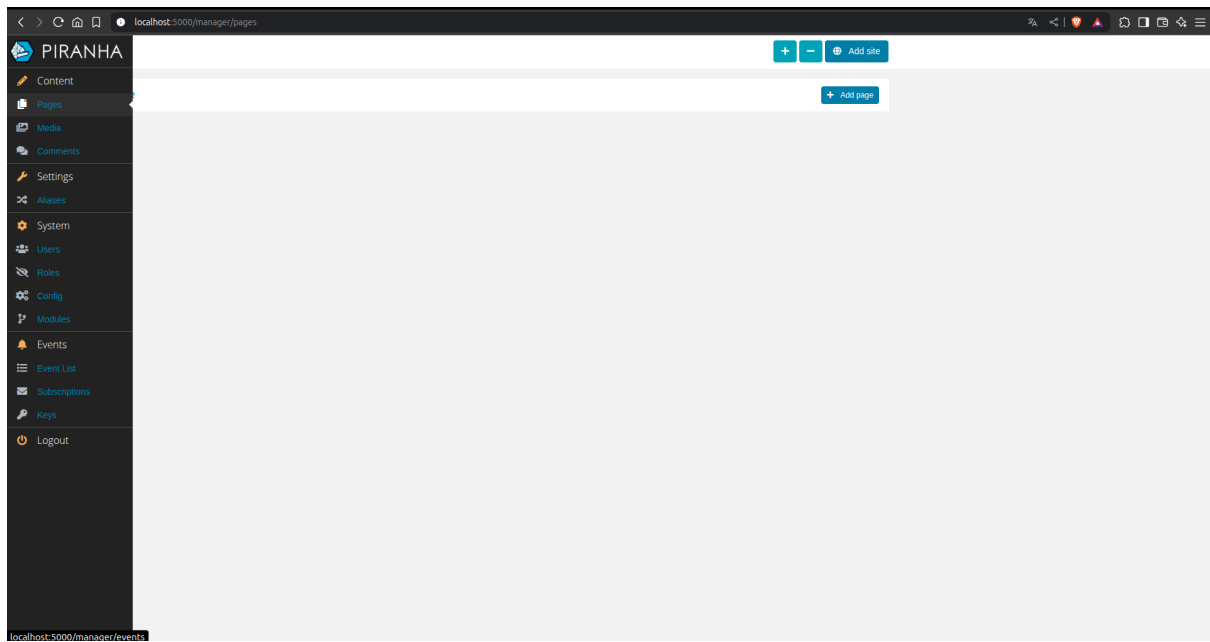


Figure 7: Side Menu with the new options

In the **Subscriptions** tab it is possible to add new subscriptions, delete all the existing ones or specific one, and update an existing subscription, has we can see in the following figures:

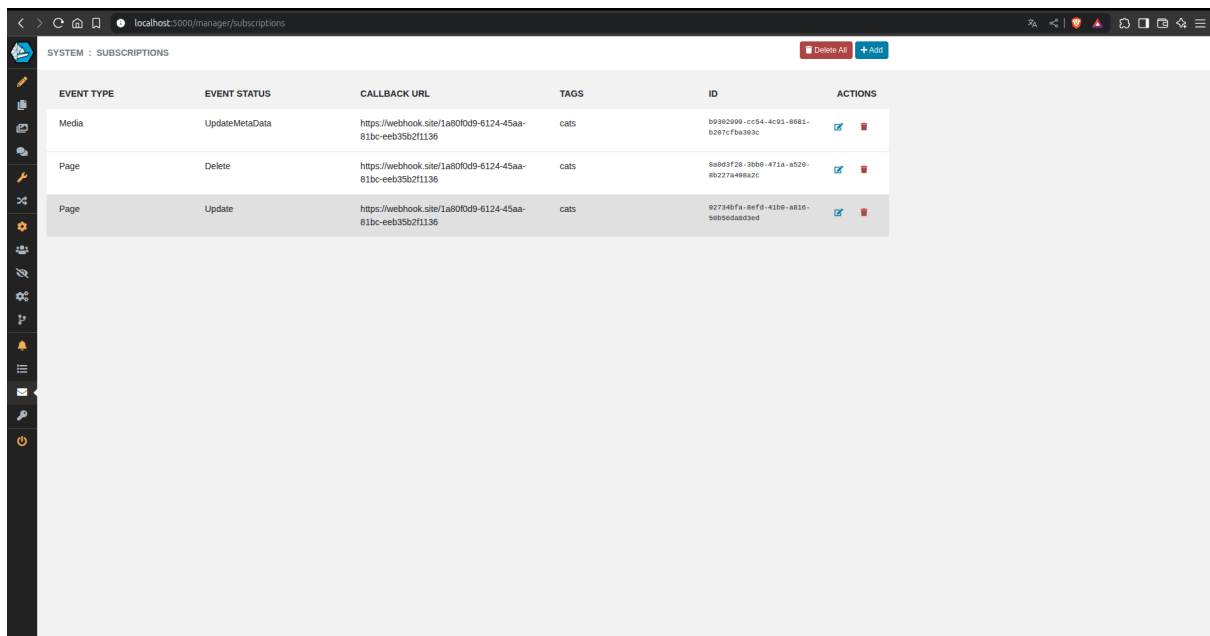


Figure 8: See all the existing subscriptions

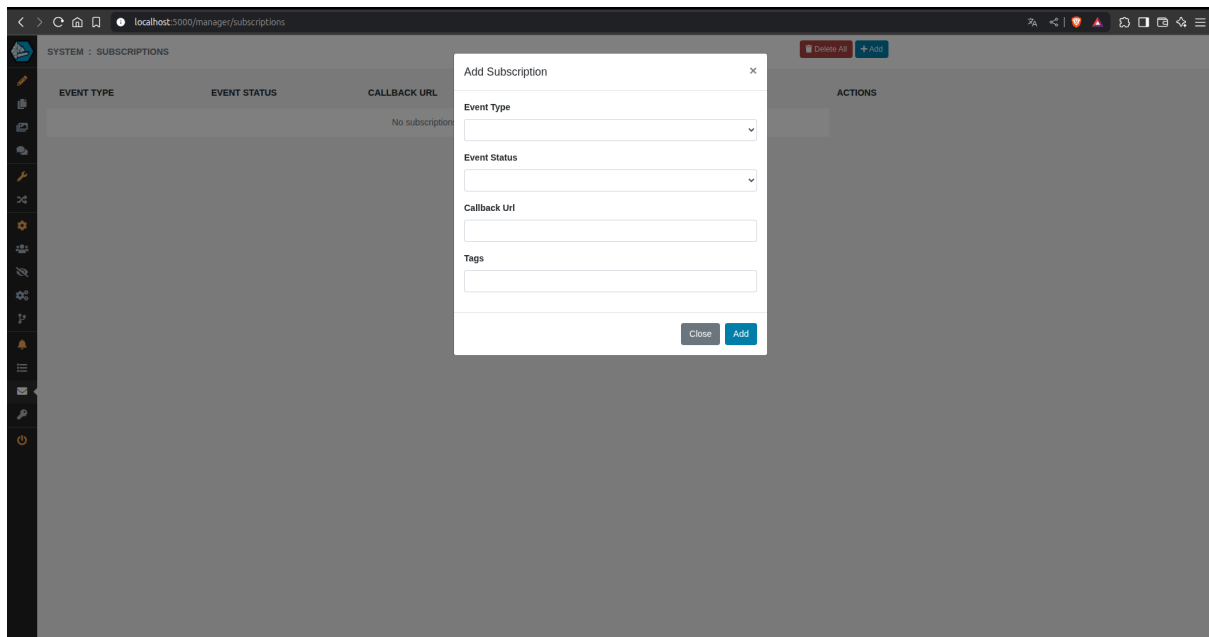


Figure 9: Add a new subscription

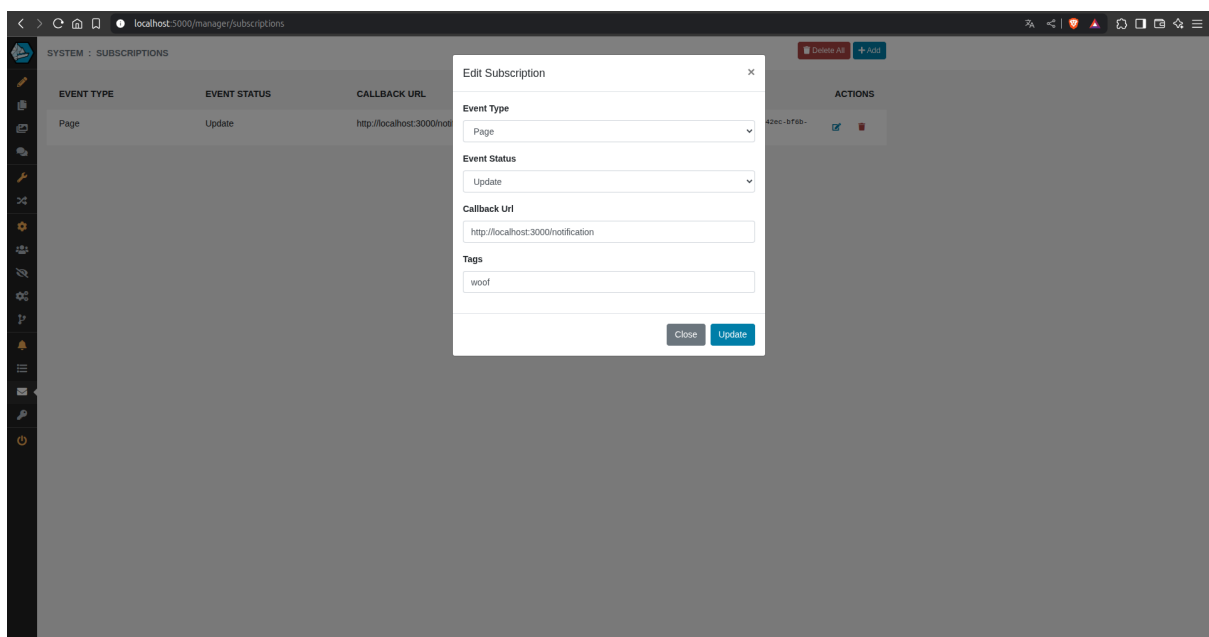


Figure 10: Update an existing subscription

In the form to update or create a subscription, some form validation has been created, where all the fields need to have values, and they cannot be equal to a subscription that already exists with those exact same values.

7.2.3 Event Publishing Infrastructure

To track and respond to key actions in the system — like creating or updating a page or media item — we needed a way to generate and handle events. Since Piranha CMS

doesn't come with a built-in event system like this, we built one ourselves by extending its infrastructure.

We started by creating a new service called the `EventConsumer`, which runs in the background. It's always active and listens for actions happening across the CMS. Whenever something important happens (like a page being saved), an event is triggered.

These events are sent through the `EventBus`, which publishes them to a message queue. For the message queue we decided to go with RabbitMQ since we are familiar with it and there are easily available packages for dotnet. Here's a quick look at what that publishing code looks like:

```
public async Task Publish(Event @event)
{
    var json = JsonSerializer.Serialize(@event);
    var body = Encoding.UTF8.GetBytes(json);

    await _channel.BasicPublishAsync(exchange: string.Empty, routingKey:
        ↪ QueueName, body: body);

    Console.WriteLine($"[EVENT PUBLISHED] Type: {@event.Type}");
}
```

At the same time, our `EventConsumer` service is listening to that queue. As soon as an event shows up, it picks it up and gets to work. First, it checks which subscribers are interested in this type of event:

```
var subscriptions = await
    ↪ _api.Subscriptions.GetByEventTypeAsync(@event.Type.ToString());
```

If there are any subscriptions, we filter them further to only include those that match the specific status of the event (like `Published` or `Draft`):

```
foreach (var sub in subscriptions)
{
    if (sub.EventStatus == @event.Status.ToString())
    {
        subs.Add(sub);
    }
}
```

Then we narrow it down even more by checking tags. This makes sure subscribers only get notified about events that are truly relevant to them:

```
foreach (var sub in subs)
{
    if (sub.Tags.Contains(','))
    {
        if (sub.Tags.Split(',').Any(tag =>
            ↪ @event.Tags.Contains(tag.Trim())))
        {

```

```

        filtered.Add(sub);
    }
}
else if (@event.Tags.Contains(sub.Tags.Trim()))
{
    filtered.Add(sub);
}
}

```

For each subscriber that matches, we grab the relevant content (like the page or media involved in the event), prepare a payload, and send it to their callback URL via a webhook:

```

var payload = new
{
    Event = new
    {
        @event.Id,
        @event.CreatedAt,
        Type = @event.Type.ToString(),
        Status = @event.Status.ToString(),
        @event.ContentId,
        @event.Tags
    },
    Content = content
};

await _api.Notifications.NotifyAsync(subscription.CallbackUrl, payload);

```

Finally, once all notifications are sent, we mark the event as "consumed" by adding it to a list. We do this inside a lock to make sure multiple events aren't written at the same time, which could cause issues:

```

lock (Lock)
{
    ConsumedEvents.Add(@event);
}

```

7.2.4 Key management module

The key management module was something that our group didn't plan to implement it, however we managed to implement all the three goals defined in the **Road Map**, and so we decided to have a simple implementation of a Key Module, where system administrators can create new keys and provide them to external publishers so that they can use this key to publish new content. This administrators can control which systems can be authorized.

To kickstart this module the first thing to do was to create a new Model, to represent an access key. We created a simple model with only tow properties:

- **Guid Id:** The key that will be provided to the external publishers.

- **string Name:** A name that is associate to the key for better key organization.

Like the subscription model, it was also required to register the Key model in Piranha's main database, for this the following lines were added to the `Db.cs`

```
mb.Entity<Models.Key>().ToTable("Piranha_Keys");
mb.Entity<Models.Key>().Property(k => k.Id).IsRequired();
mb.Entity<Models.Key>().Property(k =>
    ↪ k.Name).IsRequired().HasMaxLength(64);
mb.Entity<Models.Key>().HasIndex(k => new { k.Id }).IsUnique();
```

Where a new table was created called **Piranhas_Keys**, a restriction was added to the **Name property** where it can only have a max length of 64 characters and the **Id** of the key needs to be unique. The line `DbSet<Models.Key> Keys get; set;` was added to the `IDb.cs` file to represent a collection of entities that were used for queries or written into the main database, this means that it allows Piranha's application to create, update or delete elements in the **Keys table**.

After this, migrations were created and the SQLite database was updated with this new table, and with this succeeded we decided to move one to the creation of the **Keys Repository**. This repository is very similar to the **Subscription Repository**, where a class interface and an implementation of the interface was created with the following methods:

- **public KeyRepository(IDb db):** Initializes the repository with a database context, just like the **Subscription Repository constructor**;
- **public Task DeleteAllAsync():** Deletes all the existing keys in Piranha's database;
- **public Task DeleteAsync(Guid id):** Deletes a specific subscription that has the ID equal to the one passed has an argument of the function;
- **public async Task<IEnumerable<Key>> GetAllAsync:** Retrieves all the existing keys in the database;
- **public async Task<Key> GetByIdAsync(Guid id):** Retrieves a key that has an ID equal to the function's argument;
- **public async Task<Key> GetByNameAsync(string name):** Retrieves a key with a name equal to the function's argument;
- **public async Task<Key> SaveAsync(Key key):** Inserts a new key into the database.

The next step is to create a service to communicate with the repository and just like in the **Subscription Module** an interface and a class to implement it was created called **IKeyService** and **KeyService**, with the following methods:

- **GetByIdAsync(), GetByNameAsync(), GetAllAsync() and DeleteAllAsync():** These methods are also thin wrappers that delegate directly to the repository. No additional logic was required to integrate;
- **public Task DeleteAsync(Guid id):** Deletes a key with the specified one in the argument. It checks if a string with the function's argument exists, and if it does it deletes it with success, if it does not it throws an **InvalidOperationException**.

- `public async Task<Key> SaveAsync(Key key):` Inserts into the **Piranha_Keys** database a new key if it is valid. It also checks if any key already exists with the same name has the one that is trying to be created.

Just like the **Subscription Service**, it was also required to register the **Keys Service** in the **IApi.cs** and **Api.cs** files, where in the **IApi.cs** the line `IKeyService Keys get;` was added to have a way to access to the Key Service methods and in the **Api.cs** file a new `KeyService` was initialized `Keys = new KeyService(keyRepository);`.

With all of this done, the final step was to create a new page for the administrator to manage the keys. The following pictures show the web pages developed and actions that are possible to do.

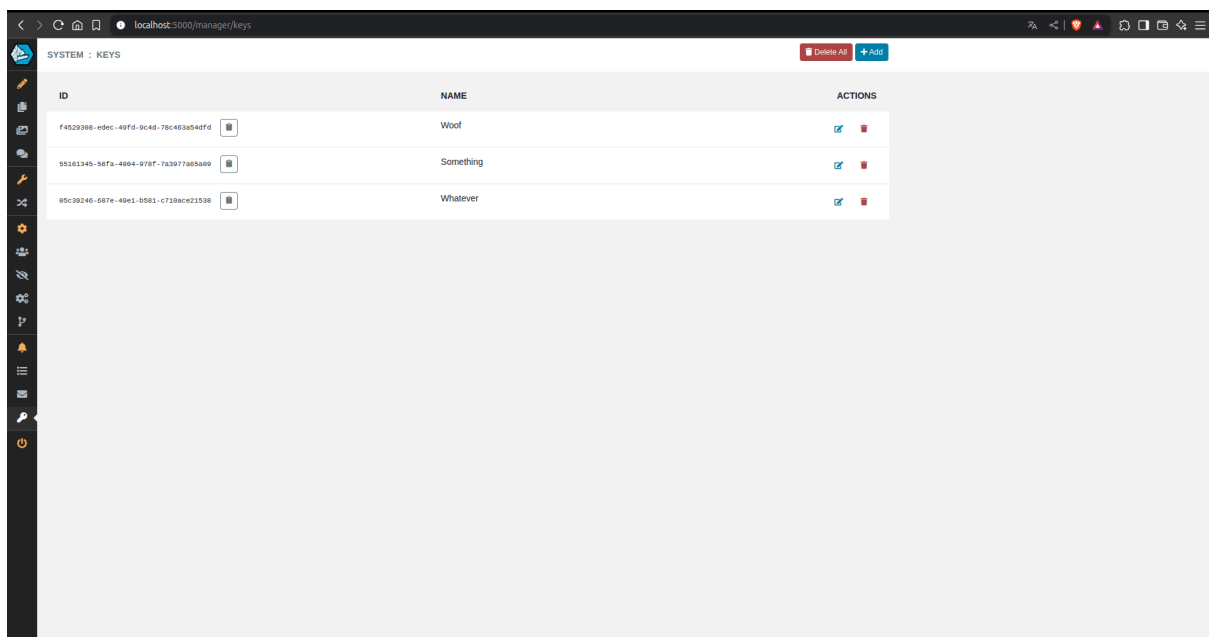


Figure 11: See all existing keys




Figure 12: Create a new key

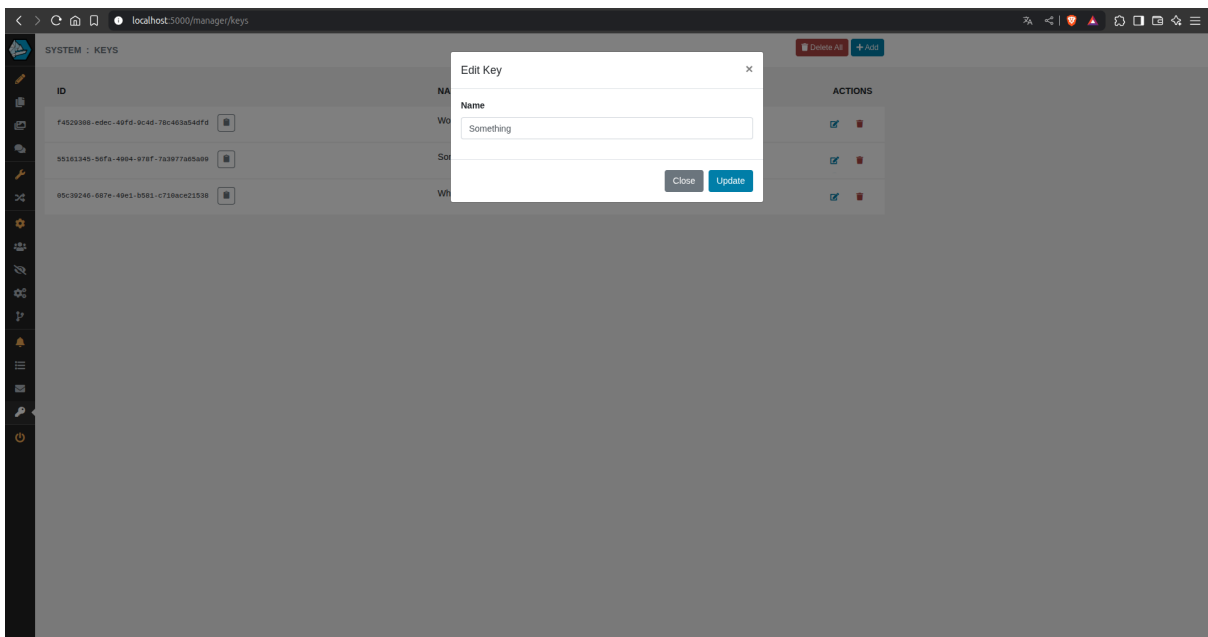


Figure 13: Update an existing key

7.3 External Publishing

Another key feature of our system is allowing trusted external parties to publish media directly into the CMS. To make this work, we had to implement two main changes:

- Enable admins to configure trusted external publishers.
- Expose a new API endpoint that allows those publishers to post content.

The first part is handled through the **Key Management module** (see Section 7.2.4). Using this module, an admin can create a new publisher entry and generate an API key for them (see Figure 12). This key must be provided by the external publisher in order to authenticate when posting content. This ensures that only authorized publishers — those with a valid key — can upload content to the CMS.

To support the actual publishing, we extended the CMS by adding a new `POST` endpoint to the `EventsApiController.cs`. This endpoint accepts media files and metadata from external sources, while checking for authorization using the API key.

We started by injecting the `KeyService` and writing a helper method to verify the provided key — either from the header or query string:

```
private async Task<bool> IsAuthorizedAsync(
    [FromHeader(Name = "X-API-Key")] string headerApiKey,
    [FromQuery(Name = "apiKey")] string queryApiKey)
{
    var apiKey = headerApiKey ?? queryApiKey;
    if (string.IsNullOrEmpty(apiKey))
        return false;

    var key = await _api.Keys.GetByIdAsync(Guid.Parse(apiKey));
    return key != null;
}
```

Next, we defined the `PublishMedia` endpoint. It accepts data in multipart form format: the media file(s), associated tags, and the API key:

```
[HttpPost("publish/")]
[Consumes("multipart/form-data")]
public async Task<IActionResult> PublishMedia(
    [FromForm] Piranha.Manager.Models.MediaUploadModel model,
    [FromForm] string Tags,
    [FromHeader(Name = "X-API-Key")] string headerApiKey,
    [FromQuery(Name = "apiKey")] string queryApiKey)
```

Once we have all the data, we validate the API key using the `IsAuthorizedAsync` method. If the key is invalid or missing, we immediately return an `Unauthorized` response. If it's valid, we process each uploaded file and save it to the CMS:

```
var uploaded = 0;
foreach (var upload in model.Uploads)
{
    if (upload.Length > 0 &&
        ↪ !string.IsNullOrEmpty(upload.ContentType))
    {
        using (var stream = upload.OpenReadStream())
        {
            await _api.Media.SaveAsync(new
                ↪ Piranha.Models.StreamMediaContent
            {

```

```

        Id = model.Uploads.Count() == 1 ? model.Id : null,
        FolderId = model.ParentId,
        Filename = System.IO.Path.GetFileName(upload.FileName),
        Data = stream,
        Tags = Tags
    });
    uploaded++;
}
}
}

```

If everything goes well, we return an HTTP 200 OK response with the number of files successfully published:

```

return Ok(new { success = uploaded > 0, message = $"{uploaded} file(s)
↪ published" });

```

Currently, we only support media uploads from external publishers due to time constraints. However, the same approach can be used to allow page publishing — the only difference would be the input format and how content is saved inside the CMS.

7.3.1 Webhook Delivery System

The final key feature in the system is the module responsible for sending notifications to clients. This module was built using the webhook pattern, which allows the system to send messages asynchronously to the client's backend whenever a new event is created. Notifications are sent via the POST URL provided by the client during the subscription process.

To implement this functionality, we created a Task in the PiranhaCMS Core System. This task consumes events from the event bus and processes each one. Notifications are only sent if the event matches both the **Type** and **Status** specified in a subscription. The first step is to retrieve all subscriptions of the relevant type, and then filter them by status—only subscriptions whose status matches the event's status will trigger a notification.

```

await _eventBus.StartConsumingAsync(async @event =>
{
    // Get all the subscriptions for the event type
    var subscriptions = await
    ↪ _api.Subscriptions.GetByEventTypeAsync(@event.Type.ToString());

    var subs = new List<Subscription>();
    foreach (var sub in subscriptions)
    {
        if (sub.EventStatus == @event.Status.ToString())
        {
            subs.Add(sub);
        }
    }
}

```

Next, we compare the tags of the event with those of each subscription to determine if there is at least one matching tag.

```
var filtered = new List<Subscription>();
foreach (var sub in subs)
{
    if (sub.Tags.Contains(','))
    {
        if (sub.Tags.Split(',').Any(tag =>
            ↪ @event.Tags.Contains(tag.Trim())))
        {
            filtered.Add(sub);
        }
    }
    else if (@event.Tags.Contains(sub.Tags.Trim()))
    {
        filtered.Add(sub);
    }
}
```

The final step is to retrieve the content associated with the event and send the notification along with that content. In the case of a file event, we call the function that generates a public URL to obtain a link to the file's content.

```
foreach (var subscription in filtered)
{
    object content = null;

    if (@event.Type == EventType.Page)
    {
        content = await
            ↪ _api.Pages.GetByIdAsync<PageInfo>(@event.ContentId);
    }
    else if (@event.Type == EventType.Media)
    {
        content = await _api.Media.GetByIdAsync(@event.ContentId);
        if (content is Media media)
        {
            // Ensure the media is fully loaded with all properties
            media.PublicUrl = _api.Media.GetPublicUrl(media);
        }
    }
    try
    {
        var sendEvent = new
        {
            @event.Id,
            @event.CreatedAt,
```

```

        Type = @event.Type.ToString(),
        Status = @event.Status.ToString(),
        @event.ContentId,
        @event.Tags
    };

    var payload = new
    {
        Event = sendEvent,
        Content = content,
    };
    await
    ↵ _api.Notifications.NotifyAsync(subscription.CallbackUrl,
    ↵ payload);

```

7.3.2 Demo Web Application

To test our solution, when considering external publishers, we decided to create a simple web application with an **Node JS** server. Where we have a simple **HTML** page and an endpoint to receive notifications.

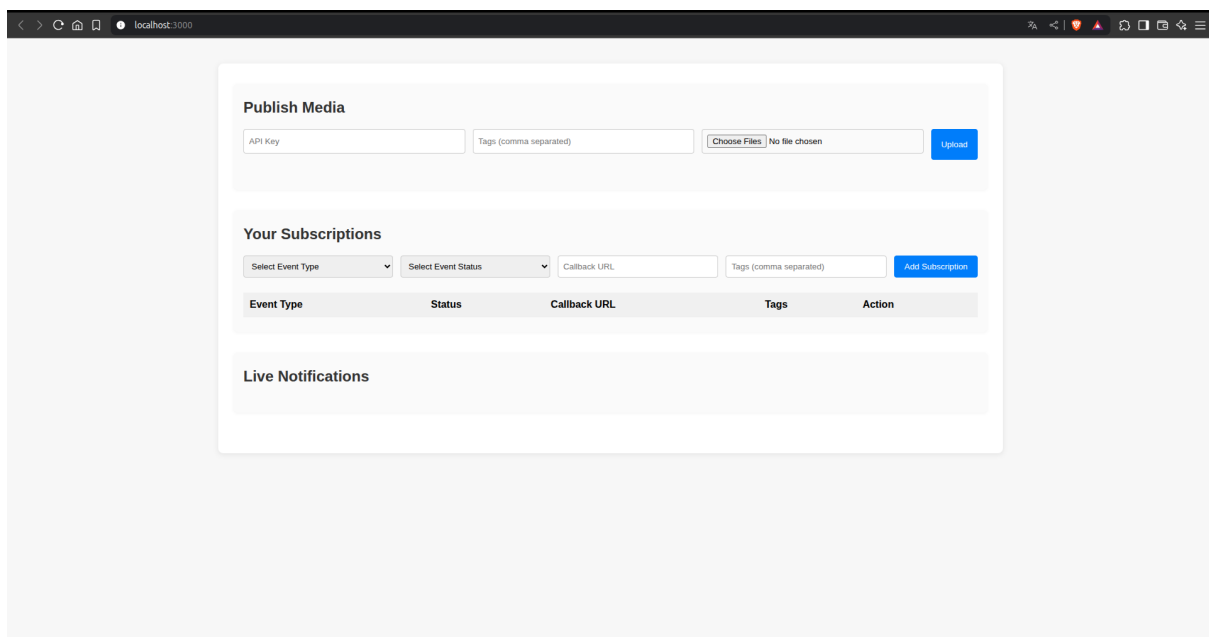


Figure 14: Demo Web Application developed

Our **HTML** page has three simple parts:

- **Publish Media:** A simple form that allows users of this application to upload media to the Piranha's system. It has three fields that need to be completed:
 1. **API Key:** This key must be provided by a system administrator and therefore it must exist in the Piranha's database.

2. **Tags:** The tags that the publisher wants to add to the media, so that the subscribers with the same tags can receive a notification about a new content published;
 3. **File:** A file to be sent and stored in the Piranha's system.
- **Your Subscriptions:** In this section it exists a form that allows a publisher to create a new subscription (equal to the one described in the **Subscription Module 7.2.2**). All fields must be completed and if everything is correct than a new subscription should appear in the table bellow the form.
 - **Live Notifications:** This is a simple section that displays all the notifications received. For notifications to appear it is required to have a subscription.

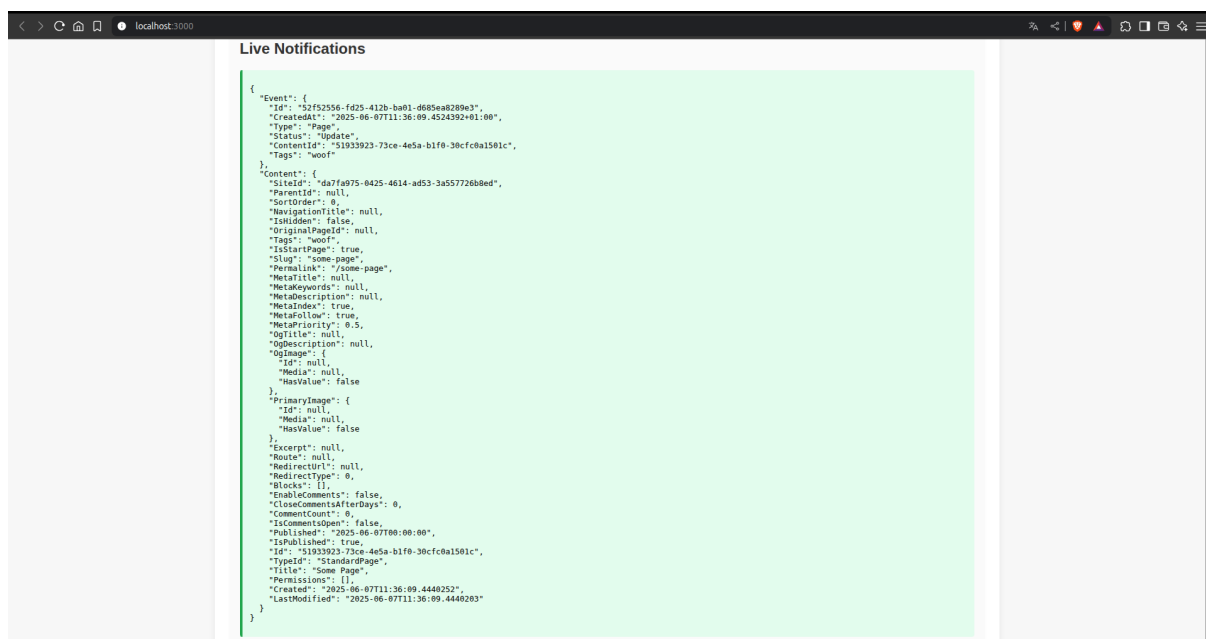


Figure 15: Example of a notification received in the Demo Web Application

The Node.js server uses **Express** and **Socket.IO** to handle real-time notifications and serve static files.

```
const express = require('express');
const http = require('http');
const socketIo = require('socket.io');
const app = express();
const server = http.createServer(app);
const io = socketIo(server);
const cors = require('cors');

app.use(express.json());
app.use(express.static('public'));
app.use(cors());
```



```

app.post('/notification', (req, res) => {
  const data = req.body;
  console.log('Notification received:', data);
  io.emit('notification', data);
  res.status(200).send({ status: 'ok' });
});

io.on('connection', (socket) => {
  console.log('A user connected');
})

server.listen(3000, () => console.log('Server is running on
↳ http://localhost:3000'));

```

Dependencies:

- **express**: Web framework used to handle HTTP requests and serve static files;
- **http**: Node's built-in module used to create the server
- **socket.io**: Enables real-time, bidirectional communication between clients and the server via WebSockets
- **cors**: Allows cross-origin requests, which is essential when the frontend and backend run on different origins.

Server Setup:

- An Express app is created and used to handle HTTP routes and middleware.
- The HTTP server is created using Node's **http** module and passed to the Socket.IO to enable WebSocket communication.
- The middleware includes:
 - **express.json()**: Parses incoming JSON requests;
 - **express.static('public')**: Serves static files from the **public** directory;
 - **cors()**: Enables Cross-Origin Resource Sharing.

Routes and WebSocket Events:

- **POST /notification**: Accepts notification data from clients. When a notification is received, it logs the data, emits it to all connected clients using **io.emit('notification', data)**, and responds with a success status.
- **io.on('connection')**: Listens for new WebSocket connections and logs a message when a user connects

Port and Startup

- The server listens on port **3000** and logs confirmation message once it's running.

7.4 Architecture

As we reached the end of the four-week period defined in the road-map, we entered the final cycle of the ADD process, reflecting on our progress and refining the system design. Through this last iteration, we analyzed the work completed, considered the constraints encountered, and incorporated the refinements made along the way.

As a result, we arrived at a **final** proposed architecture—an outcome shaped by all the development efforts, feedback loops, and design decisions made throughout the timeline of the road-map. This final architecture is presented as follows:

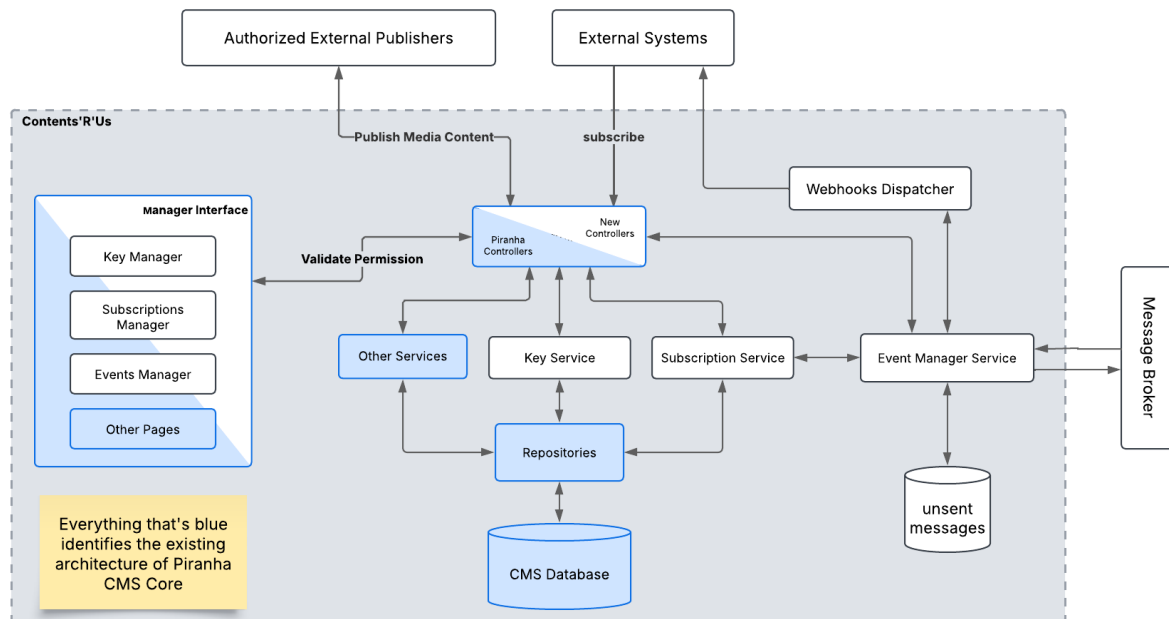


Figure 16: Proposed architecture

We've opted to revise our approach to the architecture by highlighting only the parts of Piranha that we modified. In the original architecture, we were still getting familiar with the internal workings of Piranha CMS, so it didn't accurately reflect the specific changes we made to Piranha Core. By omitting irrelevant sections, the new architecture is more streamlined and clearer, providing a better representation of our modifications. By positioning our changes alongside the existing Piranha structure, it's easier to see how our new services and admin pages fit into the system and what roles they serve.

In this version, we've also simplified the entities that interact with our system, now that we've clearly defined which entities will actually subscribe to and publish content. This updated architecture makes it easier to understand how external systems subscribe to our platform and receive webhooks, while external publishers simply add new content without subscribing.

References

- [1] Microsoft. *ASP.NET WebHooks Overview*. <https://learn.microsoft.com/en-us/aspnet/webhooks/>. Accessed: 2025-05-11. n.d.