

# Computação em Larga Escala

## Message Passage Interface (MPI)

Eurico Pedrosa

António Rui Borges

Universidade de Aveiro - DETI

2025-03-30

# Collective Communication


---

**Collective communication** refers to operations that involve **multiple processes** working together as part of a defined **communication group**. Unlike point-to-point communication, these operations enable more structured and efficient **data distribution and collection** among processes.

There are several types of collective communication in MPI:

- **Broadcast (MPI\_Bcast)**: A message is sent from a **root process** to **all other processes** within the communication group, including the root itself. **Used when the same data must be distributed to all processes.**
- **Scatter (MPI\_Scatter)**: A **root process** sends **distinct segments of data** to each process in the group, including itself. **Useful for parallelizing work where each process handles a portion of the data.**

- **Gather (MPI\_Gather):** Each process in the group sends its data to the **root process**, which collects all pieces into a single structure. **Ideal for collecting results computed in parallel.**



The **broadcast** operation (`MPI_Bcast`) enables a **single process**, identified as the **root**, to send the **same message (M)** to **all other processes** in a communication group, including itself.

### Key characteristics:

- The message  $M$  is sent by the **root process** (identified by its **rank**) to **all participating processes**.
- In the **standard implementation**, `MPI_Bcast` is a **blocking operation**: All processes **wait** until the message is fully received before proceeding.
- Conceptually, the operation involves:
  - The **root process** performing a **send**
  - All other processes performing a **receive**

### MPI\_Bcast – Function Signature

```
int MPI_Bcast(void *buffer, int count, MPI_Datatype datatype, int root, MPI_Comm comm);
```

This function broadcasts a message from the **root process** to **all processes** in the given **communicator**.

#### Parameters:

- **buffer**: A pointer to the **memory region** where the message data is stored (in the root), or where it will be received (in other processes).
- **count**: The **number of elements** in the message buffer.
- **datatype**: The **MPI data type** of the message content (e.g., MPI\_INT, MPI\_FLOAT).
- **root**: The **rank** of the process that will act as the broadcaster.
- **comm**: The **communicator** that defines the **group of processes** participating in the broadcast (typically MPI\_COMM\_WORLD).

#### Returns

- Returns MPI\_SUCCESS on success, or an error code otherwise.



```
#include <mpi.h>
#include <iostream>

int main(int argc, char** argv) {
    int rank, size, data;


    MPI_Init(&argc, &argv);           // Initialize the MPI environment
    MPI_Comm_rank(MPI_COMM_WORLD, &rank); // Get the rank of the process
    MPI_Comm_size(MPI_COMM_WORLD, &size); // Get the total number of processes

    if (rank == 0) {
        data = 42; // Root process sets the data
        std::cout << "Process " << rank << " broadcasting data = " << data << std::endl;
    }

    // Broadcast the value of 'data' from root (rank 0) to all processes
    MPI_Bcast(&data, 1, MPI_INT, 0, MPI_COMM_WORLD);

    // All processes (including root) now have the same value of 'data'
    std::cout << "Process " << rank << " received data = " << data << std::endl;

    MPI_Finalize(); // Finalize the MPI environment
    return 0;
}
```



The **scatter** operation (`MPI_Scatter`) allows the **root process** to distribute **distinct parts of a message** to **each process** in a communication group, including itself.

### Key Characteristics:

- The root process holds a **collection of k+1 messages**:

$$M_0, M_1, M_2, \dots, M_k$$

where  $k + 1$  is the **number of processes** in the communication group.

- Each process receives **one unique chunk** of the data:
  - Process 0 gets  $M_0$ , process 1 gets  $M_1$ , ..., process  $k$  gets  $M_k$ .
- By default, MPI\_Scatter is a **blocking operation**: All processes **wait** until their respective message part is received.
- Conceptually:
  - The **root process** performs multiple **sends** (one to each process).
  - All processes perform a **receive**.

**Example Use Case:** Distributing rows of a matrix or chunks of a large array to parallel workers for **independent computation**.

### MPI\_Scatter – Function Signature

```
int MPI_Scatter(  
    void *sendbuf, int sendcount, MPI_Datatype sendtype,  
    void *recvbuf, int recvcount, MPI_Datatype recvtype,  
    int root, MPI_Comm comm  
);
```

### Parameters:

- **sendbuf**: Pointer to the buffer holding all data to be sent. (**Significant only at the root process.**)
- **sendcount**: Number of elements to send **to each process**.
- **sendtype**: Data type of elements in the send buffer.
- **recvbuf**: Pointer to the buffer where each process will **store its received data**.
- **recvcount**: Maximum number of elements each process expects to receive.
- **recvtype**: Data type of elements in the receive buffer.
- **root**: Rank of the **root process**.
- **comm**: Communicator that defines the **group of processes**.

### MPI\_Scatterv – Function Signature

```
int MPI_Scatterv(  
    void *sendbuf, int *sendcounts, int *displs, MPI_Datatype sendtype,  
    void *recvbuf, int recvcount, MPI_Datatype recvtype,  
    int root, MPI_Comm comm  
);
```

### Extended Parameters:

- **sendcounts**: Array specifying the number of elements **to send to each process**.
- **displs**: Array specifying the **displacement** (offset) in sendbuf for each message. **(Each displacement is relative to the start of sendbuf.)**

This variant supports **non-uniform message sizes**, making it useful for **irregular data distributions**.

```
int main(int argc, char** argv) {
    MPI_Init(&argc, &argv);
    int rank, size;
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);

    const int chunk_size = 2;
    std::vector<int> recvbuf(chunk_size);

    std::vector<int> sendbuf;
    if (rank == 0) {
        sendbuf.resize(size * chunk_size);
        for (int i = 0; i < size * chunk_size; ++i)
            sendbuf[i] = i + 1;
    }
    MPI_Scatter(sendbuf.data(), chunk_size, MPI_INT,
               recvbuf.data(), chunk_size, MPI_INT,
               0, MPI_COMM_WORLD);

    std::cout << "Process " << rank << " received:";
    for (int val : recvbuf) std::cout << " " << val;
    std::cout << std::endl;

    MPI_Finalize();
    return 0;
}
```

```
int main(int argc, char** argv) {
    MPI_Init(&argc, &argv);
    int rank, size;
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);

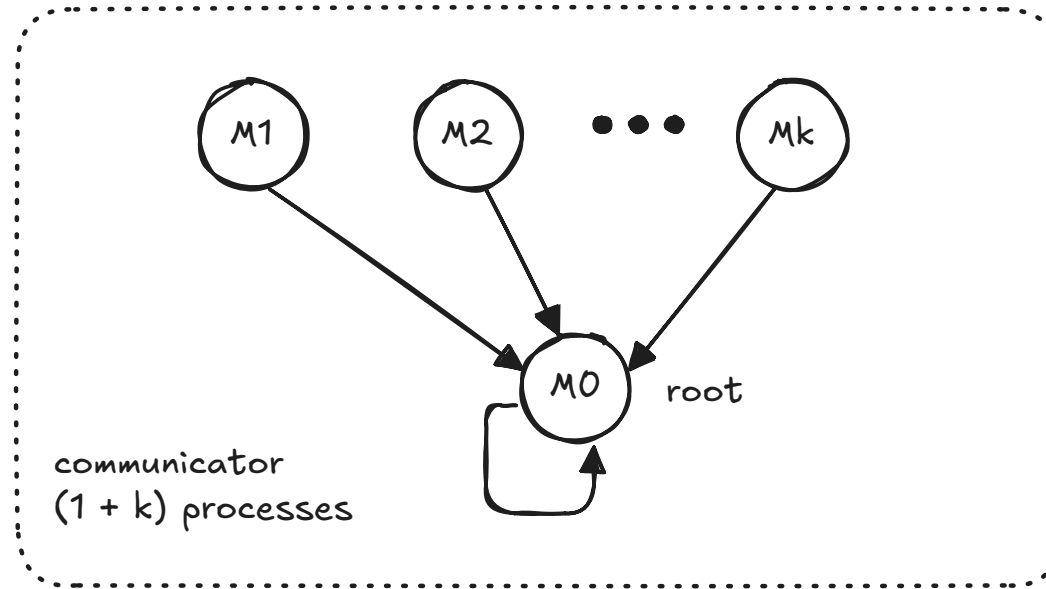
    int recvcount;
    std::vector<int> sendbuf, sendcounts(size), displs(size);
    std::vector<int> recvbuf;

    if (rank == 0) {
        int total = 0;
        for (int i = 0; i < size; ++i) {
            sendcounts[i] = i + 1; // different amount for each process
            displs[i] = total;
            total += sendcounts[i];
        }
        sendbuf.resize(total);
        for (int i = 0; i < total; ++i) sendbuf[i] = i + 1;
    }

    MPI_Scatter( sendcounts.data(), 1, MPI_INT, &recvcount, 1, MPI_INT, 0, MPI_COMM_WORLD);
    recvbuf.resize(recvcount);
    MPI_Scatterv(sendbuf.data(), sendcounts.data(), displs.data(), MPI_INT,
                recvbuf.data(), recvcount, MPI_INT, 0, MPI_COMM_WORLD);

    std::cout << "Process " << rank << " received:";
    for (int val : recvbuf) std::cout << " " << val;
    std::cout << std::endl;

    MPI_Finalize();
    return 0;
}
```



The **gather** operation (`MPI_Gather`) collects individual messages from **all processes** in a communication group and assembles them into a single buffer on the **root process**.



### Key Characteristics:

- Each process (including the root) sends a **message**:

$$M_0, M_1, M_2, \dots, M_k$$

where  $k + 1$  is the **group size**.

- The **root process** collects all messages and stores them **in order of process ranks** in a receive buffer.
- By default, MPI\_Gather is a **blocking operation**: The root process waits until it has **received all messages**.
- Conceptually:
  - All processes perform a **send**.
  - The **root process** performs a **receive** from each.

**Example Use Case:** Aggregating partial results (e.g., local sums, vectors) from all processes into a final array on the root for final processing or output.

### MPI\_Gather – Function Signature

```
int MPI_Gather(  
    void *sendbuf, int sendcount, MPI_Datatype sendtype,  
    void *recvbuf, int recvcount, MPI_Datatype recvtype,  
    int root, MPI_Comm comm  
);
```

### Parameters:

- **sendbuf**: Pointer to the message to be sent by each process.
- **sendcount**: Number of elements each process sends.
- **sendtype**: Data type of elements in sendbuf.
- **recvbuf**: Pointer to the buffer where the **root** will store the collected data. (**Significant only at root.**)
- **recvcount**: Number of elements received **from each process**. (**Significant only at root.**)
- **recvtype**: Data type of the receive buffer elements. (**Significant only at root.**)
- **root**: Rank of the root process.
- **comm**: MPI communicator.

### MPI\_Gatherv – Function Signature

```
int MPI_Gatherv(  
    void *sendbuf, int sendcount, MPI_Datatype sendtype,  
    void *recvbuf, int *recvcounts, int *displs, MPI_Datatype recvtype,  
    int root, MPI_Comm comm  
);
```

### Extended Parameters:

- **recvcounts**: Array specifying how many elements are received from each process.
- **displs**: Array specifying **displacements** (offsets) in recvbuf where each message should be stored.

MPI\_Gatherv is used for **non-uniform data gathering**, where each process might send a **different number of elements**.

```
int main(int argc, char** argv) {
    MPI_Init(&argc, &argv);
    int rank, size;

    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);

    int sendval = rank + 1;
    std::vector<int> recvbuf;

    if (rank == 0) recvbuf.resize(size); // collect one int from each process

    MPI_Gather(&sendval, 1, MPI_INT,                // each process sends one int
              recvbuf.data(), 1, MPI_INT,            // root receives one int from each process
              0, MPI_COMM_WORLD);

    if (rank == 0) {
        std::cout << "Gathered values:";
        for (int val : recvbuf) std::cout << " " << val;
        std::cout << std::endl;
    }

    MPI_Finalize();
    return 0;
}
```

```
int main(int argc, char** argv) {
    MPI_Init(&argc, &argv);
    int rank, size;
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);


    // Each process sends (rank + 1) values
    std::vector<int> sendbuf(rank + 1);
    for (int i = 0; i <= rank; ++i) sendbuf[i] = (rank + 1) * 10 + i;

    std::vector<int> recvcnts, displs, recvbuf;
    if (rank == 0) {
        recvcnts.resize(size);
        displs.resize(size);
        int offset = 0;
        for (int i = 0; i < size; ++i) {
            recvcnts[i] = i + 1;
            displs[i] = offset;
            offset += recvcnts[i];
        }
        recvbuf.resize(offset); // total size for root to collect everything
    }

    MPI_Gatherv(sendbuf.data(), sendbuf.size(), MPI_INT,
                recvbuf.data(), recvcnts.data(), displs.data(), MPI_INT,
                0, MPI_COMM_WORLD);

    if (rank == 0) {
        std::cout << "Gatherv result:";
        for (int val : recvbuf) std::cout << " " << val;
        std::cout << std::endl;
    }

    MPI_Finalize();
    return 0;
}
```



The **MPI\_Reduce** operation allows processes to perform a **global element-wise computation** on their data and deliver the final result to a designated **root process**.

### Key Characteristics:

- Each process provides a **value array**:

$$V_0, V_1, V_2, \dots, V_k$$

where  $k + 1$  is the **number of processes** in the communication group.

- These arrays are **combined element-wise** using a **commutative binary operator**, such as:
  - MPI\_SUM, MPI\_MAX, MPI\_MIN, MPI\_PROD, etc.
- The **resulting array** is sent to the **root process**.

### Blocking Semantics:

- The operation is **blocking by default**: The **root process** waits until **all values have been received and reduced**.

### Conceptual Model:

- All processes **send** their local arrays.
- The **root process** performs a **receive**, then reduces:

$$\text{Result}[i] = V_0[i] \text{ op } V_1[i] \text{ op } \dots \text{ op } V_{k[i]}$$

Where op is a **commutative binary operation**, applied element-wise.



### MPI\_Reduce – Function Signature

```
int MPI_Reduce(  
    void *sendbuf, void *recvbuf, int count,  
    MPI_Datatype datatype, MPI_Op op,  
    int root, MPI_Comm comm  
);
```

### Parameters:

- **sendbuf**: Pointer to the local data (input values) to be reduced. (**Each process sends its own buffer.**)
- **recvbuf**: Pointer to the buffer where the **root process** stores the result. (**Ignored by non-root processes.**)
- **count**: Number of elements in the array being reduced.
- **datatype**: MPI data type of the elements (e.g., MPI\_INT, MPI\_FLOAT).
- **op**: The **reduction operation** to apply. Must be a predefined **commutative** operator like:
  - MPI\_SUM, MPI\_PROD, MPI\_MAX, MPI\_MIN
  - Bitwise: MPI\_BAND, MPI\_BOR, MPI\_BXOR
- **root**: Rank of the process that receives the **final reduced result**.
- **comm**: The MPI communicator (e.g., MPI\_COMM\_WORLD) that defines the group of processes.

```
int main(int argc, char** argv) {
    MPI_Init(&argc, &argv);
    int rank, size;
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);

    const int N = 5;
    std::vector<int> local_array(N);
    std::vector<int> result_array(N); // only used at root
    // Seed the random number generator differently per process
    std::srand(static_cast<unsigned int>(std::time(0)) + rank);
    // Fill local array with random values between 1 and 10
    for (int i = 0; i < N; ++i)
        local_array[i] = std::rand() % 10 + 1;

    // Print each process's local data
    std::cout << "Process " << rank << " local array:";
    for (int val : local_array) std::cout << " " << val;
    std::cout << std::endl;

    // Perform reduction (sum) across all arrays
    MPI_Reduce(local_array.data(), result_array.data(),
               N, MPI_INT, MPI_SUM, 0, MPI_COMM_WORLD);

    if (rank == 0) {
        std::cout << "Root received reduced array (sum):";
        for (int val : result_array) std::cout << " " << val;
        std::cout << std::endl;
    }

    MPI_Finalize();
    return 0;
}
```

# Suggested Reading

---

- The Art of HPC by Victor Eijkhout of TACC
  - Volume 2: Parallel Programming for Science and Engineering
- Rookie HPC (MPI Documentation)
  - <https://rookiehpc.org/mpi/docs/index.html>