

Container Virtualization

Gestão de Infraestruturas de Computação

deti universidade de aveiro
departamento de eletrónica,
telecomunicações e informática

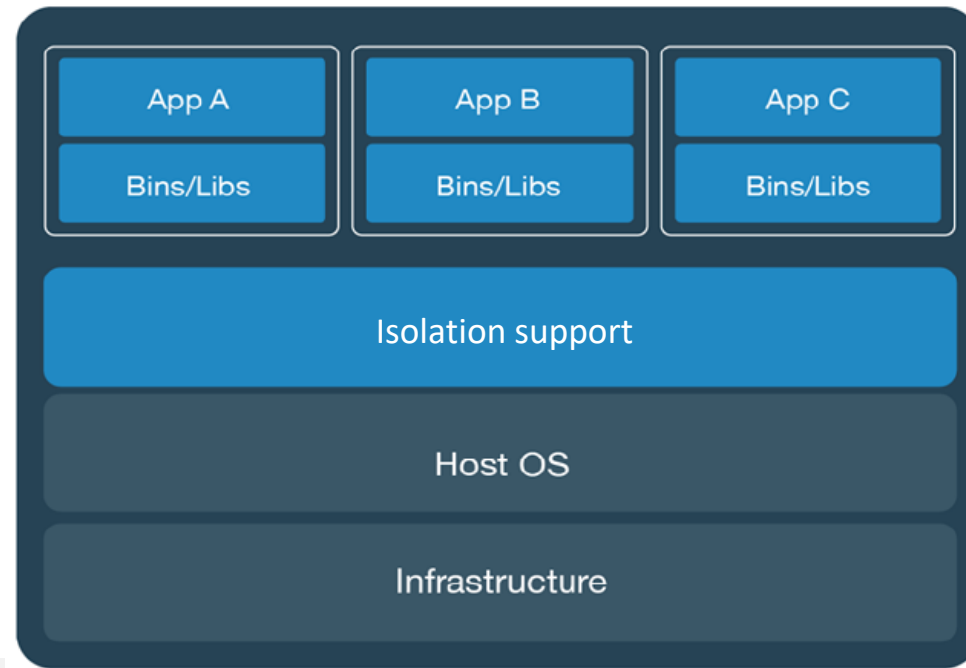
João Paulo Barraca

OS Level Virtual Machines

- **Do not provide a full virtualized hardware**
 - Only OS services are virtualized
 - Host kernel: virtualize its services in an isolated manner to different guests
- **Container: isolated execution environment to encapsulate one or more processes**
- **Relevant aspects:**
 - Isolation/Visibility: limit what can be seen by tenants
 - Resource control: limit resource consumption
 - Portability: reconstruct the same environment in multiple hosts

Container based virtualization

- Container based virtualization relies on OS mechanisms for enhanced isolation



OS Support

- **Linux has no support for containers!**
 - That is: the Linux OS does not know what is a container
- **Containers are a set of configurations, creating an environment, applied to a process (and its children)**
 - Namespaces: isolation and virtualization of each tenant
 - Cgroups: resource control of CPU, RAM, IO of each tenant

Underlying technologies

- **Namespaces**
 - (mnt, pid, net, ipc, uts/hostname, user ids)
- **cgroups**
 - (cpu, memory, disk, i/o - resource management)
- **AppArmor, SELinux**
 - (security/access control)
- **seccomp**
 - (computation isolation)
- **chroot**
 - (file system isolation)

Namespaces

- **Mechanism used to isolate and virtualize system resources**
 - Processes in a namespace are unable to see remaining resources
 - Their namespace looks like the entire host
 - Processes in a namespace access resources as if there is no concurrency
- **Network namespace:**
 - Mechanism to create a somewhat independent set of networking resources
 - Network interfaces, routing tables
 - Network interfaces can only belong to a single namespace
- **PID namespace:**
 - Restricted namespace with private process identifiers
 - Host processes are invisible

cgroups

- **Mechanism used to restrict (limit, control) or monitor the amount of resources used by “groups of processes”**
 - Processes can be organized in groups, to control their accesses to resources
- **Example: CPU control groups for scheduling**
 - Limit the amount of CPU time that processes can use, etc...
- **Similar cgroups for other resources**
 - memory, IO, pids, network, ..

Setting up a container

1. Setup all the needed namespaces and control groups

- Usually according to some template

2. Create a “disk image” for the container

- directory containing the container’s filesystem
- will exist in the host (it’s a directory)

3. Chroot to the container filesystem

- Must contain all the libraries/files needed to execute the program

4. Start the program to containerize

- This process it will have PID 1 in the container
- Note: this process can mount procfs or other pseudo-fileSYSTEMS
 - Namespaces allow to control the information exported in those pseudofilesystems

Setting up a container

- **Due to network namespace, containerized processes do not see the host's network interfaces**
 - But we usually do not assign a real network interface to a namespace
 - Would have impact to the host
- **Networking with containers:**
 - Create a virtual ethernet pair: Two virtual interfaces, connected point-to-point
 - Packets sent on one interface are received on the other, and vice-versa
 - Associate one virtual interfaces to the network namespace of the container
 - Bind the other one to a software bridge

Containers are not Hardware VMs

- Core idea: a container is a **Process in a sandbox**, not a fully emulated virtual hardware

| Virtual Machine | Container |
|---|---|
| Each VM runs a different OS | All containers share the same OS |
| Booting the VM involves the Full boot of the guest HW | Containers boot in seconds (start service) |
| VM snapshots are made on demand | Containers use multiple images in layers |
| VM description is a XML document | Container description is a sequence of instructions |
| VM will pre-consume RAM and CPU | Containers have almost zero overhead |
| A VM description will result in one VM | A dockerfile may be used to start many containers |
| Has limits of tens of VMs per CPU core | Can have hundreds of containers per CPU core |

Tools

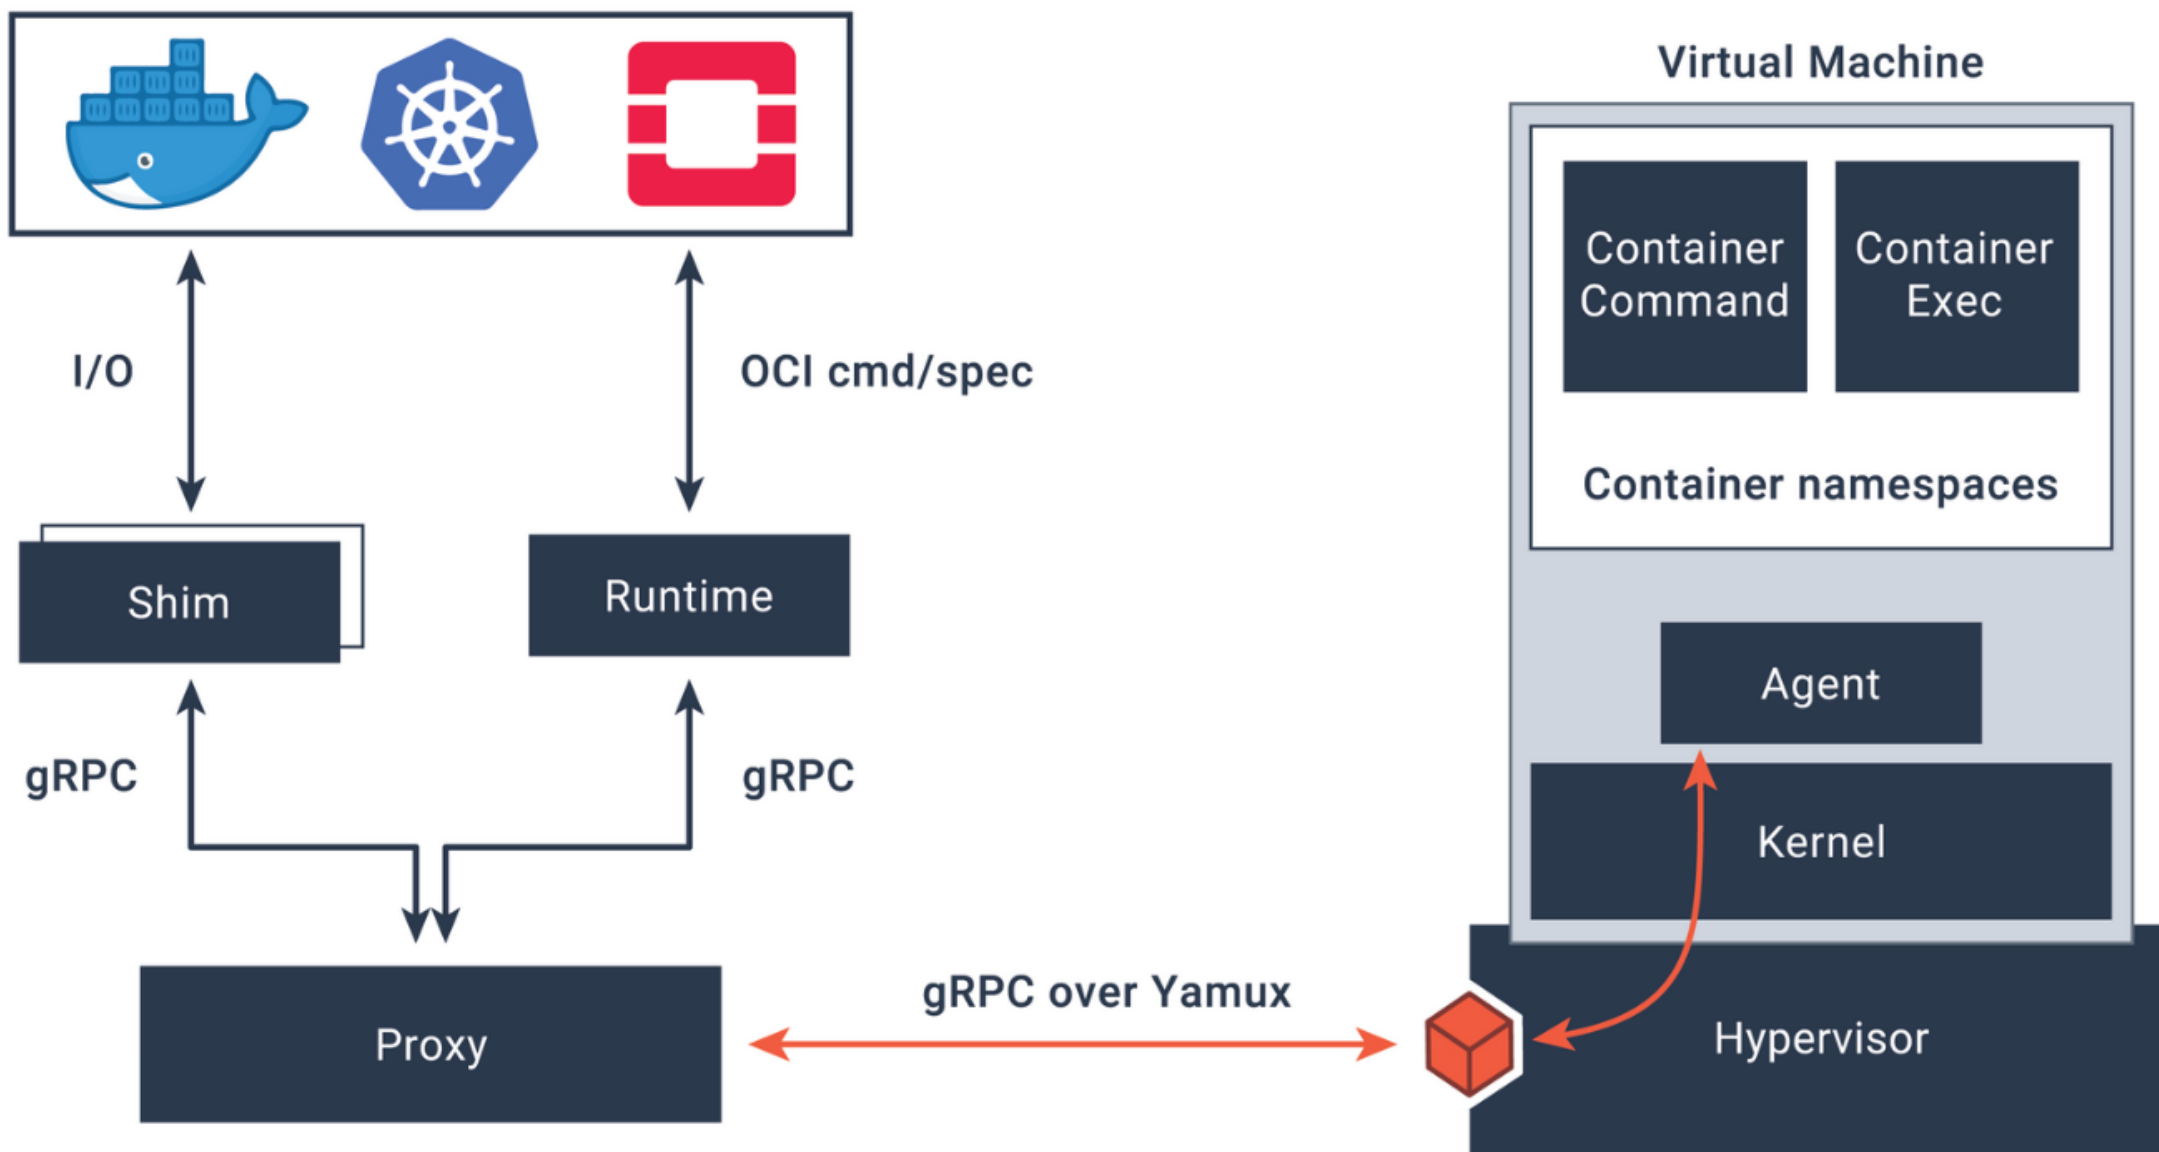
- **Creating containers by hand is a difficult task**
 - Reasonably high set of instructions to create namespaces, cgroups, chroots, interfaces, bridges
- **User space tools facilitate manipulation of containers**
 - LXC: Linux Containers
 - Docker
 - Singularity
 - Kubernetes


Portability

- **Containers are all about creating replicable execution environments**
 - Same environment (namespaces, chroot, cgroups) that can be used in different hosts
 - Important to have a common description of containers
- **Open Container Initiative (OCI): <https://www.opencontainers.org>**
 - Defines standards for user-space tools
 - Runtime specification: configuration, execution, and lifecycle of a container
 - Image specification: how to represent the data of a container

Portability

- **Containers are evolving beyond OS level virtualization**
 - If runtime and images are standardized, implementation doesn't matter.
 - It is not relevant if the container will run on OS Level Virtualization or on a HW VM
- **Containers are evolving to Light Weight Virtual Machines**
 - Namespaces: can be implemented as a Virtual Machines
 - Cgroups: can be implemented as the Virtual Machines with control tools
 - Images: can be implemented as virtual disks
 - Example: Kata containers



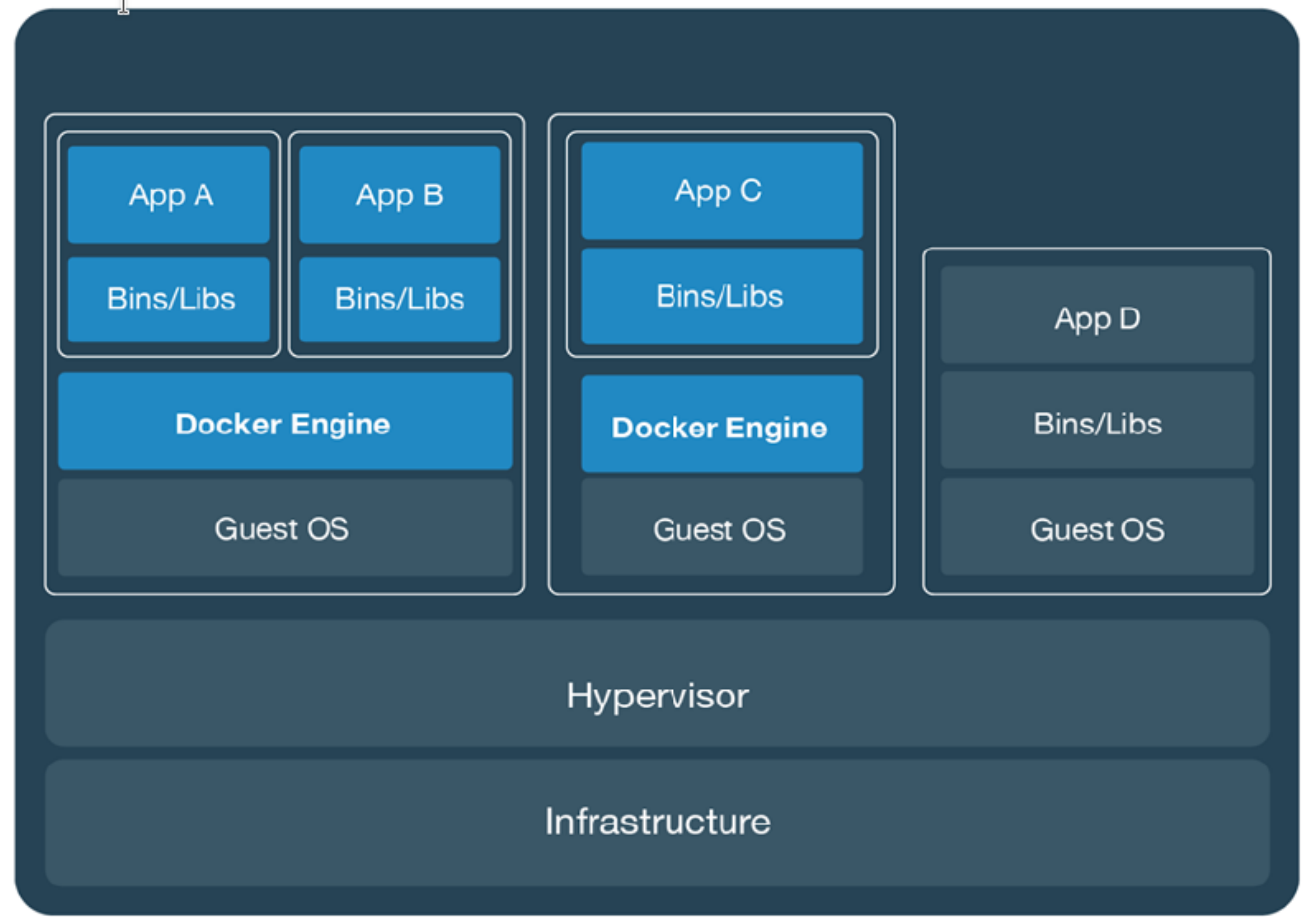
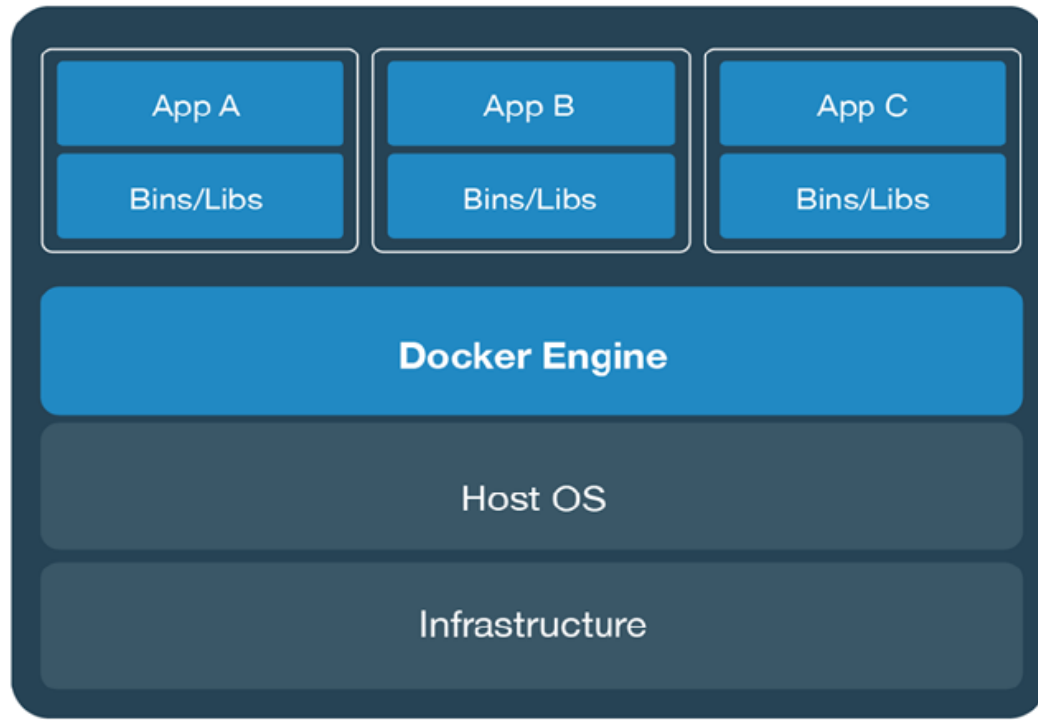


- **Commercial Product**
 - “Evolution” from Solaris Containers and Linux Containers (LXC)
 - At a conceptual level
 - Released as Open Source in 2013
- **Core Concept: Container**
 - App level construct (unrelated to infrastructure)
 - Composed by multiple functionality (namespaces) in coordination
 - One Container runs One Application (which can spawn child apps)
- **Containers pack an application to run in any underlying hardware**
 - including configurations, dependencies, auxiliary data, etc...

Docker Concepts

- Image: **the data** of a container (application, libraries, images, etc...)
- Container: A **running instance** of an application.
 - Composed by one or more images. Each image adds a specific functionality
- Engine: Software that executes commands for containers
- Registry: **Repository** of docker images
- Control Plane: infrastructure to manage containers and images

Containers can run with VMs





Linux Containers



liblxc

namespaces

cgroups

SELinux/AppArmor

Linux kernel



Docker 1.10 and later



runC

runC

runC

containerd-shim

containerd-shim

containerd-shim

containerd

Docker Engine

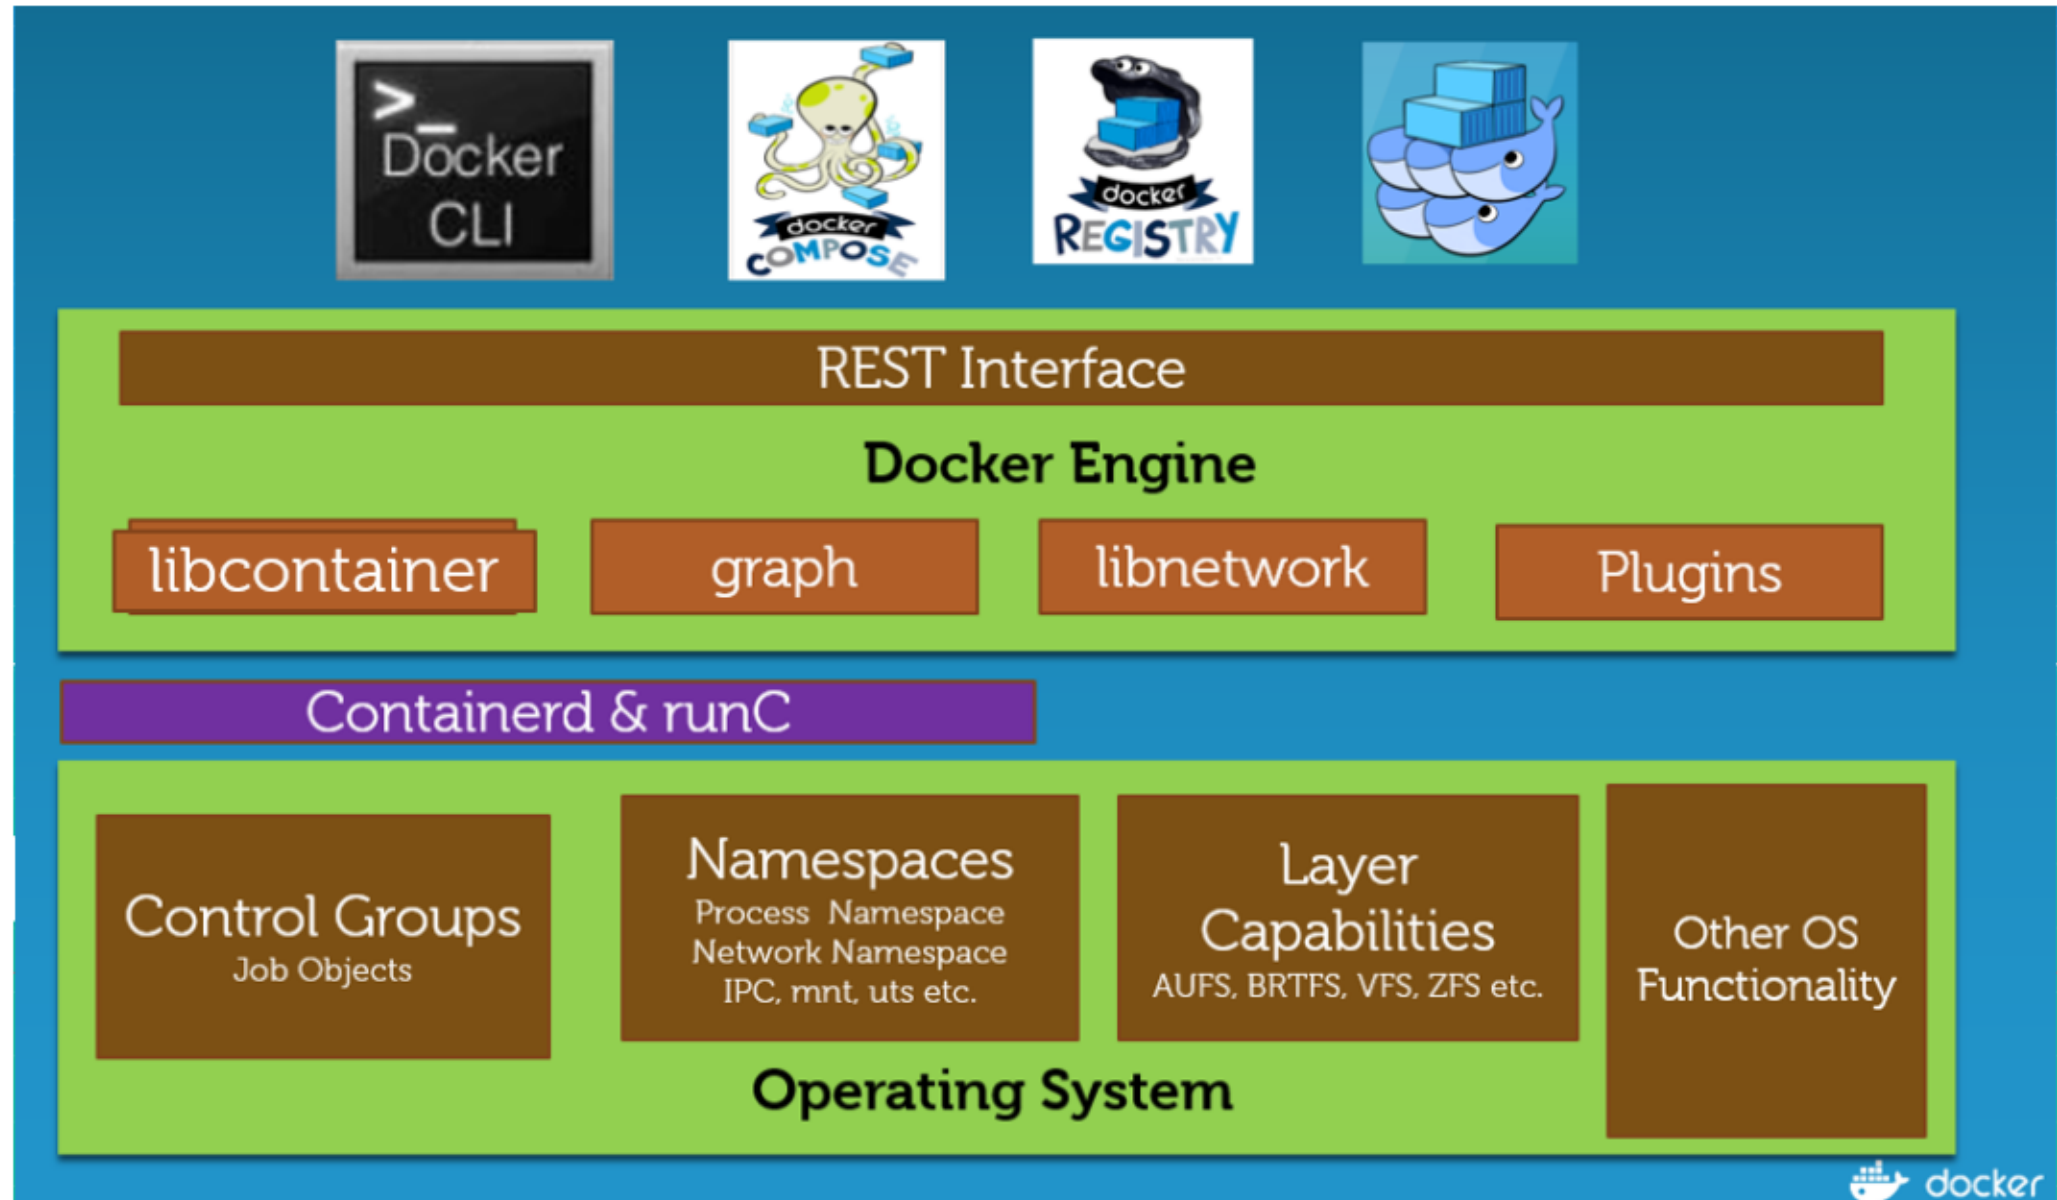
namespaces

cgroups

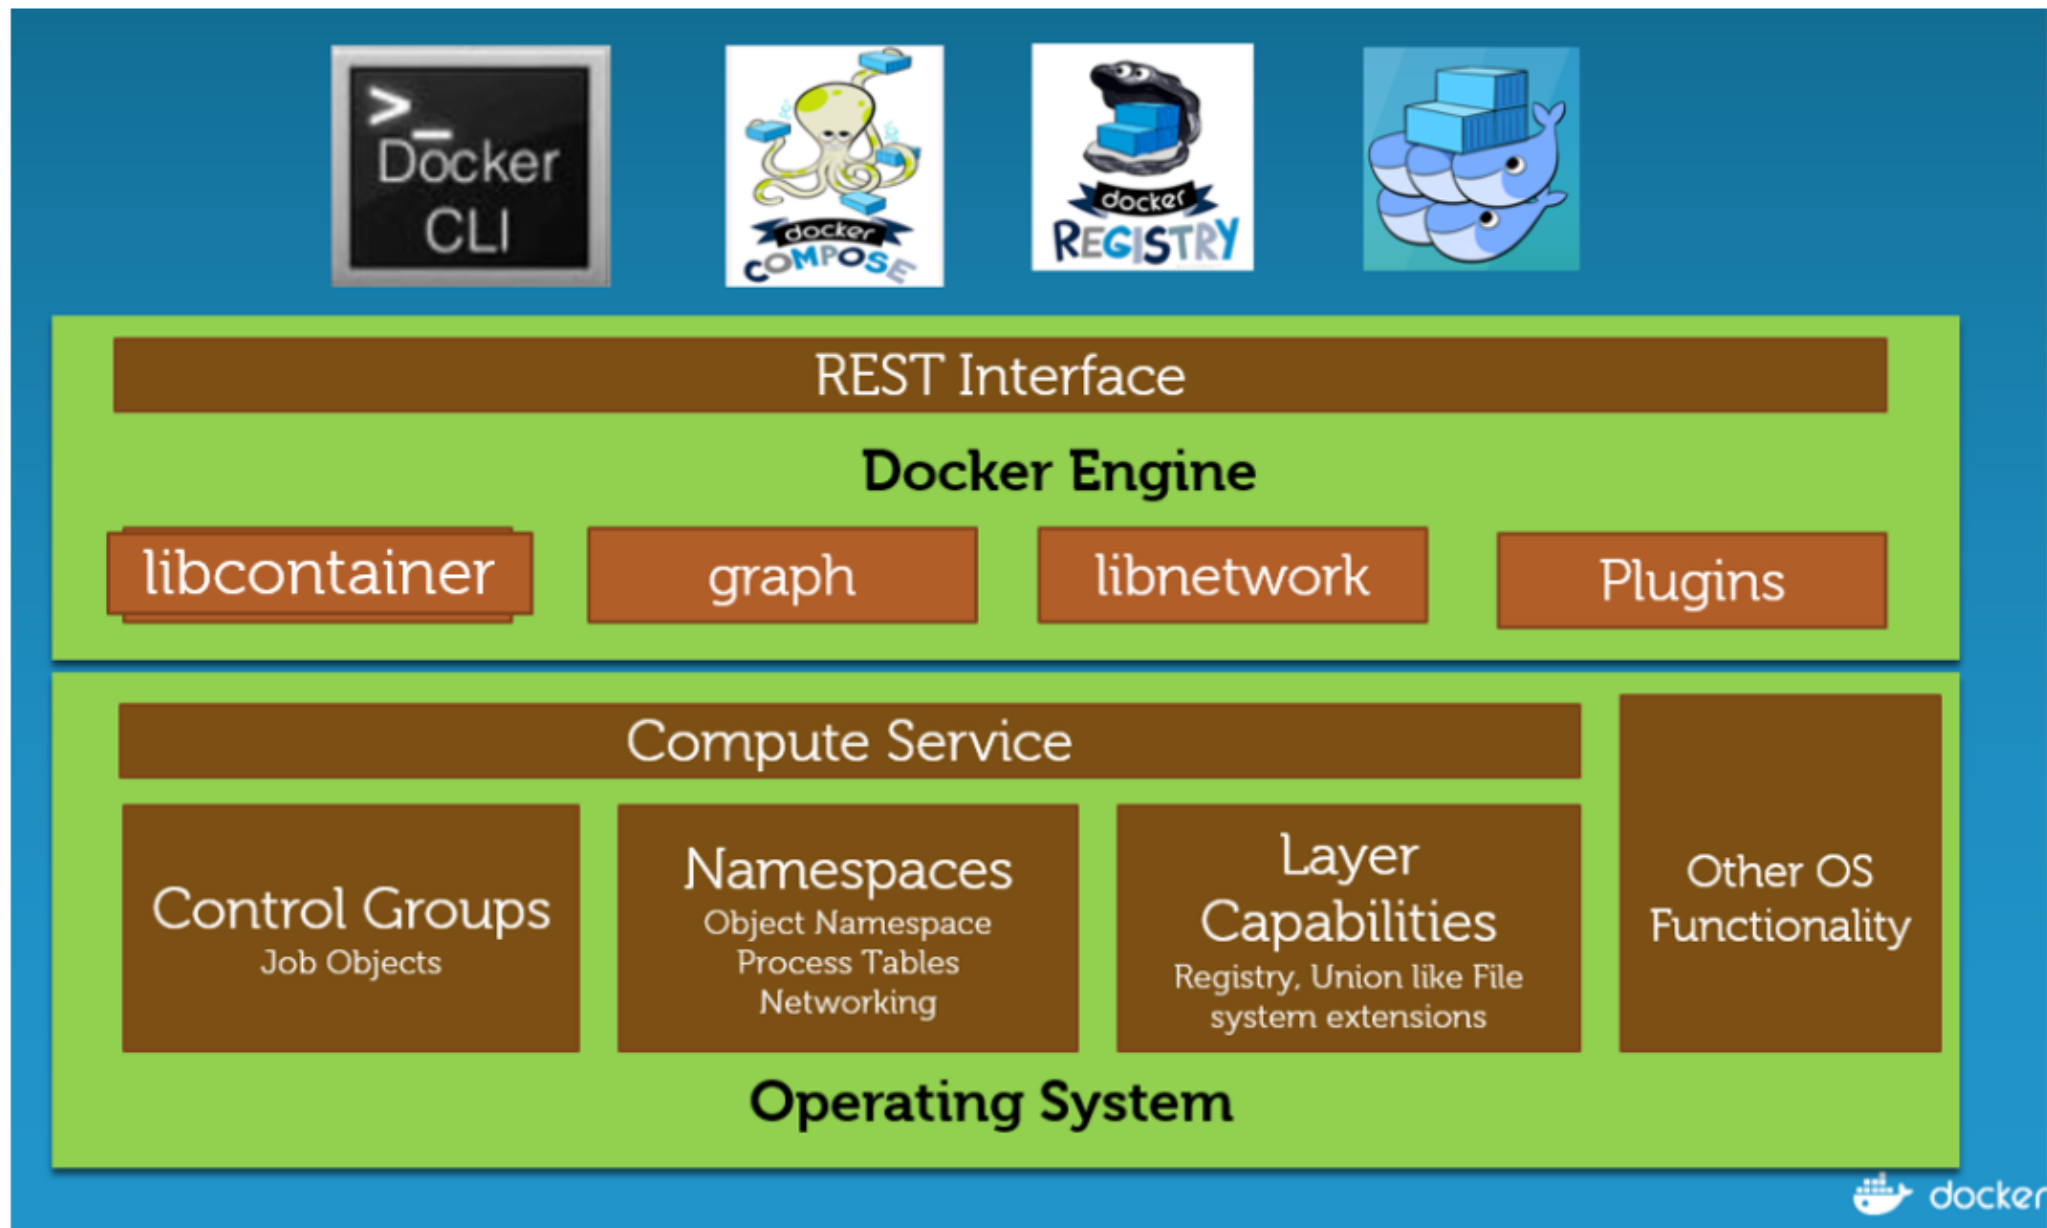
SELinux/AppArmor

Linux kernel

Docker Engine on Linux Platform



Docker Engine on Windows Platform



Docker Registry

- **Central repository to store and deliver images**
 - indexed by **name** and **tag**
 - can be private to an infrastructure or public (Docker Hub)
- **Clients push images to the registry**
 - Local client
- **Docker Engine pulls images and executes the container**
 - Running on servers
- **Layers are hashed and indexed allowing reuse**
 - Pulling an image only requires the layers that do not exist locally

Docker Hub - Largest Public Registry

Name

nginx ☆

NGINX Docker Official Images
Official build of Nginx.

1B+

Container Linux 386 x86-64 ARM 64 IBM Z ARM PowerPC 64 LE Application Infrastructure

Official Image

Linux - x86 (latest)

Copy and paste to pull this image

```
docker pull nginx
```

[View Available Tags](#)

Description Reviews Tags

Supported tags and respective Dockerfile links

- 1.17.9, mainline, 1, 1.17, latest
- 1.17.9-perl, mainline-perl, 1-perl, 1.17-perl, perl
- 1.17.9-alpine, mainline-alpine, 1-alpine, 1.17-alpine, alpine

Tags

root@server:~# docker search nginx

| NAME | DESCRIPTION | STARS | OFFICIAL | AUTOMATED |
|----------------------------------|---|-------|----------|-----------|
| nginx | Official build of Nginx. | 12788 | [OK] | |
| jwilder/nginx-proxy | Automated Nginx reverse proxy for docker con... | 1750 | | [OK] |
| richarvey/nginx-php-fpm | Container running Nginx + PHP-FPM capable of... | 758 | | [OK] |
| linuxserver/nginx | An Nginx container, brought to you by LinuxS... | 95 | | |
| bitnami/nginx | Bitnami nginx Docker Image | 77 | | [OK] |
| tiangolo/nginx-rtmp | Docker image with Nginx using the nginx-rtmp... | 64 | | [OK] |
| jc21/nginx-proxy-manager | Docker container for managing Nginx proxy ho... | 45 | | |
| nginxdemos/hello | NGINX webserver that serves a simple page co... | 41 | | [OK] |
| nginx/unit | NGINX Unit is a dynamic web and application ... | 36 | | |
| jlesage/nginx-proxy-manager | Docker container for Nginx Proxy Manager | 35 | | [OK] |
| nginx/nginx-ingress | NGINX Ingress Controller for Kubernetes | 28 | | |
| privatebin/nginx-fpm-alpine | PrivateBin running on an Nginx, php-fpm & Al... | 22 | | [OK] |
| schmunk42/nginx-redirect | A very simple container to redirect HTTP tra... | 18 | | [OK] |
| blacklabelops/nginx | Dockerized Nginx Reverse Proxy Server. | 13 | | [OK] |
| nginxinc/nginx-unprivileged | Unprivileged NGINX Dockerfiles | 13 | | |
| centos/nginx-112-centos7 | Platform for running nginx 1.12 or building ... | 12 | | |
| raulr/nginx-wordpress | Nginx front-end for the official wordpress:f... | 12 | | [OK] |
| centos/nginx-18-centos7 | Platform for running nginx 1.8 or building n... | 12 | | |
| nginx/nginx-prometheus-exporter | NGINX Prometheus Exporter | 9 | | |
| sophos/nginx-vts-exporter | Simple server that scrapes Nginx vts stats a... | 7 | | [OK] |
| mailu/nginx | Mailu nginx frontend | 6 | | [OK] |
| bitnami/nginx-ingress-controller | Bitnami Docker Image for NGINX Ingress Contr... | 4 | | [OK] |
| ansibleplaybookbundle/nginx-apb | An APB to deploy NGINX | 1 | | [OK] |
| wodby/nginx | Generic nginx | 0 | | [OK] |
| centos/nginx-110-centos7 | Platform for running nginx 1.10 or building ... | 0 | | |

Workflow

- Search docker image

\$ docker search nginx

- Pull docker image

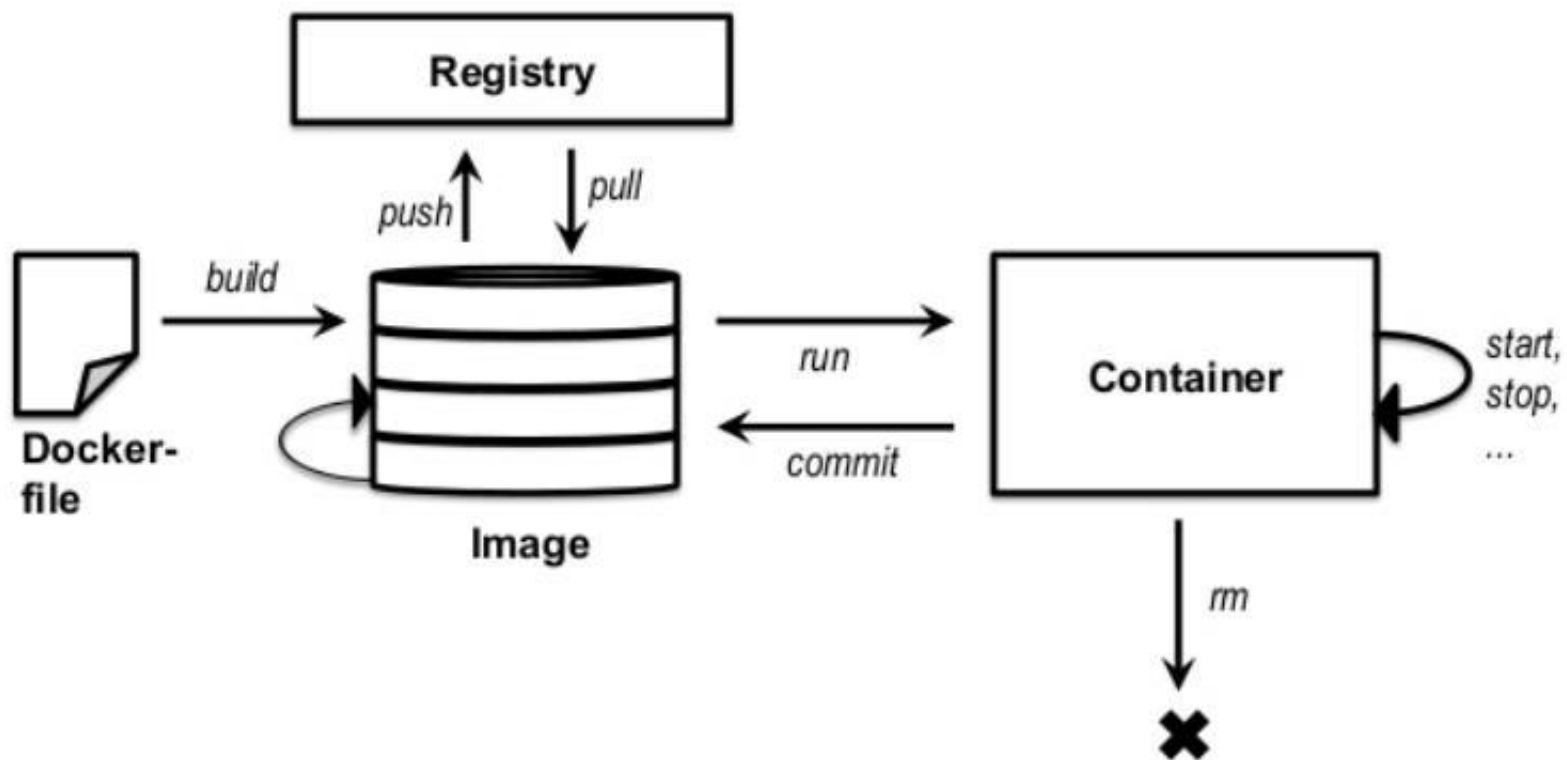
\$ docker pull nginx

- List local docker images

\$ docker image ls

- Run container

\$ docker run -d --name frontend-server nginx



Docker file

- **Sequence of commands to prepare the container for execution**
 - Lists dependencies
 - Lists commands to be executed inside container
 - Sets the command to execute on start
- **Used locally or distributed in registry**
 - Public registry: Docker hub
- **May define versions for same software**
 - nginx:latest, nginx:perl, nginx:alpine...

```
FROM debian:stretch-slim

LABEL maintainer="NGINX Docker Maintainers <docker-maint@nginx.com>"

ENV NGINX_VERSION 1.15.9-1~stretch
ENV NJS_VERSION 1.15.9.0.2.8-1~stretch

RUN set -x \
    && apt-get update \
    && apt-get install --no-install-recommends --no-install-
suggests -y gnupg1 apt-transport-https ca-certificates \
    && \

...OMMITED...

RUN ln -sf /dev/stdout /var/log/nginx/access.log \
    && ln -sf /dev/stderr /var/log/nginx/error.log

EXPOSE 80

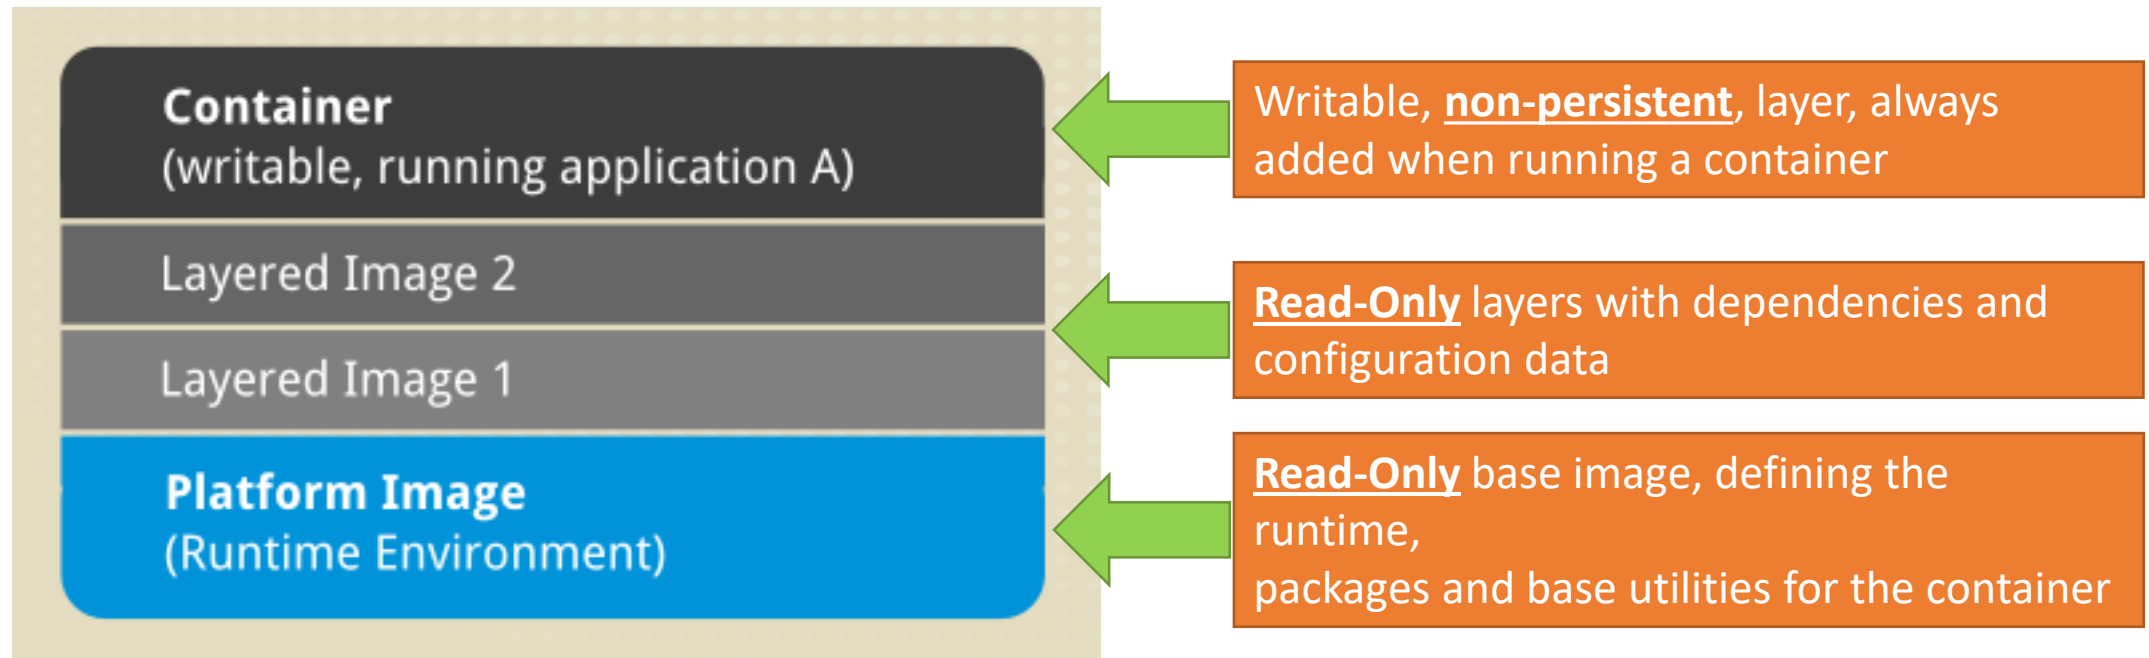
STOPSIGNAL SIGTERM

CMD ["nginx", "-g", "daemon off;"]
```

Images

- **Containers have their initial data in images**
 - Contain data, applications, libraries
 - Contain an entry point to be executed when the container is initiated
- **Composed by multiple layers**
 - A base image
 - Several images, changing the content (with differences)
 - Exploits overlayfs (a unionfs similar to AUFS)
 - Can use others like btrfs, ZFS
- **Read Only and Non persistent**
 - Serve as templates for containers to start
 - Multiple containers will use the same images (containers are ephemeral)

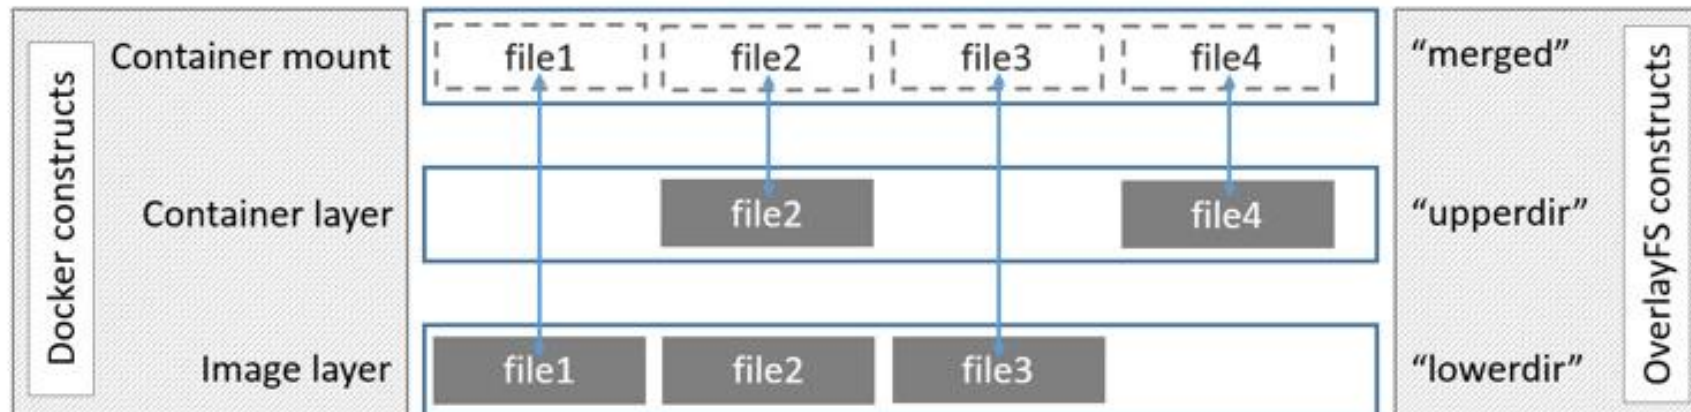
Image layering



Available at `/var/lib/docker/...`

Images

- **It's important to keep images small**
 - Dont: base your image on a full blown distro (e.g. debian)
 - Do: base your image on a small distro (e.g. Alpine)
- **Presented on the host**
 - `/var/lib/docker/overlay2/<ID>/`
 - `/merged`: FS seem inside container
 - `/diff`: differences from base


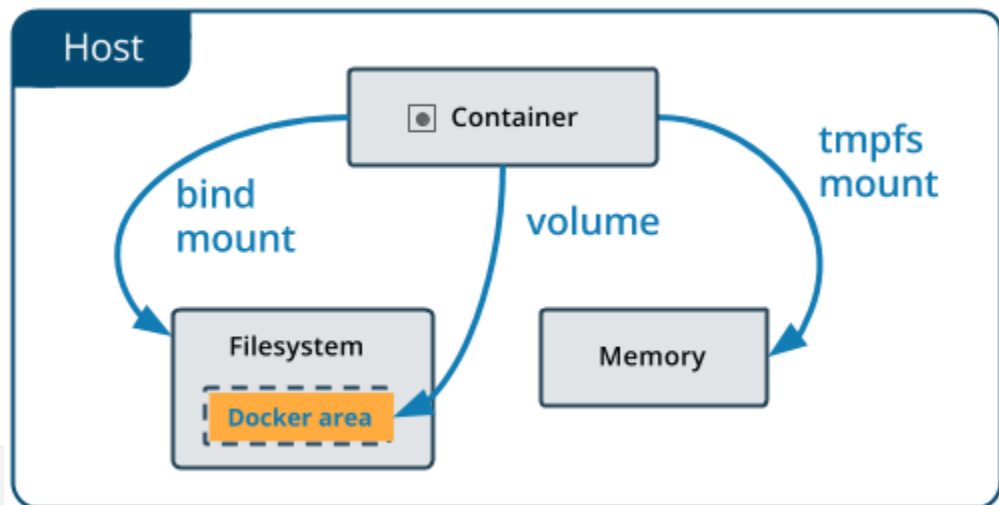


Images

- **OverlayFS has limitations!**
 - Only a subset of the calls is supported
 - May create issues with some applications or decrease performance
- **Issues:**
 - open: opening a file RW after the same file is open as R will fail.
 - first open (R) will refer to a file in the base image
 - second open (RW) will result in a copy of the file being created (stored in the diff folder)
 - File descriptions will not point to the same file
 - rename: not supported
 - applications must copy and then unlink
 - much slower and will take additional disk space

Container Persistence

- **Top layer is writable, but not persisted!**
 - Stopping the container will result in data loss
- **Persistence must be achieved through external resources**
 - Bind mount: Host path
 - Volume: Docker managed resource
 - ... or external, database cluster



Container persistence

- **Bind mounts:**

- mount an existing path into a path inside the container

- **Example: mount *app* dir from host, into */var/www* inside the container**

- Usefulness: persistence or add shared/specific data into container (ex, web page, static resources)

```
$ docker run -d\  
  --mount source=/app,target=/var/www\  
  --name frontend-server  
  nginx
```

Container persistence

- **Volume mounts:**

- mount a specific storage object (managed by docker) at a given path
- volumes are not deleted with containers

- **Example: mount *app* volume, into */var/www* inside the container as read-only**

- Usefulness: persistence or add shared/specific data into container (ex, web page, static resources)

```
$ docker run -d\  
  --mount source=app:/var/www:ro\  
  --name frontend-server  
  nginx
```


Container Network

- **Containers run with specific network configuration**
 - No ingress network access added by default
- **Standard physical networks (others can be created)**
 - bridge: virtual switch
 - host: network only connected to host OS
 - none: isolated network (only with loopback, and can be customized)
- **Others**
 - overlay: used to provide connectivity over multiple engines (swarm mode)
 - macvlan: assigns specific MAC address to container

Container Network Ports

- **Services inside container cannot provide services to other applications**
 - No network services by default
 - **Run command must expose ports**
 - Outside ports mapped to inside ports
 - **E.g. map external port 8000 to internal (inside container) port 80**
- ```
$ docker run -d -p 8000:80 --name frontend-server nginx
```

# Sensitive data

- **Frequently, there is sensitive data that must be available to containers**
  - Database/ext service access credentials, certificates, admin credentials
  - Sensitive data is container specific, or shareable among several containers
- **Direct (no so correct) approach: Build images with credentials**
  - must build different images for each credential pair
  - must rebuild images when credentials change
  - credentials frequently end in the teams source control (e.g. Git repository)

# Docker Compose



- **Compose services created from multiple containers**

- that is... consider the specification of a complete application
- docker-compose.yml file

- **Example:**

- specifies service name: nginx
- container name and build ref
- volumes to add
  - uses an environmental variable
- ports to expose to outside

```
version: '3'
services:
 nginx:
 container_name: nginx
 build:
 context: ./nginx
 dockerfile: Dockerfile
 volumes:
 - app:/var/www
 - ${NGINX_LOG_PATH}:/var/log/nginx
 ports:
 - 80:80
 - 443:443
```

# Sensitive data: Docker Secrets

- **Allows creating blobs of data with restricted access**
  - Exist independently of the containers
  - Allow changing the secret without changing the container
    - ex: database access creds for dev, staging and prod
  - Only accessed inside specific containers
    - at `/run/secrets/secret_name`
- **Manage secret:** `docker secret create/inspect/list/rm`
- **Add to service:** `docker service... --secret secret_name`
  - Will expose `secret_name` inside container

# Docker Compose

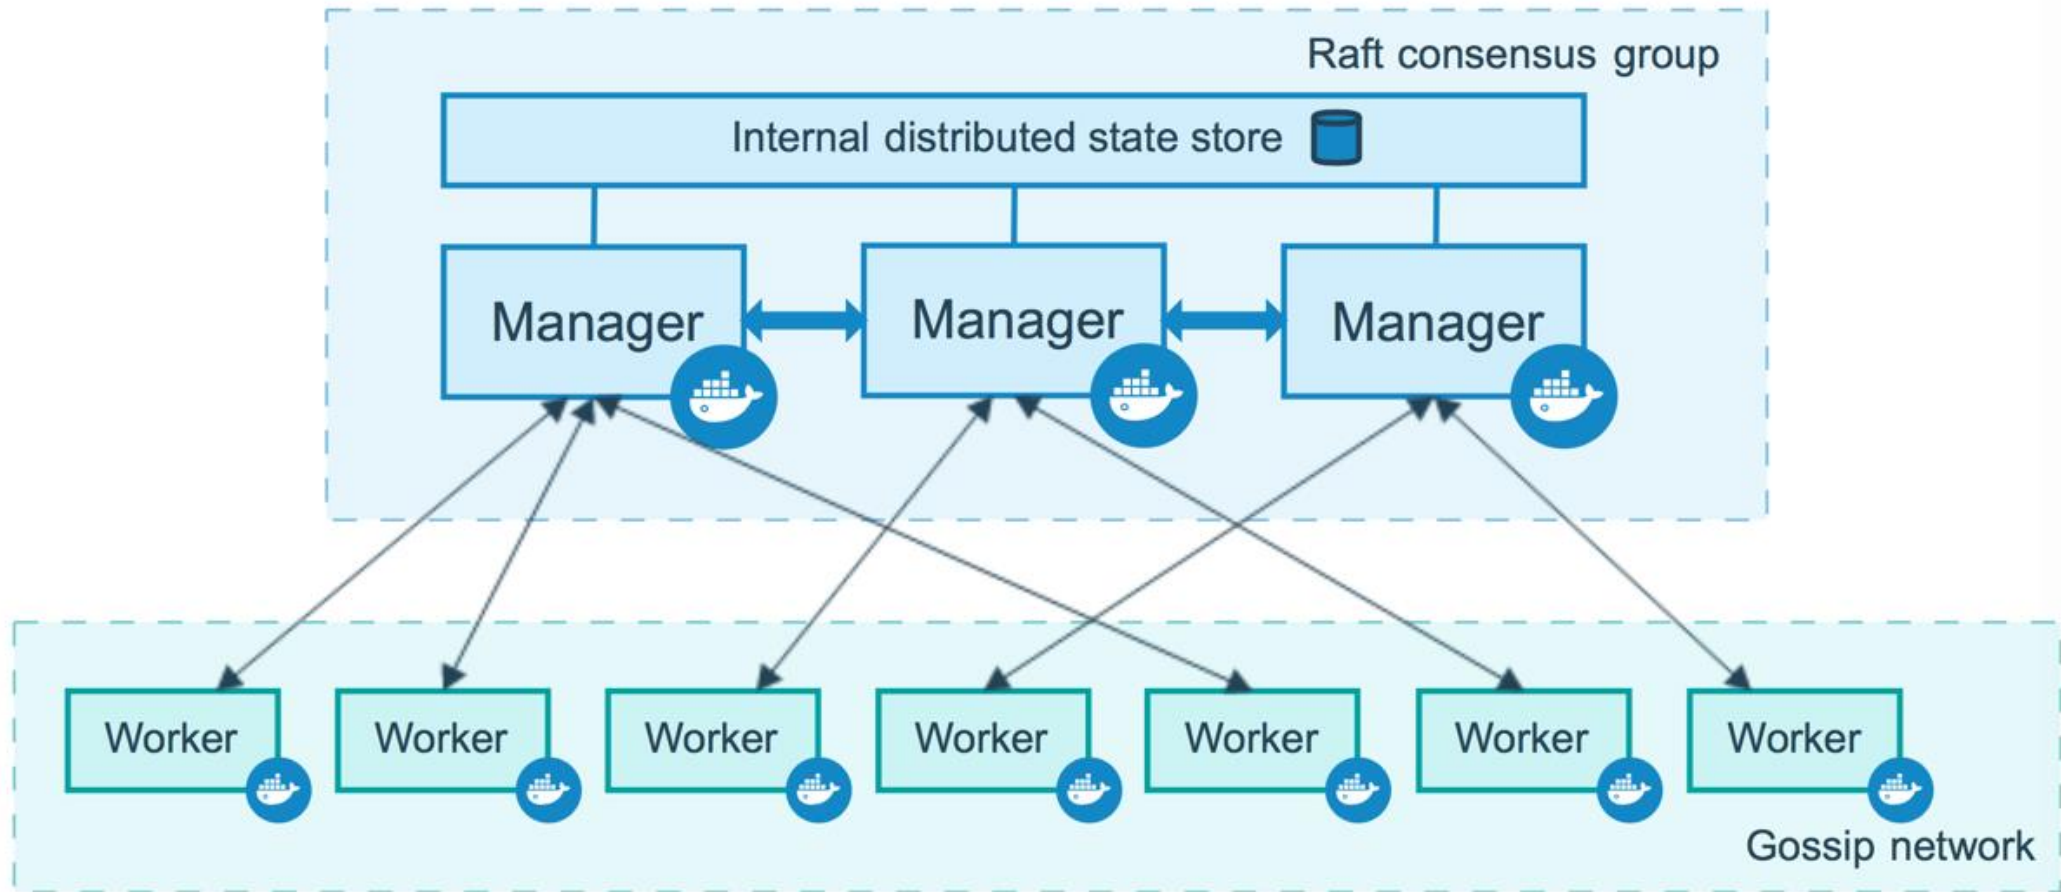


```
version: '3'
services:
 nginx:
 container_name: nginx
 build:
 context: ./nginx
 dockerfile: Dockerfile
 ...
 secrets:
 key_file:
 file: server-key.pem
 cert_file:
 file: server-cert.pem
```

# Docker Swarm Mode

- **Docker engine mode allowing integrated management of services over a cluster**
  - Instead of running containers at a host, containers run in a cluster
  - Managed automatically by docker
- **Swarm masters: manage the services running in the cluster (and execute containers)**
  - Act as traffic ingress points, and will accept workers (using auth)
  - Cluster may have multiple masters
- **Swarm workers: execute containers**
  - Must join a swarm master explicitly
  - Authentication is required

# Docker Swarm Mode





# Docker Swarm Features

- **Decentralized Cluster Management**
  - with multiple masters
  - internal consensus to keep the cluster operational
- **Declarative service model**
  - Manage stacks with services instead of individual containers

Create 15 instances of a Wordpress service:

```
$ docker service create --replicas 15 --name wpinst wordpress
```

Instances are distributed to the available nodes (managers and workers)

# Docker Swarm Features

- **Scaling**

- Possible to define the number of replicas
- Managers will distribute instances automatically
- Number can be modified with service running:
  - 10 instances (destroy 5): `docker service scale wpinst=10`
  - 100 instances (create 90): `docker service scale wpinst=100`

- **Desired state reconciliation:**

- Cluster is reconfigured to achieve the desired state required for the services
  - instances, network, ...

# Docker Swarm Features

- **Multi-host network:**

- Services can have networks associated
- Instances will see other instances of same service, even if running on different nodes
- Network is transparent and overlayed over existing infrastructure

```
$ docker service create --replicas 15 --network wpnet --name wpinst wordpress
```

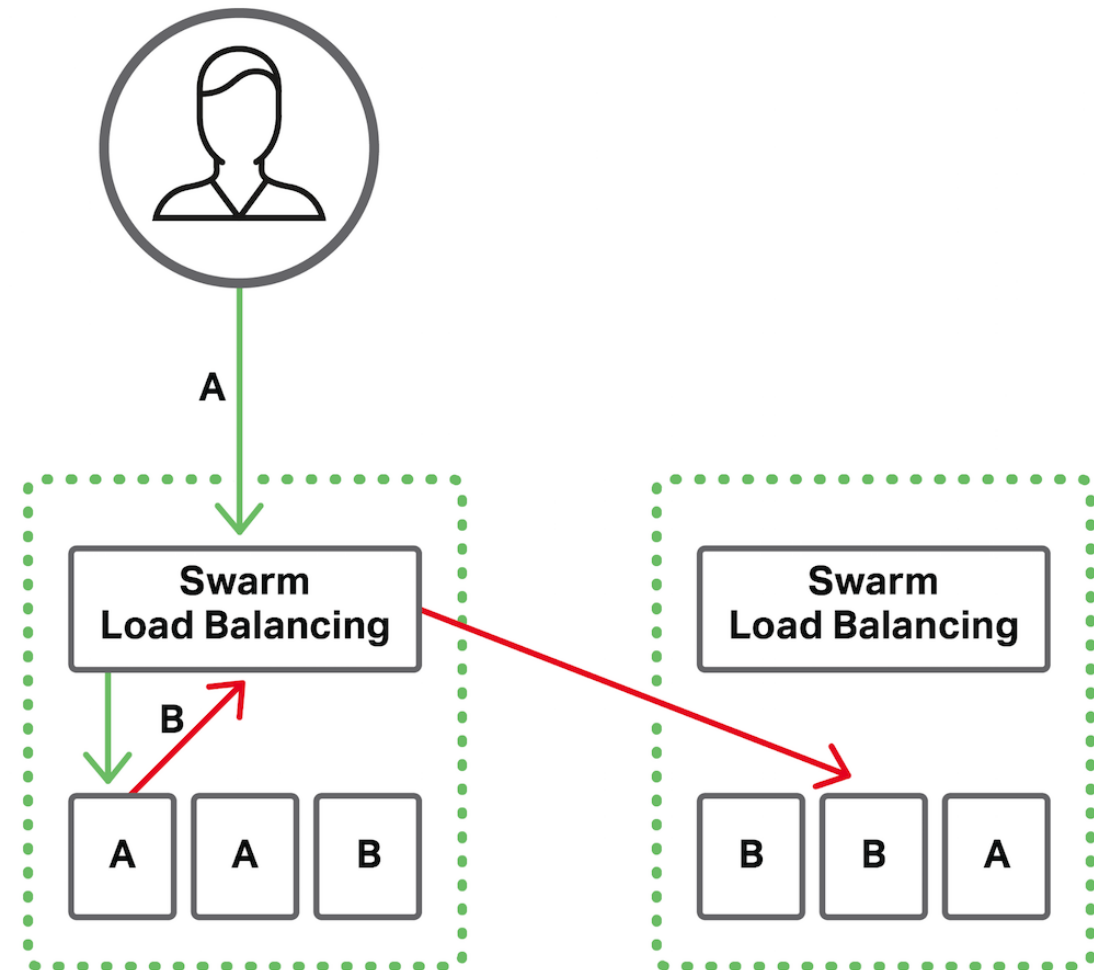
# Docker Swarm Features

- **Service discovery:**
  - Docker swarm implements a DNS service
  - All services are available by its name
- **Example: consider a WP service and a Database service**
  - In the WP instance, the Database connection string can refer to host “database” (no IP address required)
  - WP Instances will be provided with the IP address of the database container

# Docker Swarm Features

## Load balancing

- Ingress traffic is load balanced to the individual instances composing a service
- Single L3 port (e.g. TCP) expose to the outside
  - Each client connection will be forwarding to a different instance
    - from the pool of available instances (can be changed)
- Works for internal connections
  - Example: A=WPInstance, B=database





# Kubernetes



- A platform for orchestrating (Docker/Containerd) containers in a clustered environment, usually with multiple hosts
- Provides:
  - **container grouping**: groups containers in same environment
  - **load balancing** distributes load across physical resources
  - **auto-healing**: recover from failures (auto start containers)
  - **scaling features**: Increase number of containers as required by load
  - ...
- Motivation: running individual containers is not enough. Kubernetes is an orchestration solution that keeps services operating

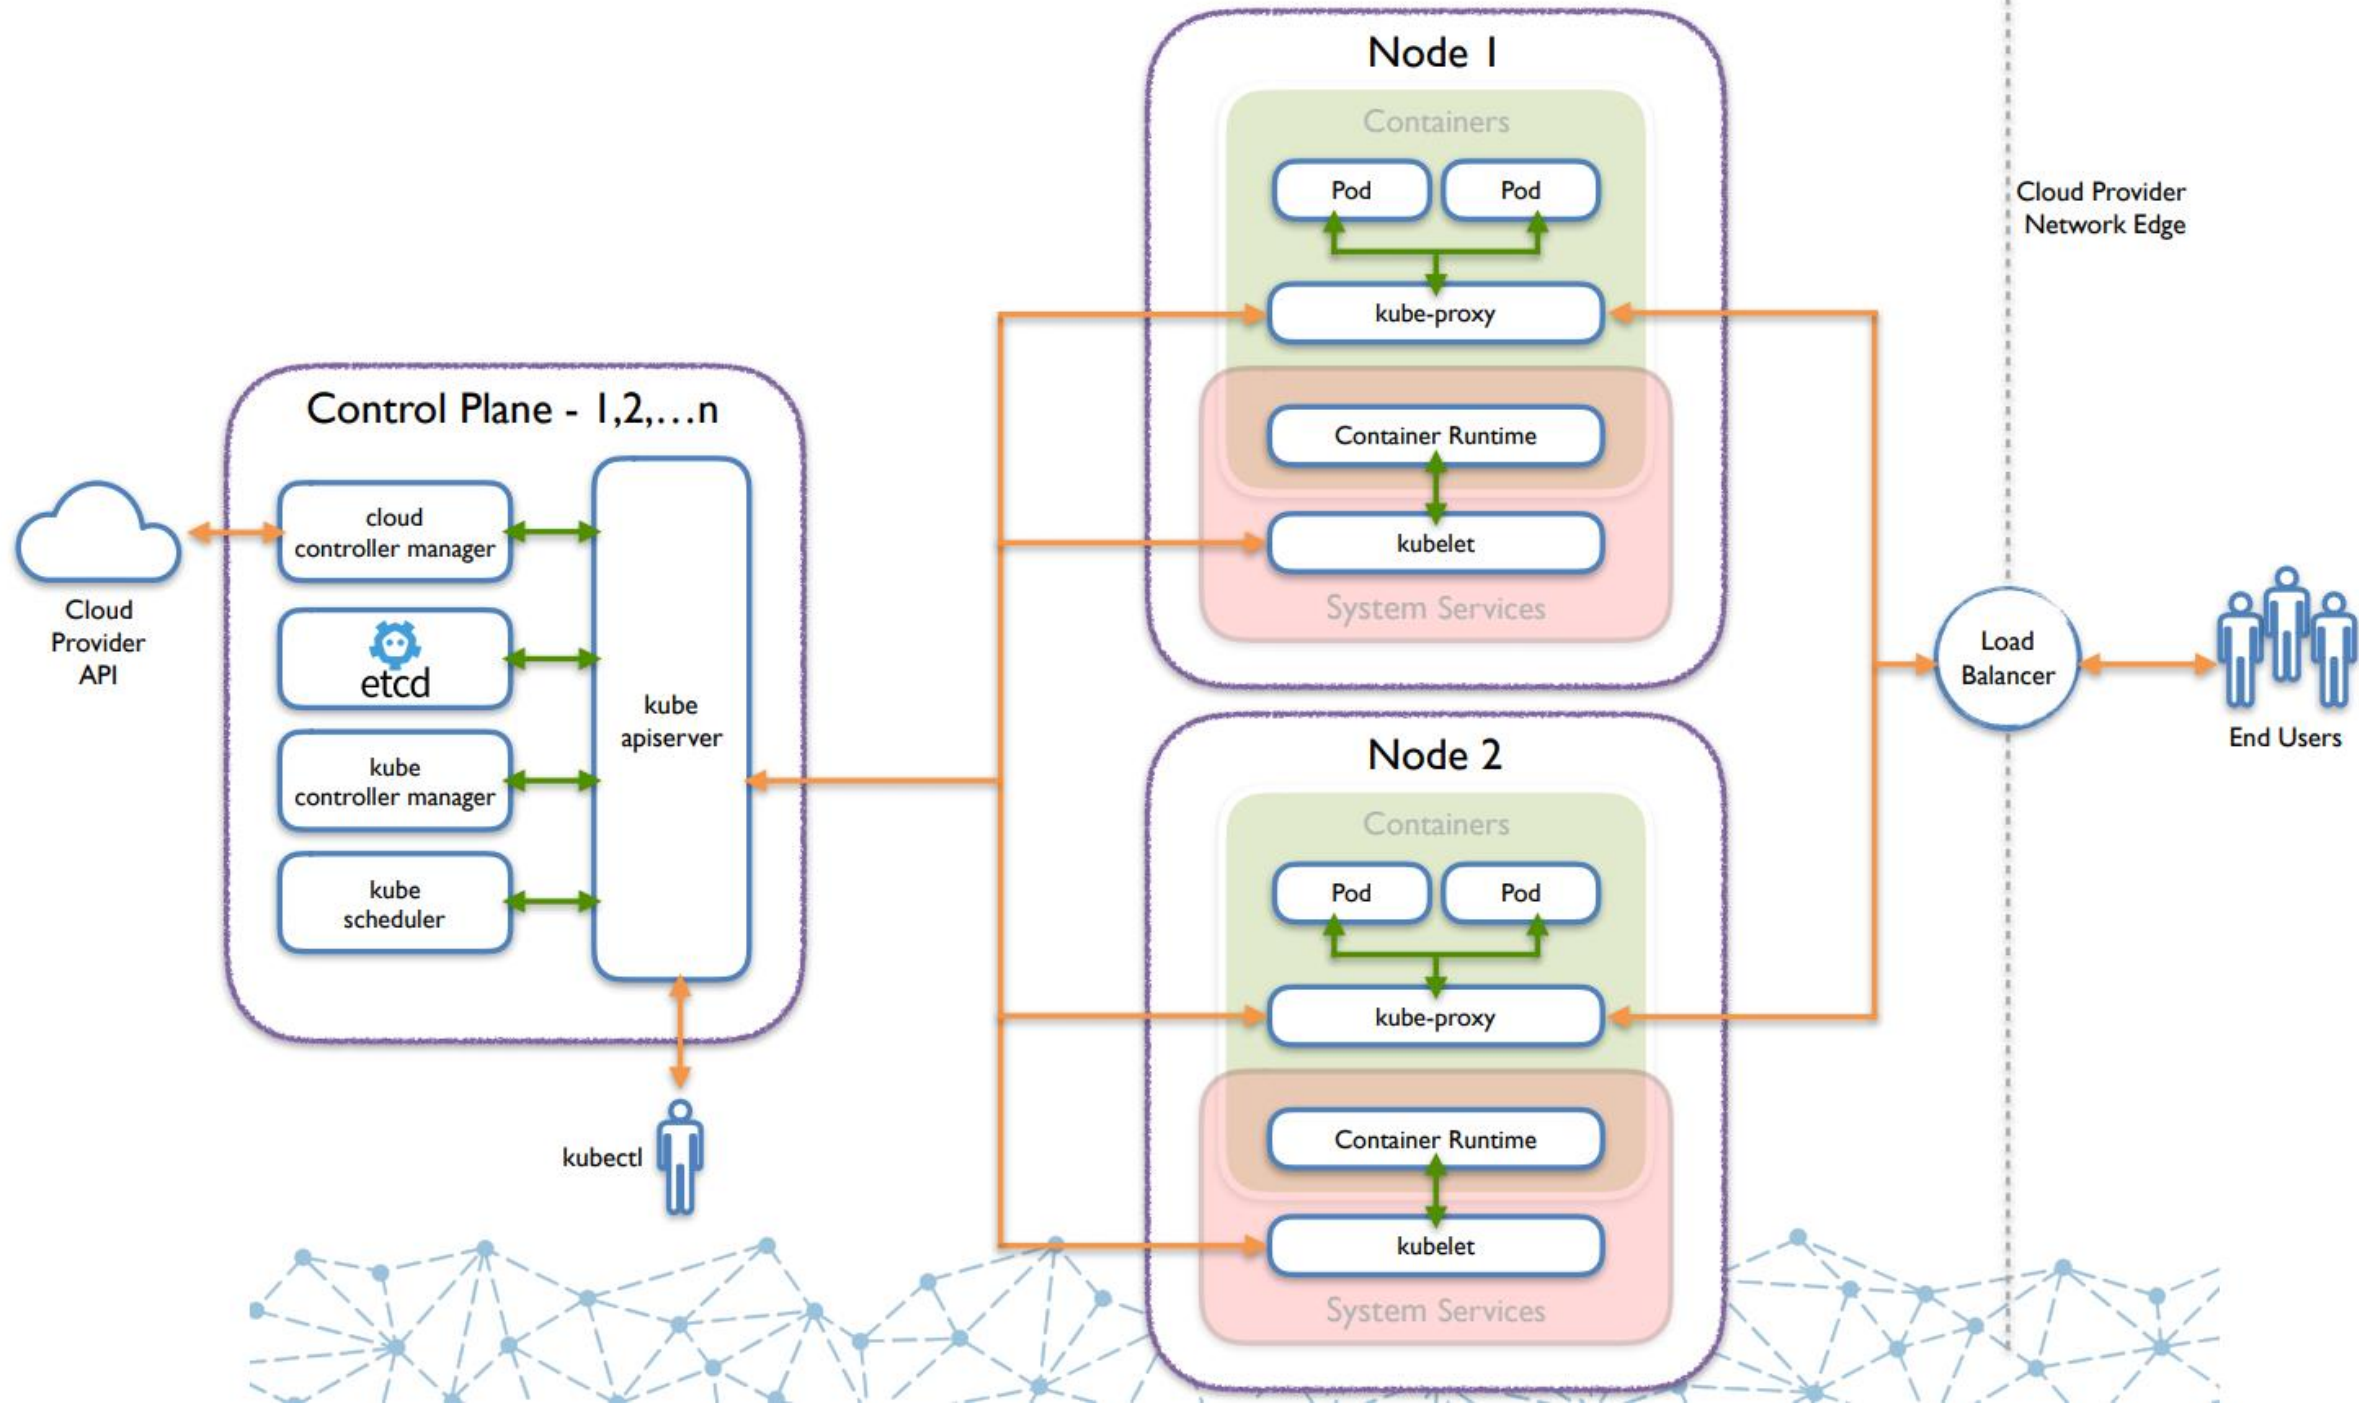
# Kubernetes

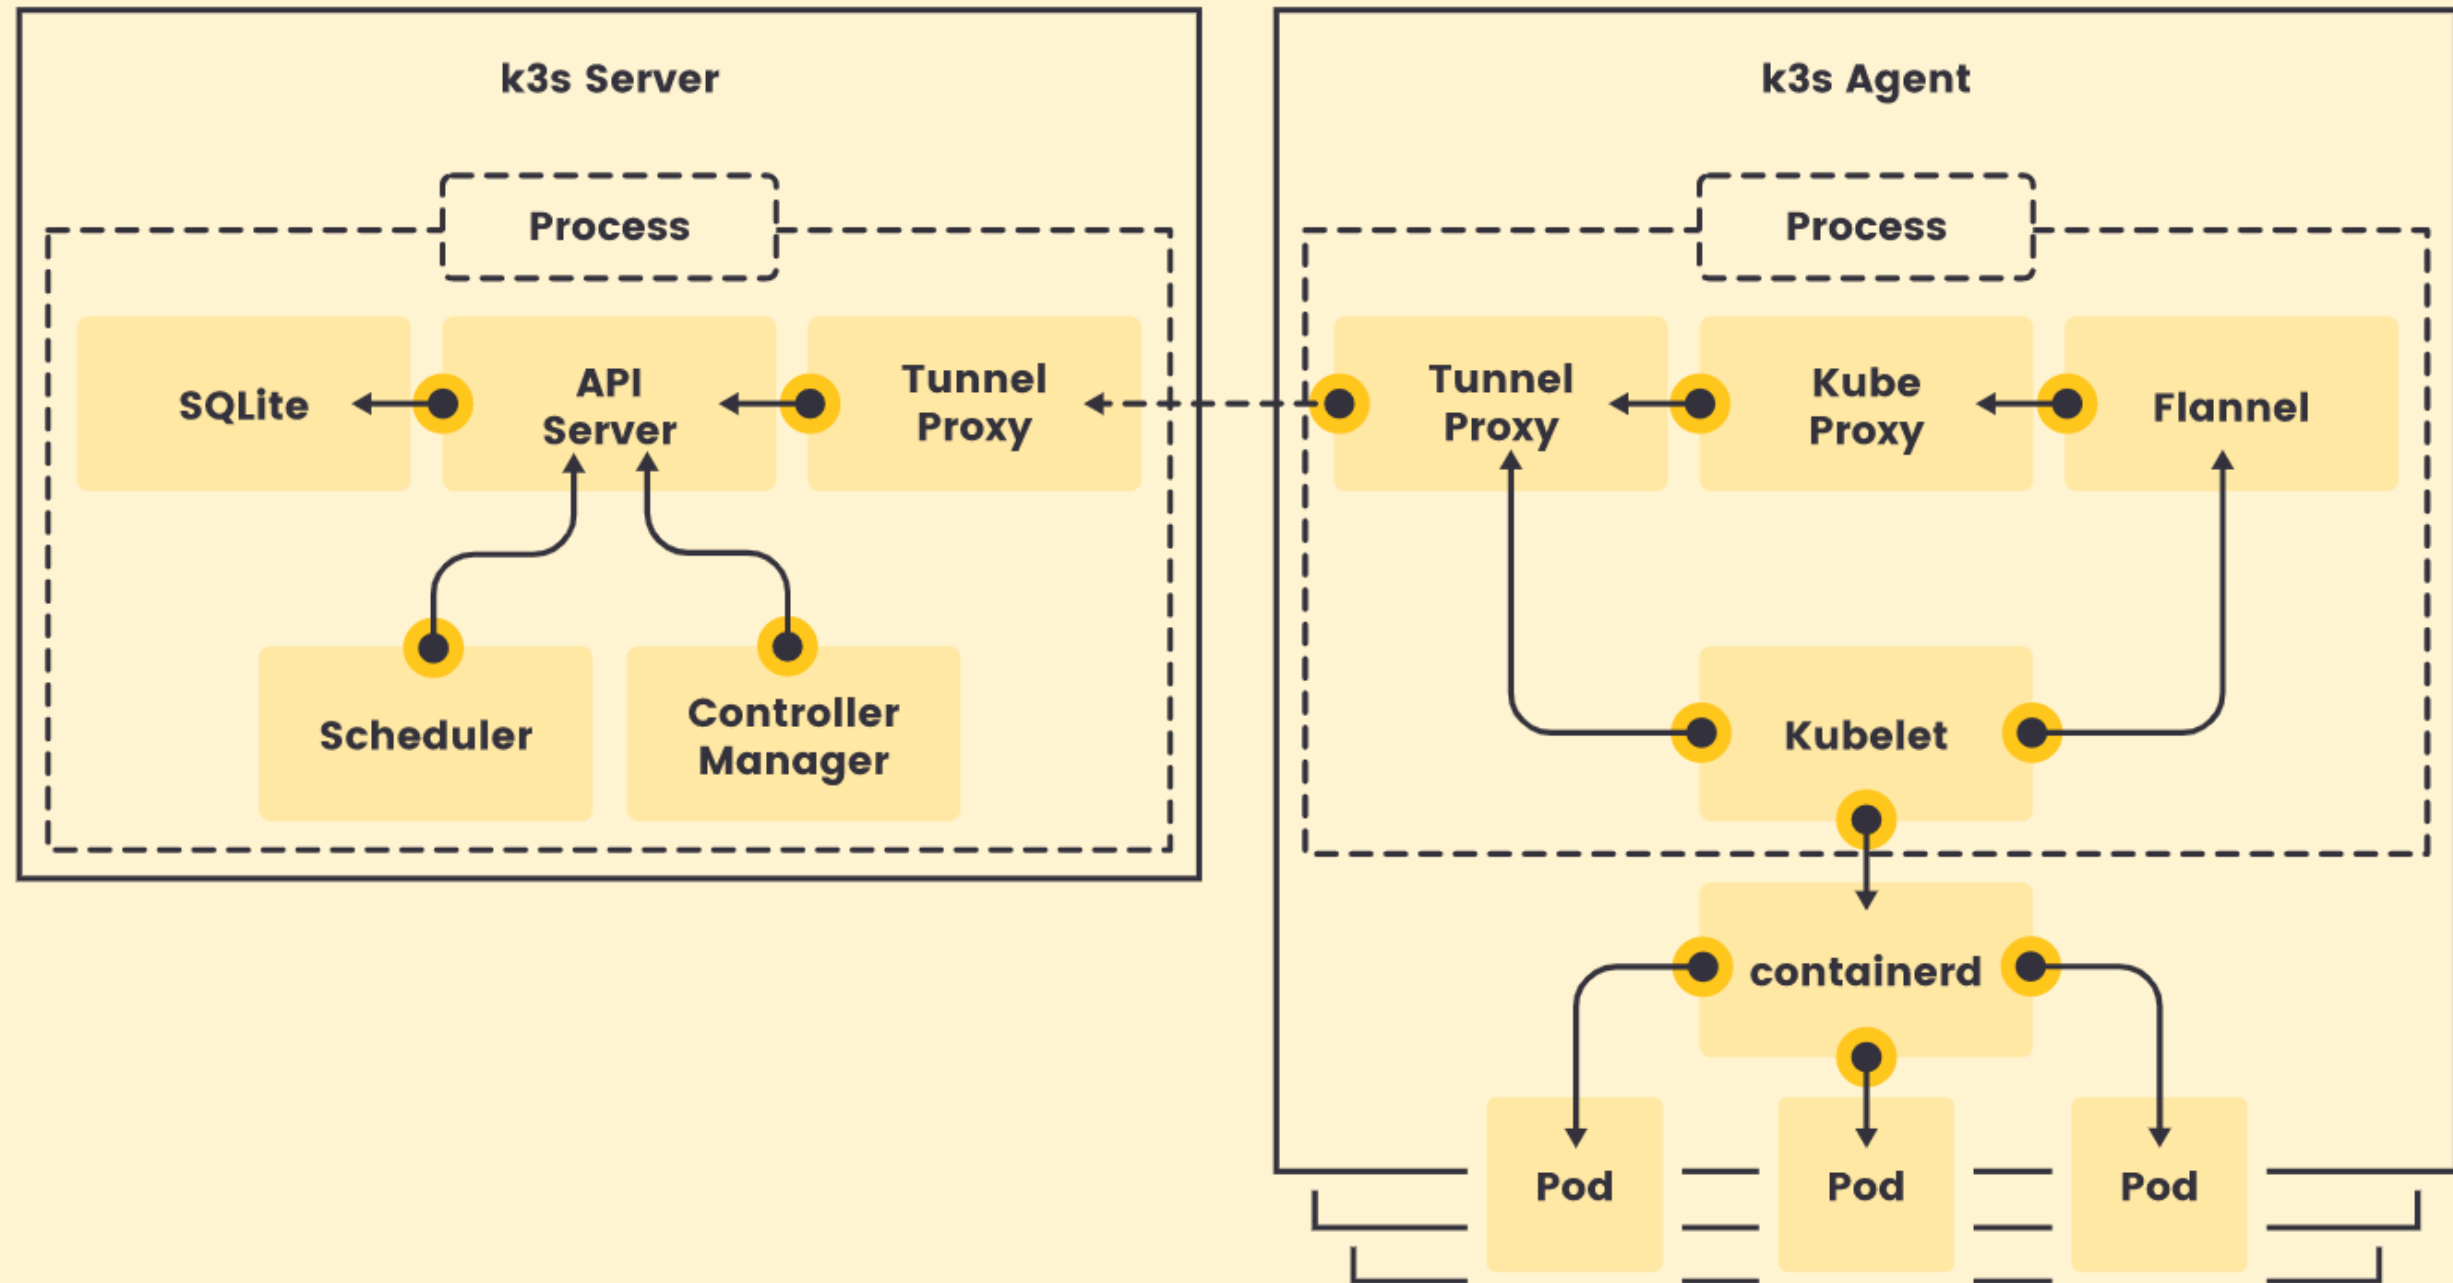
- **An orchestrator manages the lifecycle of an entity (VM, container, etc..)**
  - In this case: the lifecycle of containers
  - How it is created
  - Where it is instantiated
  - How it is balanced to new nodes
- **Orchestrators usually also manage related resources**
  - Data volumes
  - Secrets
  - Configuration values
  - External resources and services
- **Orchestrators are not simple automatic deployment tools**
  - They have an integrated view of the resources and make decisions to optimize placement



# Kubernetes

- **Focus on managing hundreds/thousands of containers, on a fleet of thousands of nodes**
  - While dealing with rollout of new versions, failure, security, ...
- **Implement a software defined platform: given the specification, the software will be kept running, as long as there are resources**





# Key software components

- **Kube-apiserver: provides a REST API for clients to interact with**
  - Control plane
  - Data Storage
- **Etcd: The Cluster Datastore, with a consistent and highly available key value storage**
  - For configuration data
- **Kube-Controller-Manager: Manages all component control loops**
  - Ensures the cluster is at the desired state
- **Kube-Scheduler: Places new Pods into Nodes**
  - Takes in consideration MANY scheduling decisions (affinity, anti-affinity, constraints, data locality....)

# Key software components

- **Kube-Proxy:** Manages network rules at each node, and does connection forwarding / load balancing
  - Similar to the docker proxy
- **Kubelet:** takes in consideration the description of each Pod (PodSpecs) and ensures (locally in the node) that the containers are running as expected
- **Container Runtime Engine:** A software that runs containers, many supported
  - Containerd (docker)
  - Rkt
  - cri-o
  - Virtlet
  - Kata

# Networking


- **Pods have an internal network namespace**
  - You have multiple pods with the same networks!
- **Pod Network: Cluster wide network for Pod-to-Pod communication**
- **Service Network: Virtual IPs (durable), managed by kube-proxy for service discovery**
  - Because Pod IP address are volatile, services provide anchor points for communication

# Namespaces

- **Logical Clusters inside the Real Cluster**
- **Provide the means for partitioning the cluster to multiple teams/projects**
- **Objects are created under a namespace and deletion of the namespace may also delete all objects**
  - Exceptions are the resources related StatefulSets (later)
- **Some namespaces are always present (default, kube-system,...)**
  - Typical namespaces: production, development, monitoring

# Pod

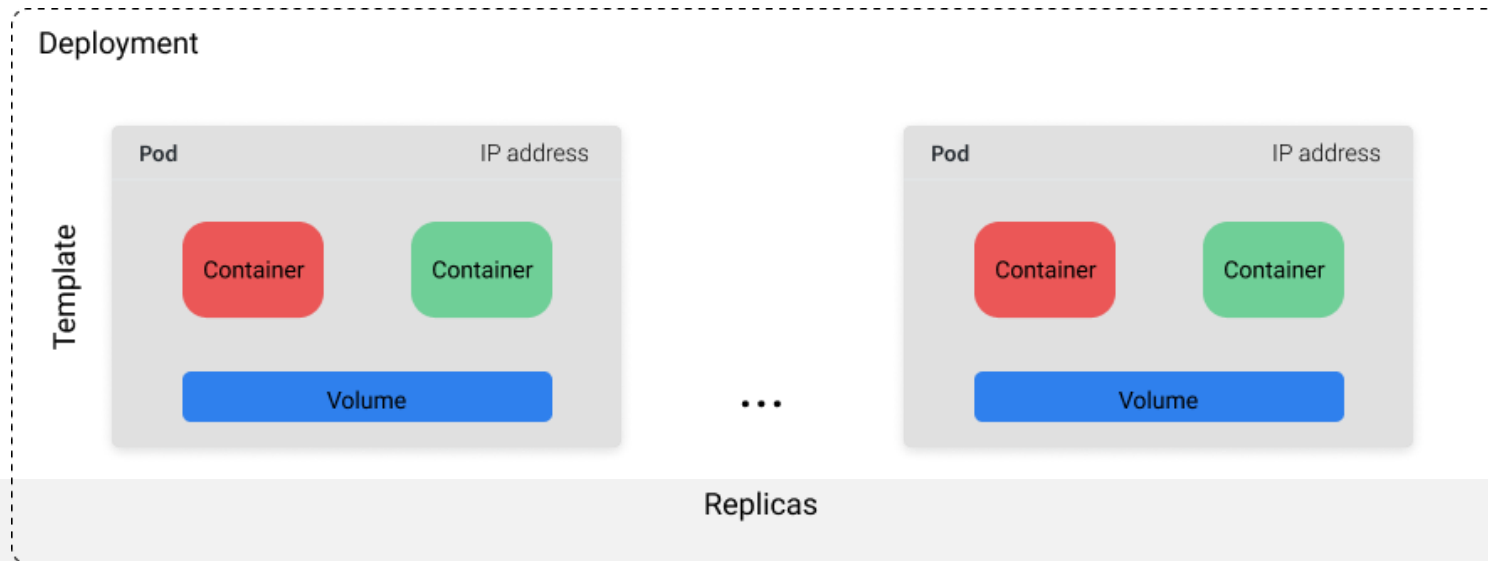
- A group of Containers. A basic management unit.
  - Includes information about what to run, and how to run something
  - Mimics a local host with applications (containers)
- Containers are ephemeral
- Containers are tagged to versions and we go changing version
- Pod contents are always co-located and co-scheduled





# Deployment

- **A set of pods that should operate together**
  - Can be a software, but also includes volumes, networks, secrets, etc..
- **Defined using a declarative format (yaml file) stating the desired state of the pods**



# Deployment

```
apiVersion: apps/v1
kind: Deployment
metadata:
 name: app
 namespace: diving-into-k3s
spec:
 replicas: 1
 selector:
 matchLabels:
 app: app
 template:
 metadata:
 labels:
 app: app
 spec:
 containers:
 - name: app
 image: 10.110.0.3:5000/app:v1
 resources:
 requests:
 memory: "32Mi"
 cpu: "10m"
 limits:
 memory: "128Mi"
 cpu: "500m"
 ports:
 - containerPort: 8080
```

What is the context

How to find pods in an infrastructure

How many instances

What to run. In this case, a Docker container

Resource limits

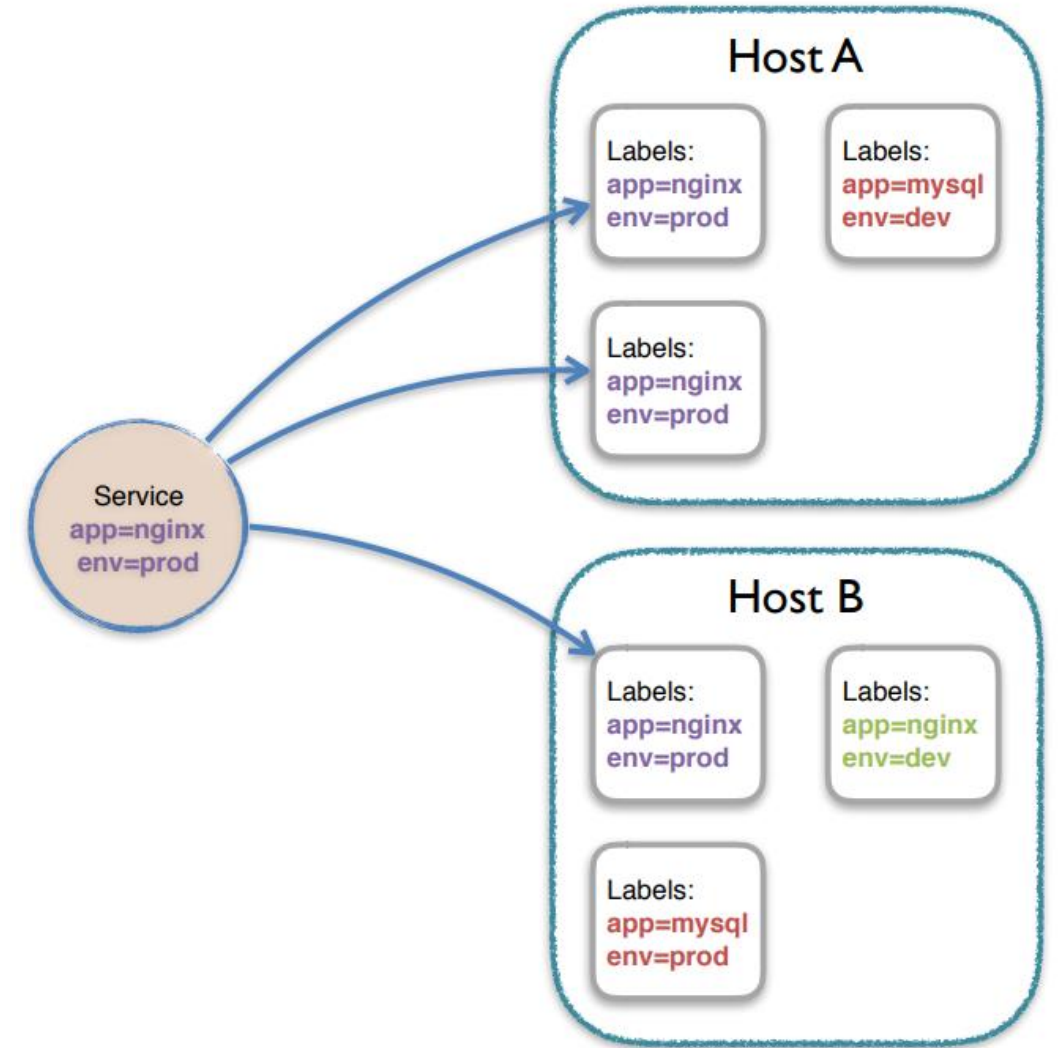
How to expose

# StatefulSet

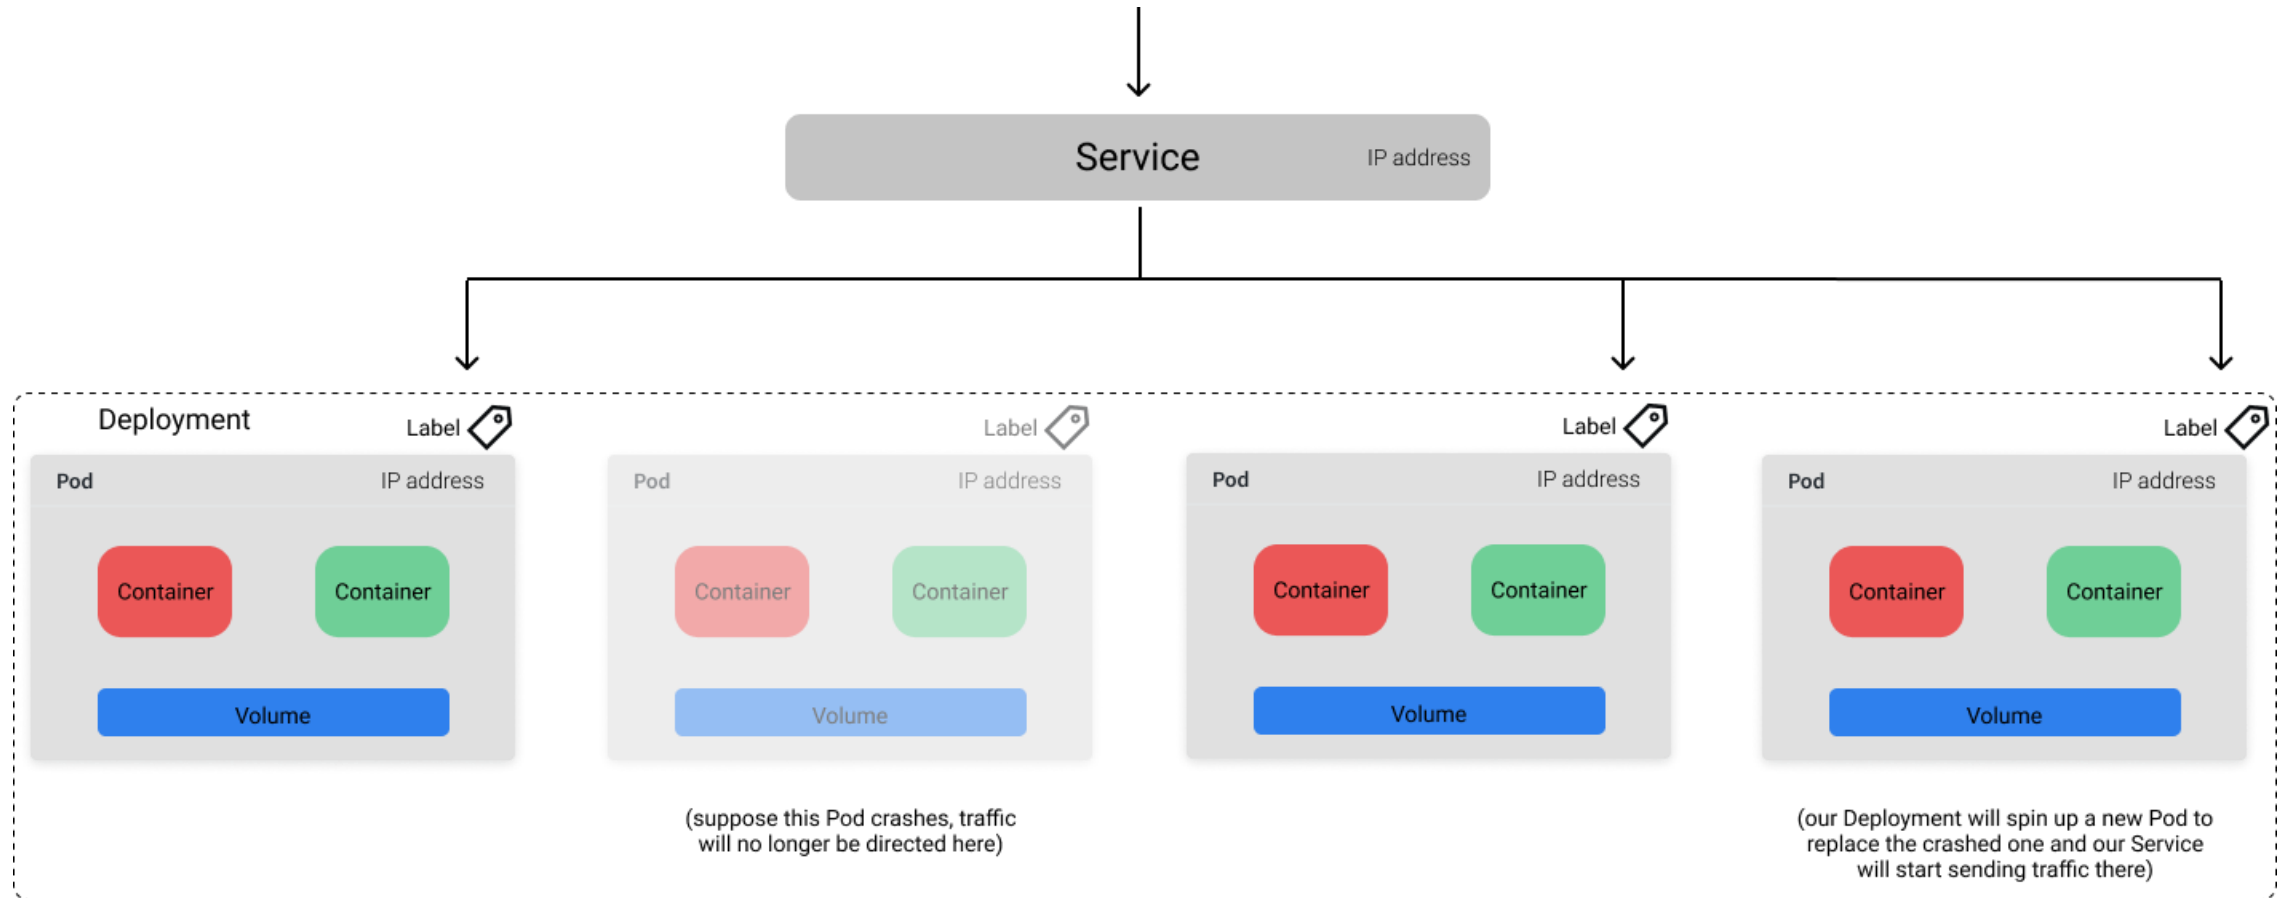
- **Similar to a deployment but Pods are stateful**
  - Have persistent and unique IDs
  - Are started by a specific order
  - Storage is persistent
    - Requires a `StorageClass` to actually store data
  - Graceful deployment and scaling
- **Useful for stateful software, whose instances have unique roles**
  - e.g. a Database

# Service

- K8s abstracts internal aspects from users, such as the IP addresses of replicas
  - Also, number and location of said replicas
- A service provides the visibility to reach a deployment from inside the cluster
  - Allows Pods to reach other pods
  - Service provides a stable endpoint, no matter how many pods are available (or where)
- Persistent
  - Static cluster IP
  - Static DNS Name

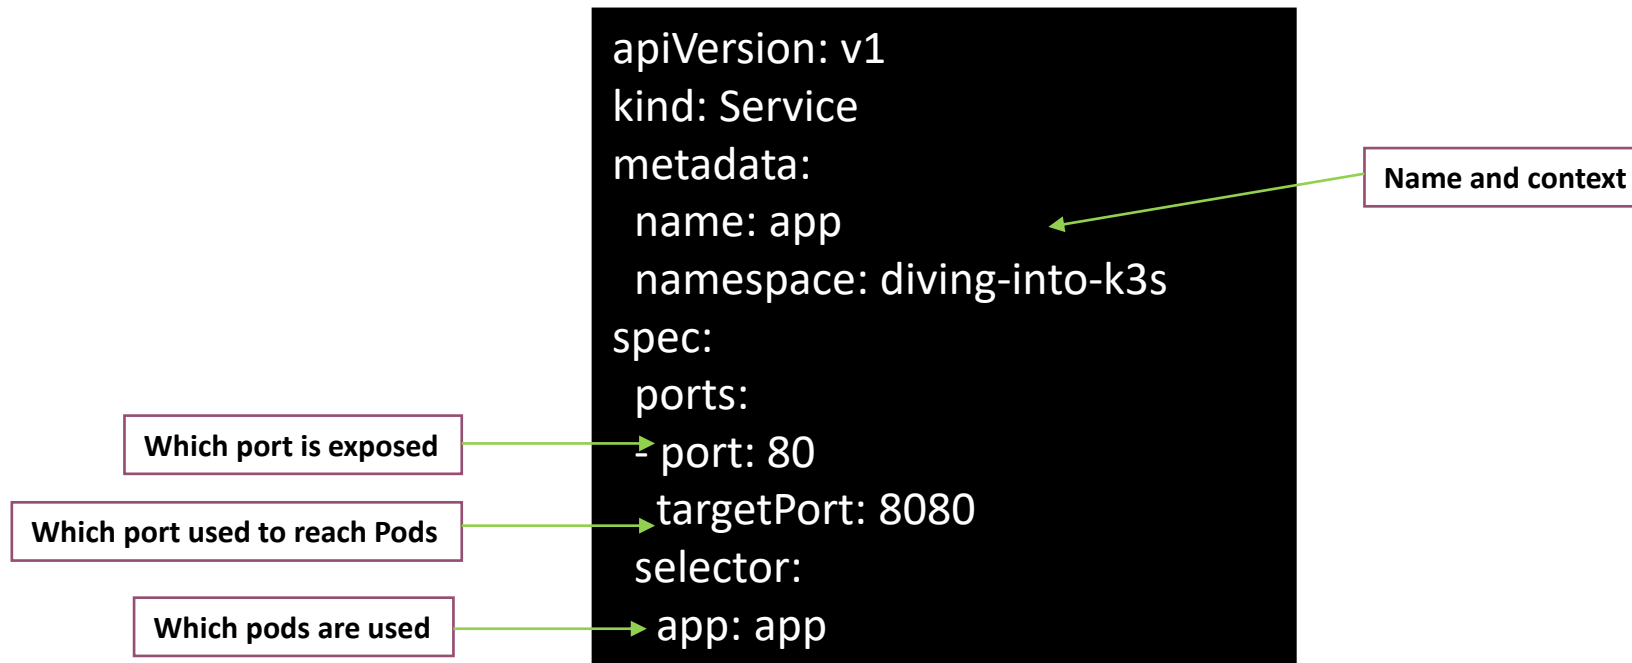


# Service



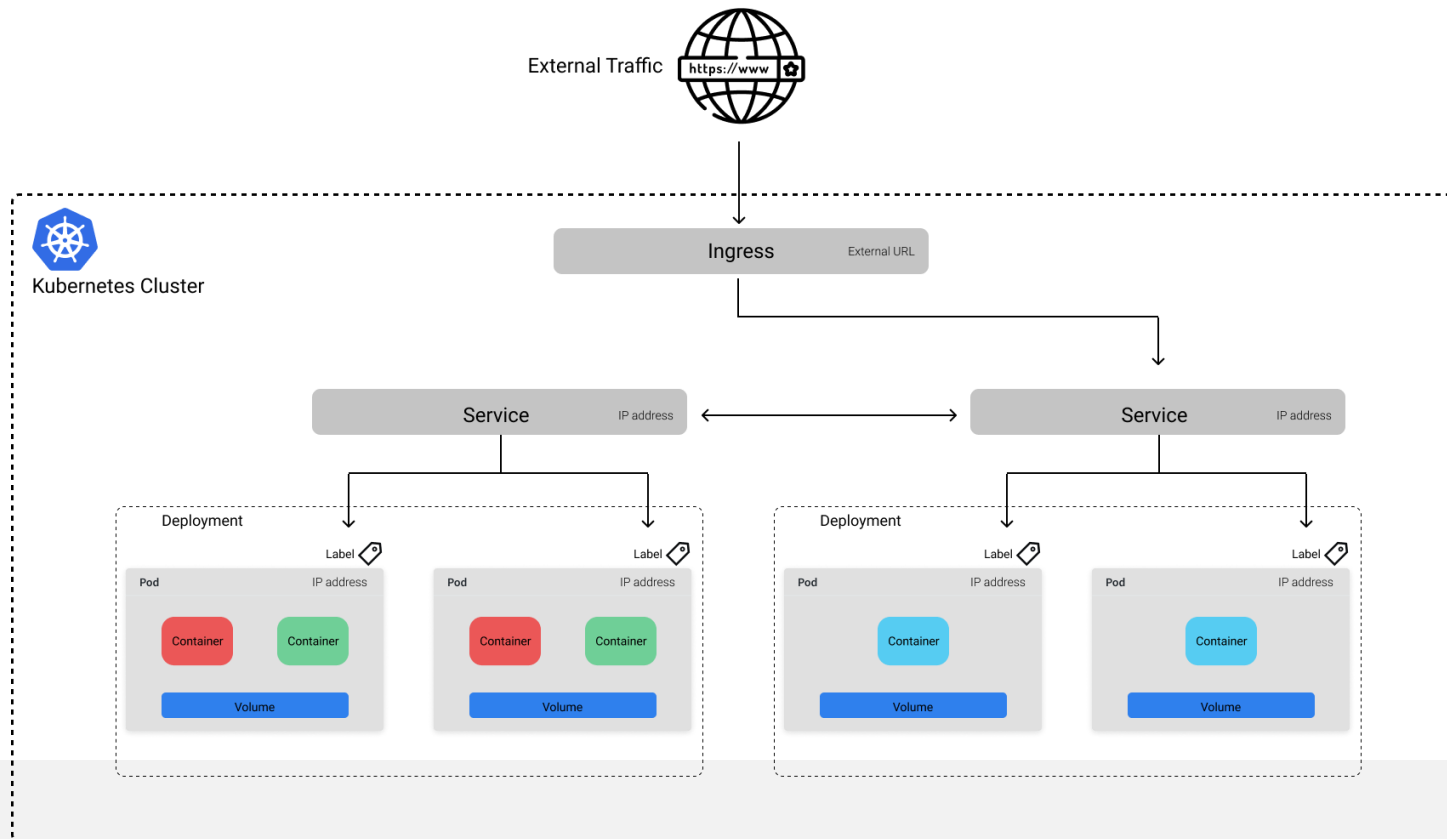
# Service

- A service named “app” allows Pods to access other pods by DNS, through the service
  - e.g. curl <http://app> would be redirected to one pod (round robin)



# Ingress

- Similar to a Service, makes some services visible to the outside world
- Frequently uses a proxy software (nginx, Traefik, apache2)



# Ingress with Traefik

Name and Namespace

Traefik Specific Configuration for This ingress

```

apiVersion: networking.k8s.io/v1beta1
kind: Ingress
metadata:
 name: app-k3s
 namespace: diving-into-k3s
 annotations:
 kubernetes.io/ingress.class: traefik
 traefik.ingress.kubernetes.io/frontend-entry-points: http,https
 traefik.ingress.kubernetes.io/redirect-entry-point: https
 traefik.ingress.kubernetes.io/redirect-permanent: "true"
spec:
 rules:
 - host: app.k3s
 http:
 paths:
 - path: /
 backend:
 serviceName: app
 servicePort: 80
```

How to expose the Pods (the hostname and port)

What Pods to use (will use round robin)



# Job

- An ephemeral execution to run one time
- In contrast to a pod in a deployment, which are expected to be running, jobs are one-time tasks.
- Kubernetes ensures the job runs as stated, dealing with availability of nodes

```
apiVersion: batch/v1
kind: Job
metadata:
 name: model-train-job
spec:
 template:
 spec:
 containers:
 - name: ml-model-train
 image: ml-development:1.2.1
 command: ["python", "train.py"]
 restartPolicy: Never
 backoffLimit: 5
```

How should we create the Pod for this Job?

How many times should we try to complete this Job before quitting?

# ConfigMaps

- **Provide the means to create generally available configurations for entities**
  - Used for non-confidential data
  - Key-Value pairs
  - Can be injected using configuration files in volumes, command line arguments or environmental variables
- **Allows decoupling the software from the execution environment specific configuration**
  - Usually, different environments (dev and prod?) may have different configurations
- **ConfigMaps are created, and then attributed to a Pod in the deployment**
  - Only the required Pods will access the configuration

# ConfigMaps

A subPath (the Item) in a Volume is mapped to the Containers.

The config file (nginx.conf) will be available at /etc/nginx/nginx.conf

ConfigMap Items are mapped to a Volume.

Definition of the ConfigMap  
(notice that it is restricted to the namespace)

ConfigMap has a name, and items inside it

```
Create nginx deployment

...
containers:
 - image: nginx:alpine
 name: nginx
 ports:
 - containerPort: 80
 resources: {}
 volumeMounts:
 - name: nginx-conf
 mountPath: /etc/nginx/nginx.conf
 subPath: nginx.conf
 readOnly: true
 volumes:
 - name: nginx-conf
 configMap:
 name: nginx-conf
 items:
 - key: nginx.conf
 path: nginx.conf

apiVersion: v1
kind: ConfigMap
metadata:
 name: nginx-conf
 namespace: diving-into-k3s
data:
 nginx.conf: |
 ...content of the nginx.conf configuration
```

# Secrets

- **Similar to ConfigMaps, but aims to store a small, sensitive information blob (password, token, key)**
  - By default are stored as a base64 string, with access limited only by the API
  - Kubernetes allows both Encryption and Role Based Access Control (RBAC)
    - Secrets will be strongly encrypted and only provided to allowed users
  - Differ from ConfigMaps as:
    - They can only be used by kubelet, an environmental variable or a volume
      - Cannot pass it through arguments
    - They have a type, defining its purpose

# Secrets

- **The Opaque type has no restrictions on content**
  - Meant for arbitrary user data loaded from a file, or specified interactively
- **The remaining types impose specific content format**

| Builtin Type                        | Usage                                 |
|-------------------------------------|---------------------------------------|
| Opaque                              | arbitrary user-defined data           |
| kubernetes.io/service-account-token | service account token                 |
| kubernetes.io/dockercfg             | serialized ~/.dockercfg file          |
| kubernetes.io/dockerconfigjson      | serialized ~/.docker/config.json file |
| kubernetes.io/basic-auth            | credentials for basic authentication  |
| kubernetes.io/ssh-auth              | credentials for SSH authentication    |
| kubernetes.io/tls                   | data for a TLS client or server       |
| bootstrap.kubernetes.io/token       | bootstrap token data                  |